

Big Data Infrastructures

Philippe Cudré-Mauroux

Fall 2018

Lecture 2 - Storage and Indexing

Outline

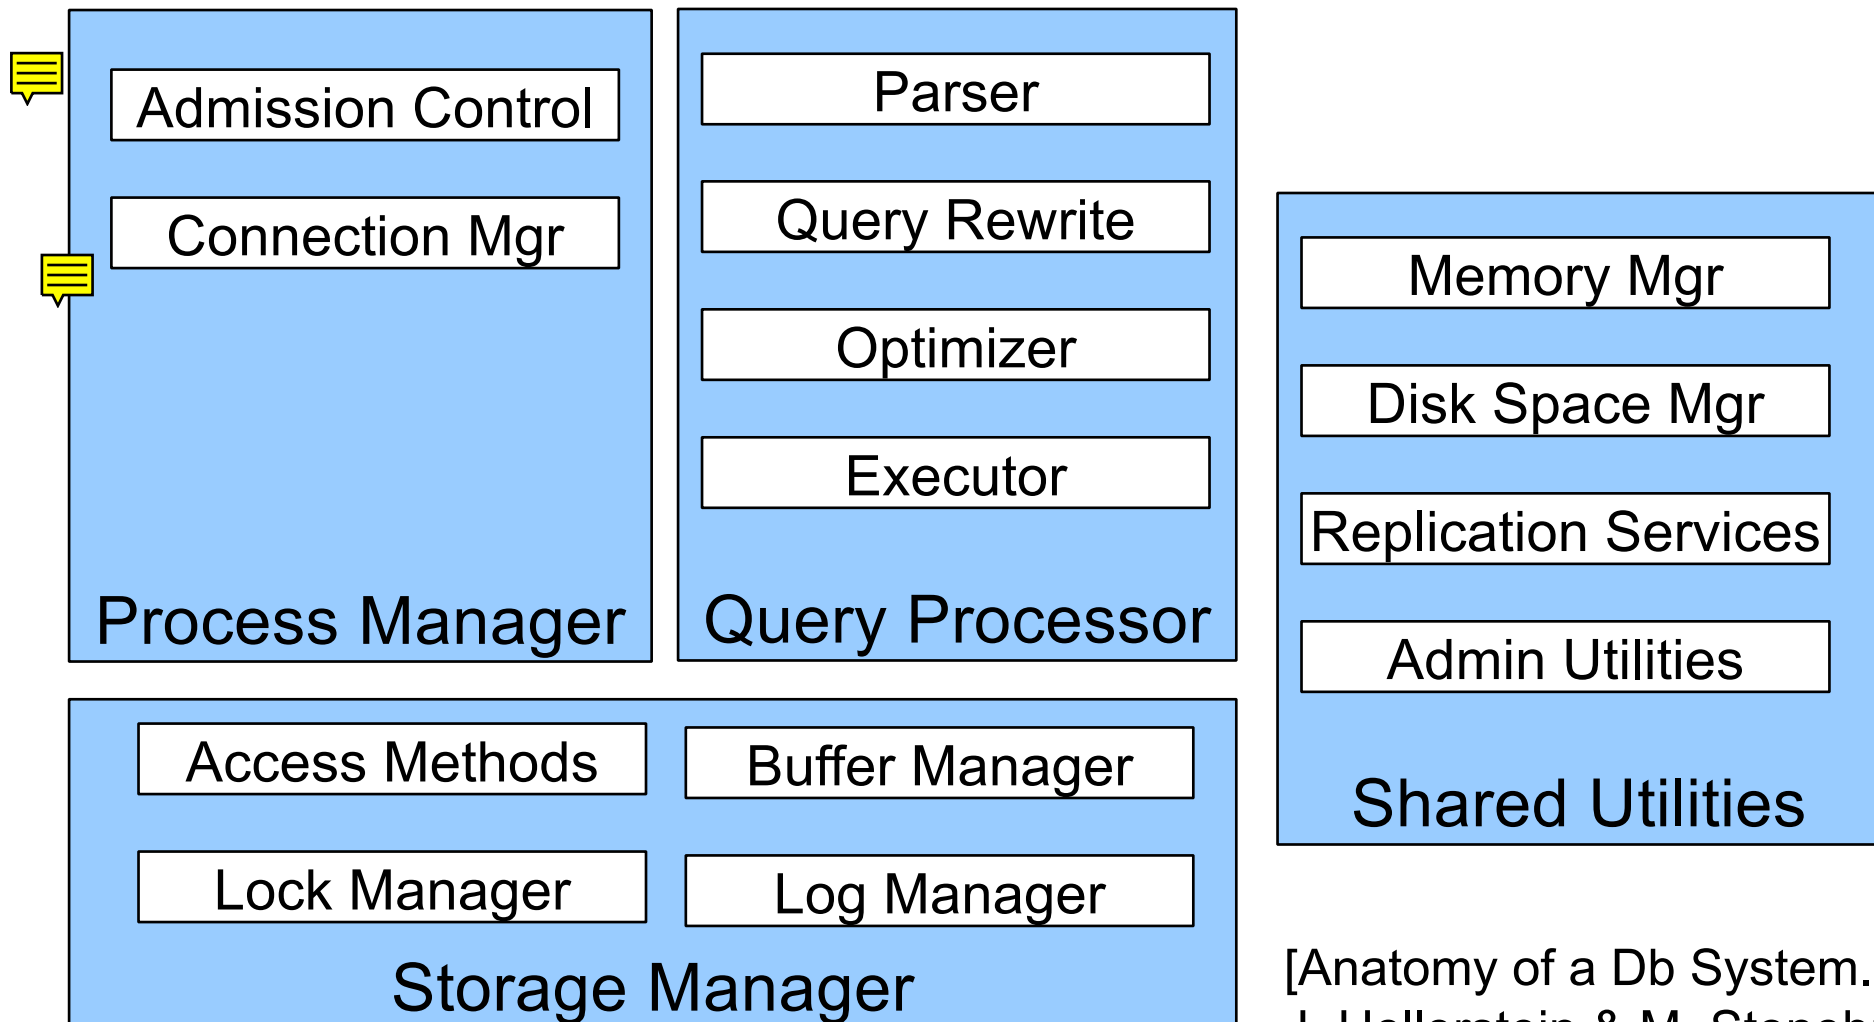
- **Case-Study**
- **Data storage**
 - Disk and files
 - Operations on files
- **Indices**
 - Index structures
 - Hash-based indices
 - B+ trees

Why is this important?

- **Data storage**
- **Indexes**

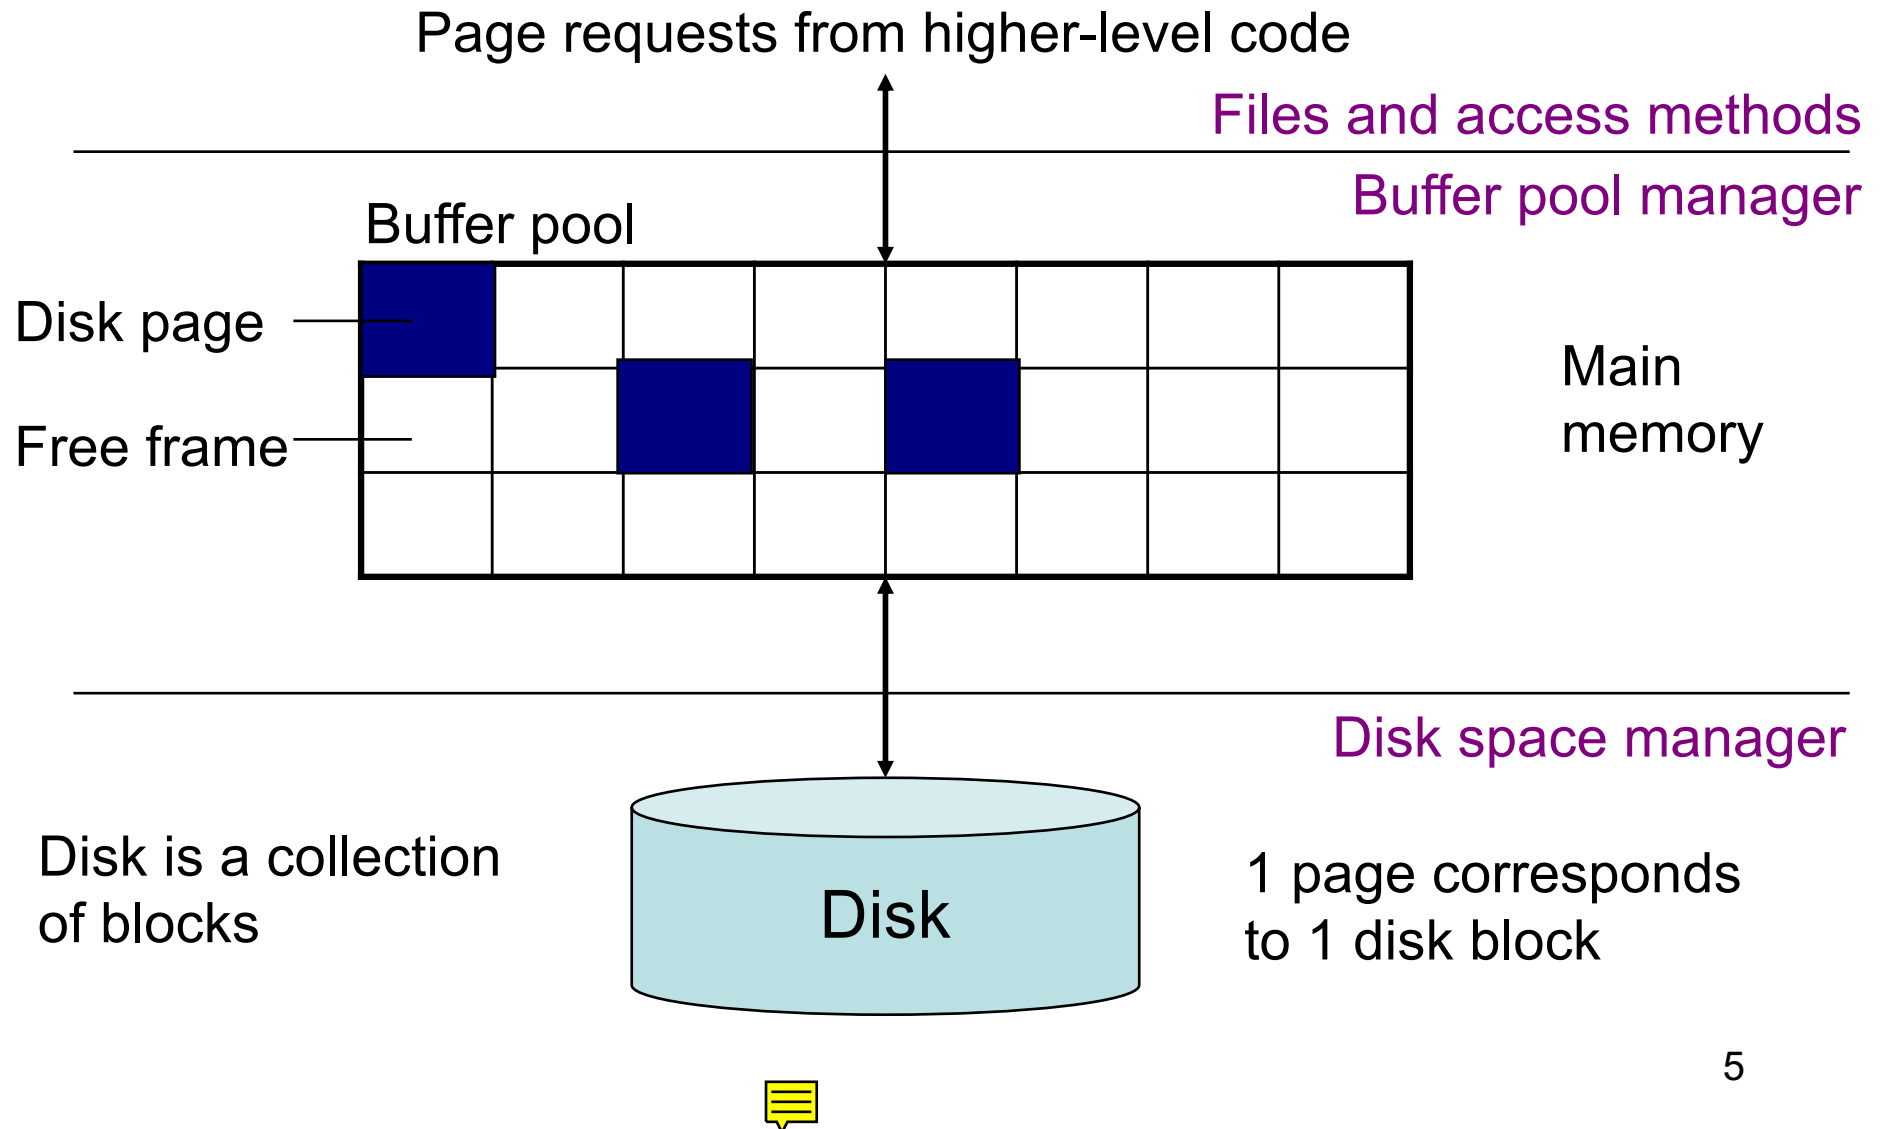


DBMS Architecture




[Anatomy of a Db System.
J. Hellerstein & M. Stonebraker.
Red Book. 4ed.]

Buffer Manager



Data Storage

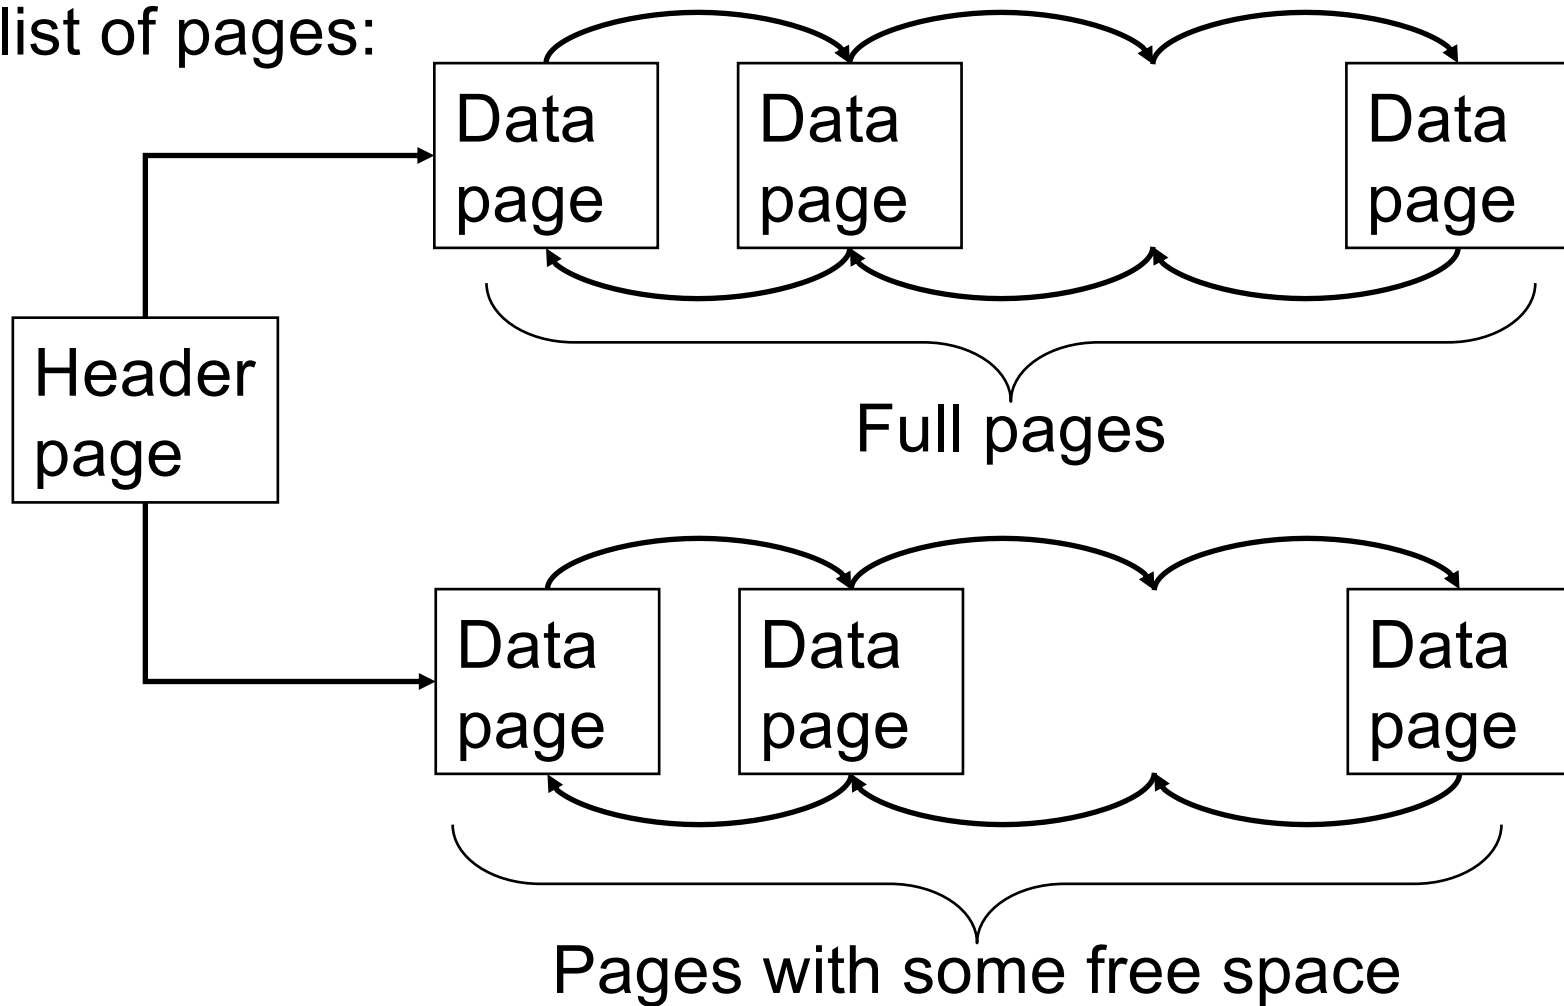
- Basic **abstraction**
 - *Collection of records or file*
 - Typically, 1 relation = 1 file
 - A file consists of *one or more pages*
- How to organize pages into files?
- How to organize records inside a file?
- Simplest approach: **heap file** (unordered) 
 - Further approaches: clustered file or sorted file

Heap File Operations

- **Create** or **destroy** a file
- **Insert** a record
- **Delete** a record with a given record id (rid)
 - rid: unique tuple identifier
 - can identify disk address of page containing record by using rid
- **Get** a record with a given rid
- **Scan** all records in the file

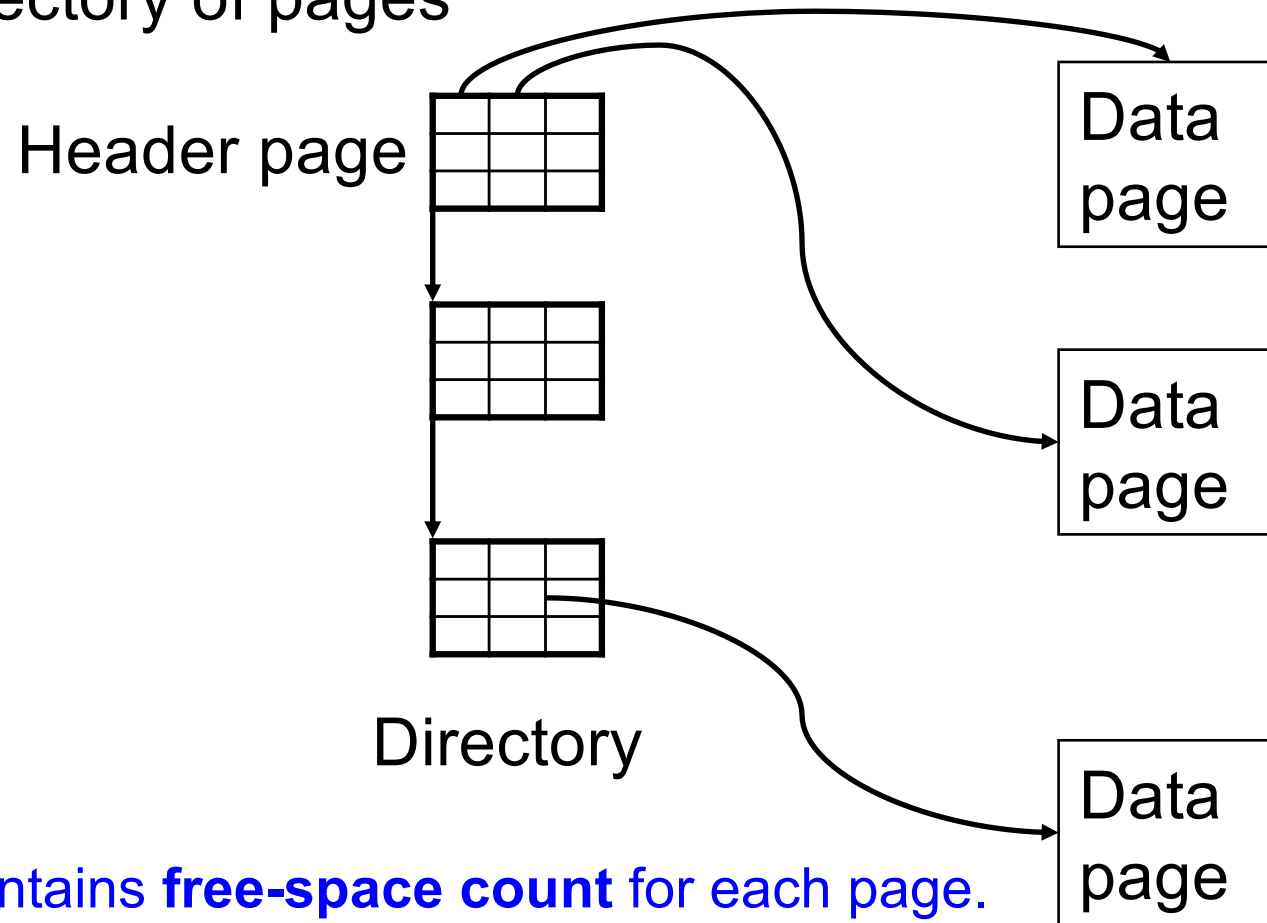
Heap File Implementation 1

Linked list of pages:



Heap File Implementation 2

Better: directory of pages



Directory contains **free-space count** for each page.
Faster inserts for variable-length records

Page Formats

Issues to consider


- 1 page = 1 disk block = fixed size (e.g. 8KB)
- Records:
 - Fixed length
 - Variable length
- Record id = RID
 - For example RID = (PageID, SlotNumber)

Why do we need RIDs in a relational DBMS?

Is the RID typically known to the end-user?

Page Format Approach 1

Fixed-length records: packed representation




Problems ?

Number of records

How to handle variable-length records?

Need to move records for each deletion, changing RIDs

Page Format Approach 2



Can handle variable-length records

Can move tuples inside a page without changing RIDs

Record Formats


Fixed-length records → Each field has a fixed length
(i.e., it has the same length in all the records)

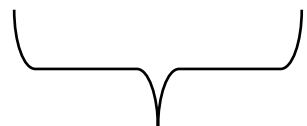
Field 1	Field 2	Field K
---------	---------	-----	-----	---------

Information about field lengths and types is in the catalog

Record Formats


Variable length records




Record header

Remark: NULLS require no space at all (why ?)

Long Records Across Pages



- When records are very large
- Or even medium size: saves space in blocks
- Commercial RDBMSs avoid this

LOB

- Large objects
 - Binary large object: BLOB
 - Character large object: CLOB
- Supported by modern database systems
- E.g. images, sounds, texts, etc.
- Storage: attempt to cluster blocks together

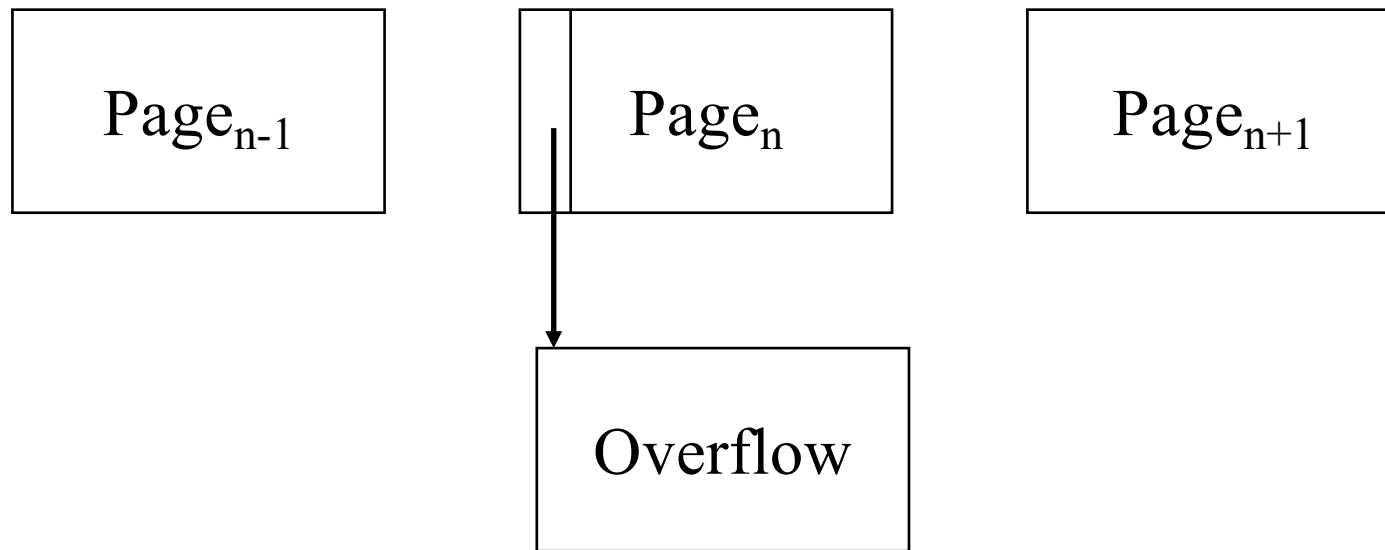
Outline

- **Data storage**
 - Disk and files
 - Operations on files
- **Indexes**
 - Index structures
 - Hash-based indexes
 - B+ trees

Modifications: Insertion

- File is unsorted (= **heap file**)
 - add it wherever there is space (easy 😊)
- File is sorted
 - Is there space on the right page ?
 - Yes: we are lucky, store it there
 - Is there space in a neighboring page ?
 - Look 1-2 pages to the left/right, shift records
 - If anything else fails, create **overflow page**

Overflow Pages



- After a while the file starts being dominated by overflow pages: time to reorganize

Modifications: Deletions

- Free space in page, shift records
 - Be careful with slots
 - RIDs for remaining tuples must NOT change
- May be able to eliminate an overflow page

Modifications: Updates

- If new record is shorter than previous, easy 😊
- If it is longer, need to shift records
 - May have to create overflow pages

Searching in a Heap File

File is **not sorted** on any attribute

`Student(sid: int, age: int, ...)`

30	18 ...
70	21

— 1 record

20	20
40	19

} 1 page

80	19
60	18

10	21
50	22

Heap File Search Example

- 10,000 students
- 10 student records per page
- Total number of pages: 1,000 pages
- Find student whose sid is 80
 - Must read on average 500 pages
- Find all students older than 20
 - Must read all 1,000 pages
- Can we do better?

Sequential File

File **sorted on an attribute**, usually on primary key

`Student(sid: int, age: int, ...)`

10	21 ...
20	20

30	18
40	19

50	22
60	18

70	21
80	19

Sequential File Example

- Total number of pages: 1,000 pages
- Find student whose sid is 80
 - Could do binary search, read $\log_2(1,000) \approx 10$ pages
- Find all students older than 20
 - Must still read all 1,000 pages
- Can we do even better?

Outline

- **Data storage**
 - Disk and files
 - Operations on files
- **Indexes**
 - Index structures
 - Hash-based indexes
 - B+ trees

Indexes

- **Index**: data structure that organizes data records on disk to optimize selections on the **search key fields** for the index
- An index contains a collection of **data entries**, and supports efficient retrieval of all data entries with a given search key value **k**

Index Classification Overview

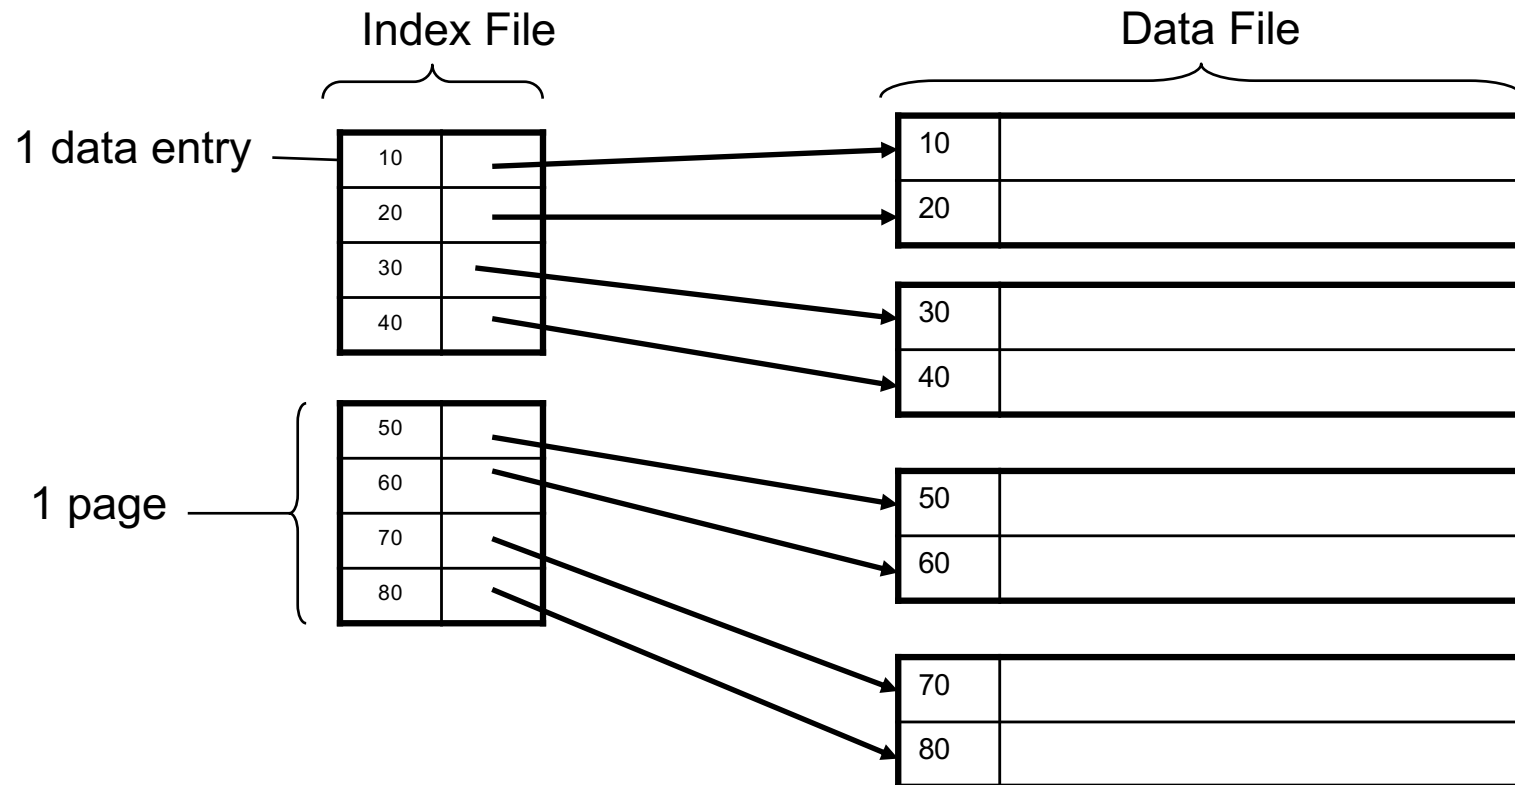
- Primary/secondary
 - Primary = determines the location of indexed records
 - Secondary = cannot reorder data, does not determine data location
- Dense/sparse
 - Dense = every key in the data appears in the index
 - Sparse = the index contains only some keys
- Clustered/unclustered
 - Clustered = records close in index are close in data
 - Unclustered = records close in index may be far in data
- B+ tree / Hash table / ...

Indexes

- **Search key** = can be any set of fields
 - not the same as the primary key, nor a key
- **Index** = collection of data entries
- **Data entry** for key k can be:
 - The actual record with key k
 - In this case, **the index is also a special file organization**
 - This type of index is also called the **primary index** of a file
 - (k, RID)
 - (k, list-of-RIDs)


Primary Index

- Index determines the location of indexed records
- Dense index: sequence of (key, pointer) pairs



Primary Index

- Sparse index




Primary Index Example

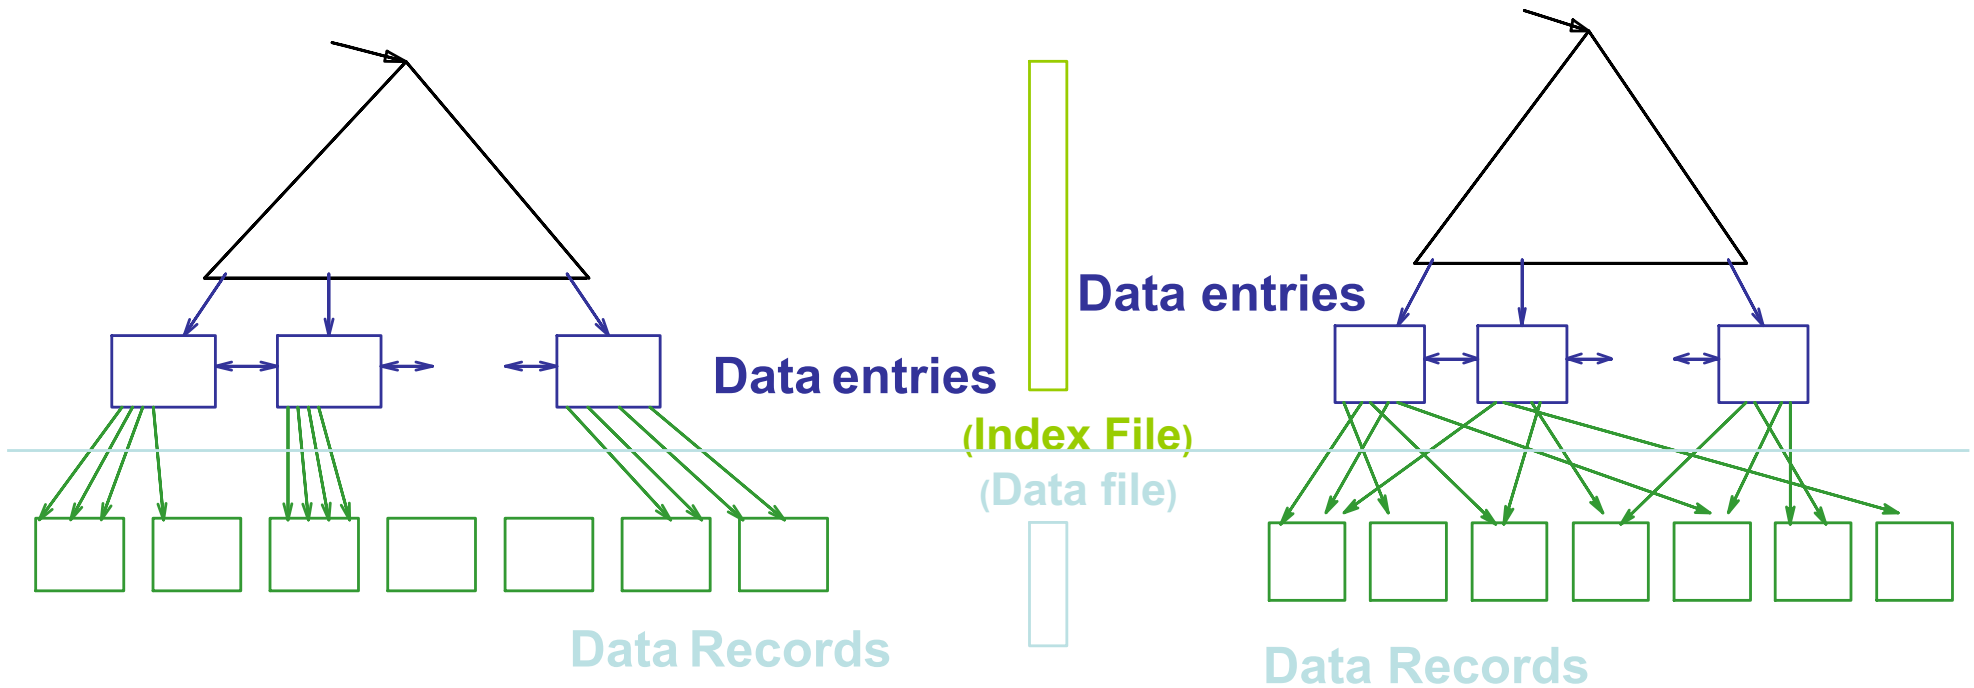
- Let's assume all pages of index fit in memory
- Find student whose sid is 80
 - Index (dense or sparse) points directly to the page
 - Only need to read 1 page from disk.
- Find all students older than 20
 - Must still read all 1,000 pages.
- How can we make *both* queries fast?

Secondary Indexes

- To index **other attributes than primary key**
- Always dense (why ?)



Clustered vs. Unclustered Index



CLUSTERED

UNCLUSTERED

Clustered = records close in index are close in data

Clustered/Unclustered

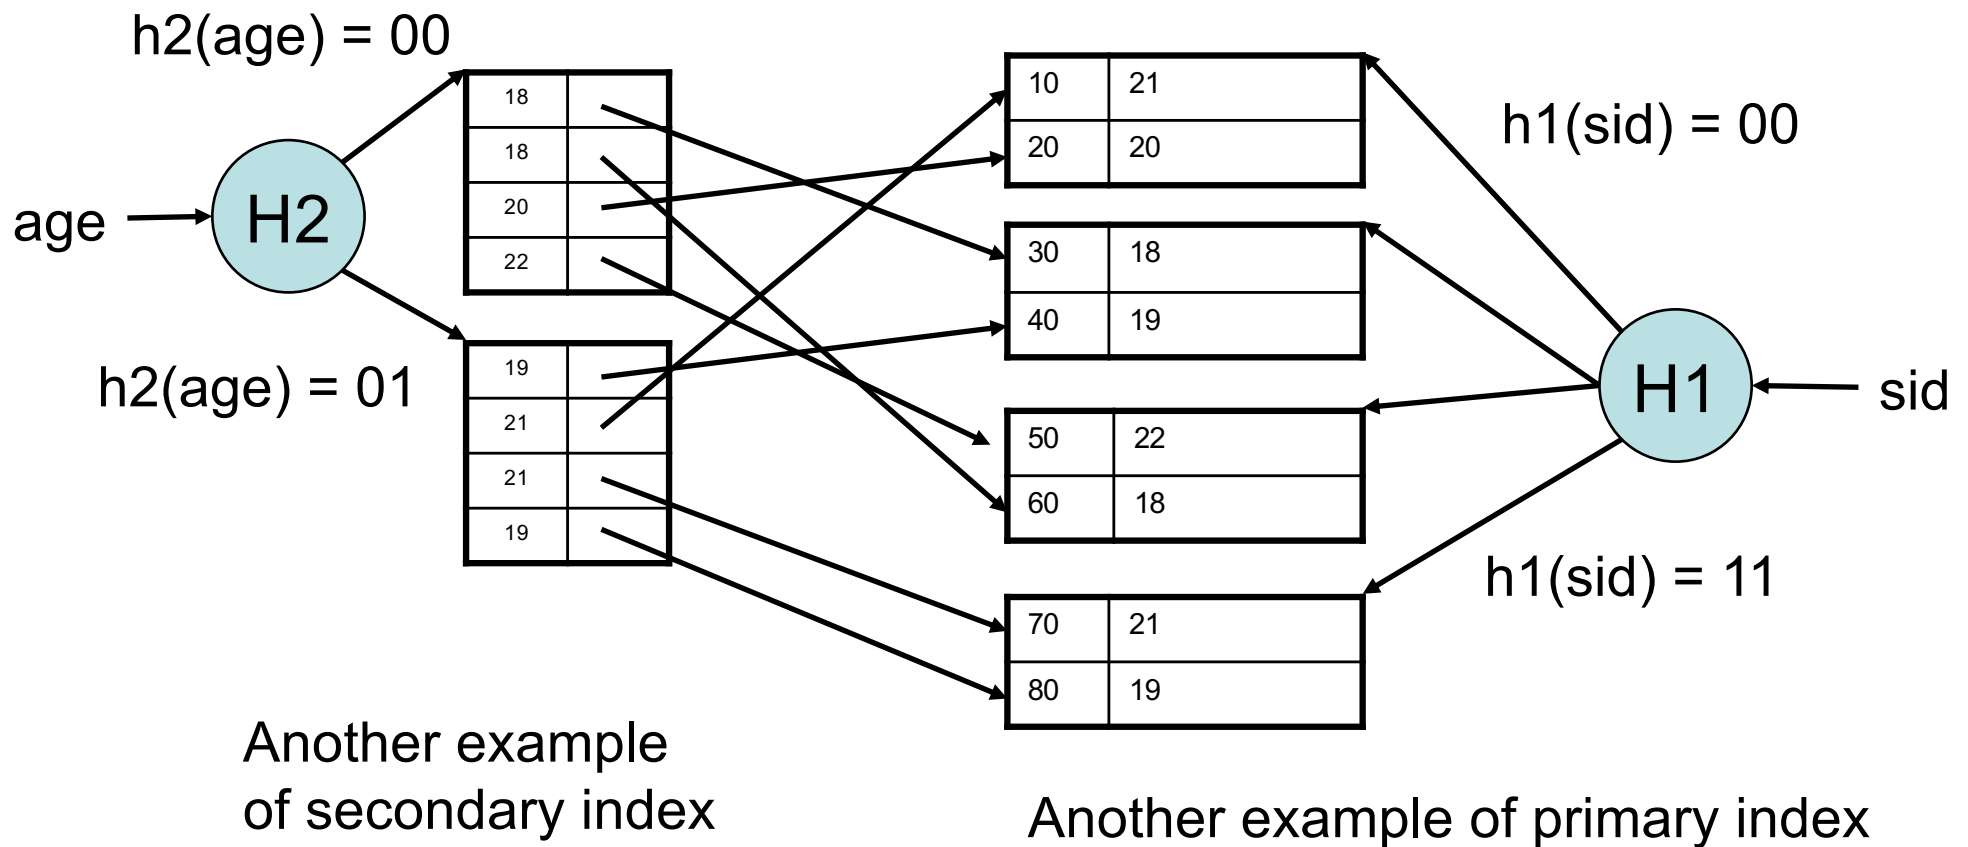
- Primary index = clustered by definition
- Secondary indexes = usually unclustered

Large Indexes

- What if index does not fit in memory?
- Would like to index the index itself
 - Hash-based index
 - Tree-based index

Hash-Based Index

Good for point queries but not range queries



Tree-Based Index


- How many index levels do we need?
- Can we create them automatically? **Yes!**
- **Can do something even more powerful!**

B+ Trees


- Search trees
- Idea in B Trees
 - Make 1 node = 1 page (= 1 block)
 - Keep tree balanced in height
- Idea in B+ Trees
 - Make leaves into a linked list : facilitates range queries

B+ Trees Basics

- Parameter **d** = the degree
- Each node has **$d \leq m \leq 2d$ keys** (except root)



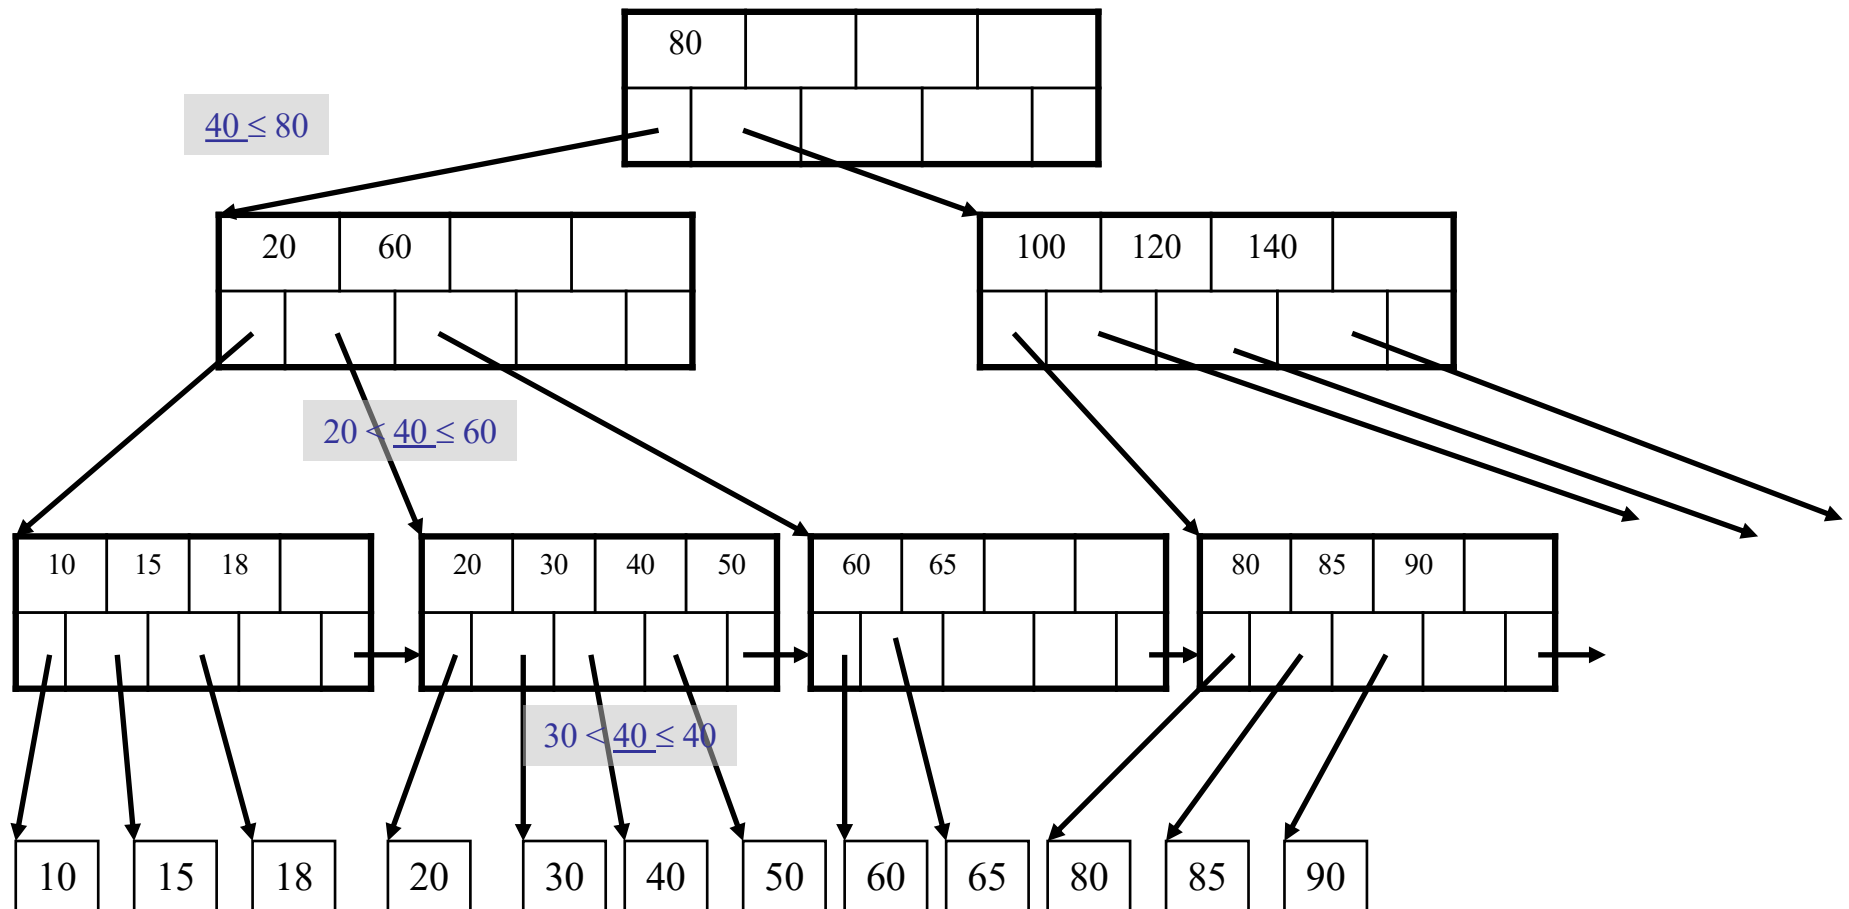
- Each leaf has **$d \leq m \leq 2d$ keys**:



B+ Tree Example

$d = 2$

Find the key 40



Searching a B+ Tree

- Exact key values:
 - Start at the root
 - Proceed down, to the leaf
- Range queries:
 - Find lowest bound as above
 - Then sequential traversal

```
Select name  
From Student  
Where age = 25
```

```
Select name  
From Student  
Where 20 <= age  
and age <= 30
```

B+ Tree Design

- How large d ?
- Example:
 - Key size = 4 bytes
 - Pointer size = 8 bytes
 - Block size = 4096 bytes
- $2d \times 4 + (2d+1) \times 8 \leq 4096$
- $d = 170$


B+ Trees in Practice

- Typical order: 100. Typical fill-factor: 67%.
 - average fanout = 133
- Typical capacities
 - Height 4: $133^4 = 312,900,700$ records
 - Height 3: $133^3 = 2,352,637$ records
- Can often hold top levels in buffer pool
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 Mbyte
 - Level 3 = 17,689 pages = 133 Mbytes

Insertion in a B+ Tree

Insert (K, P)

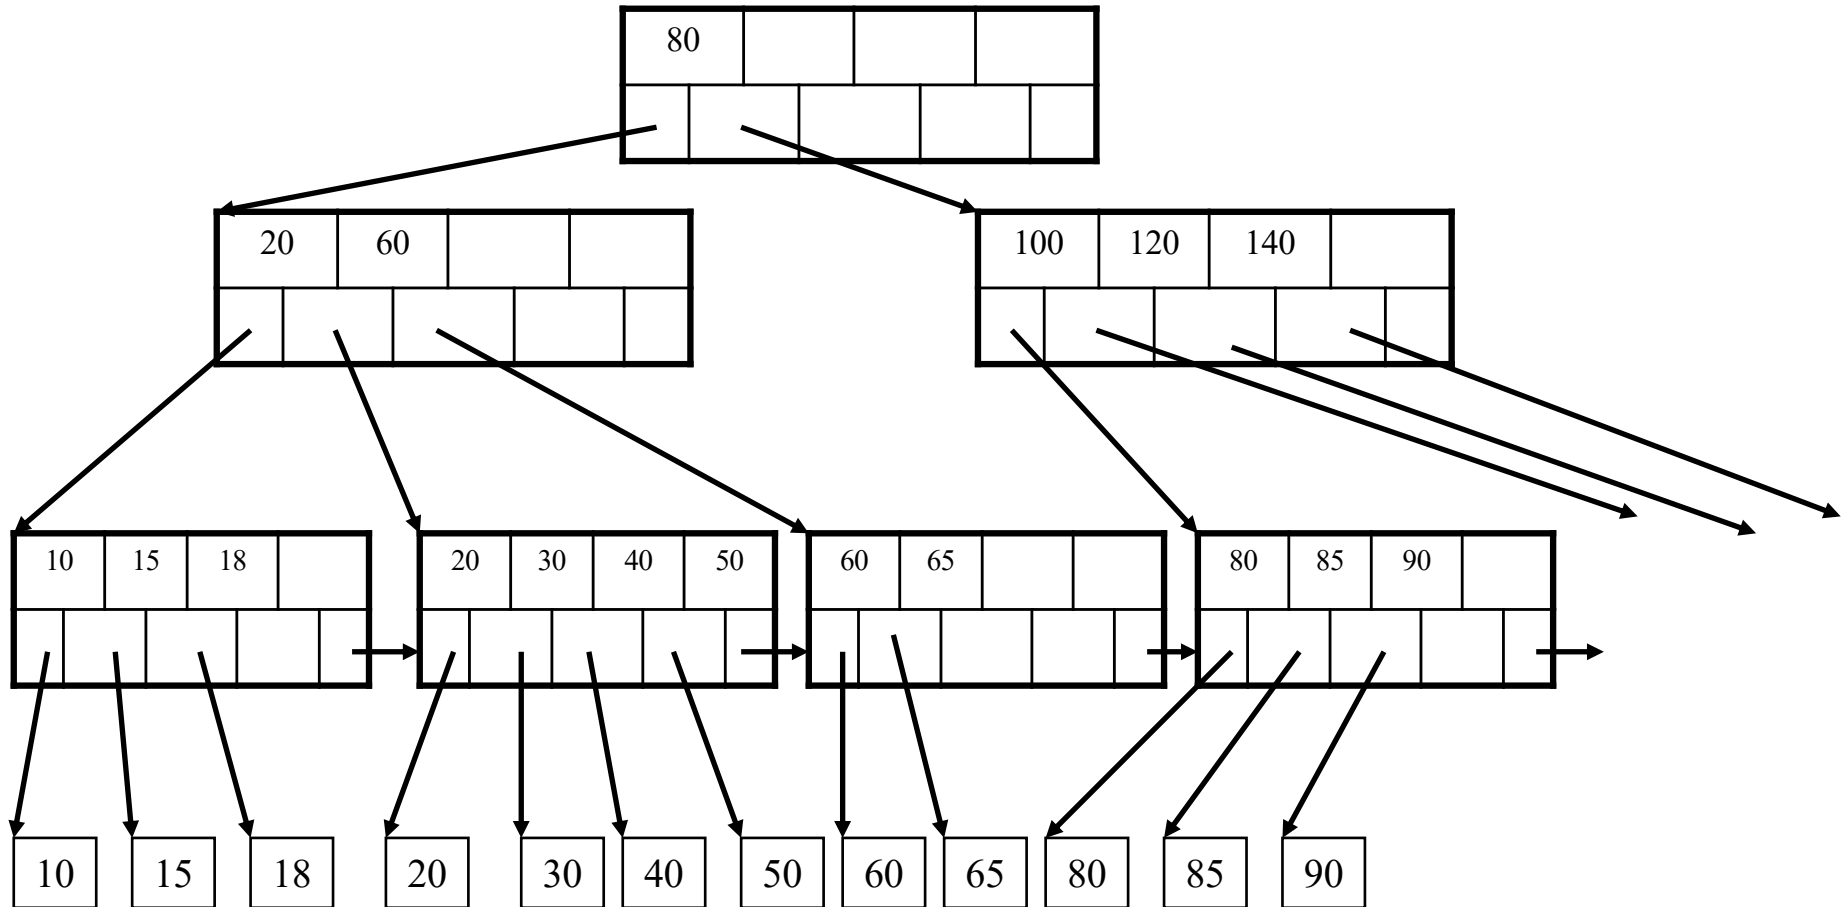
- Find leaf where K belongs, insert
- If no overflow ($2d$ keys or less), halt
- If overflow ($2d+1$ keys), split node, insert in parent:



- If leaf, also keep K_3 in right node
- When root splits, new root has 1 key only

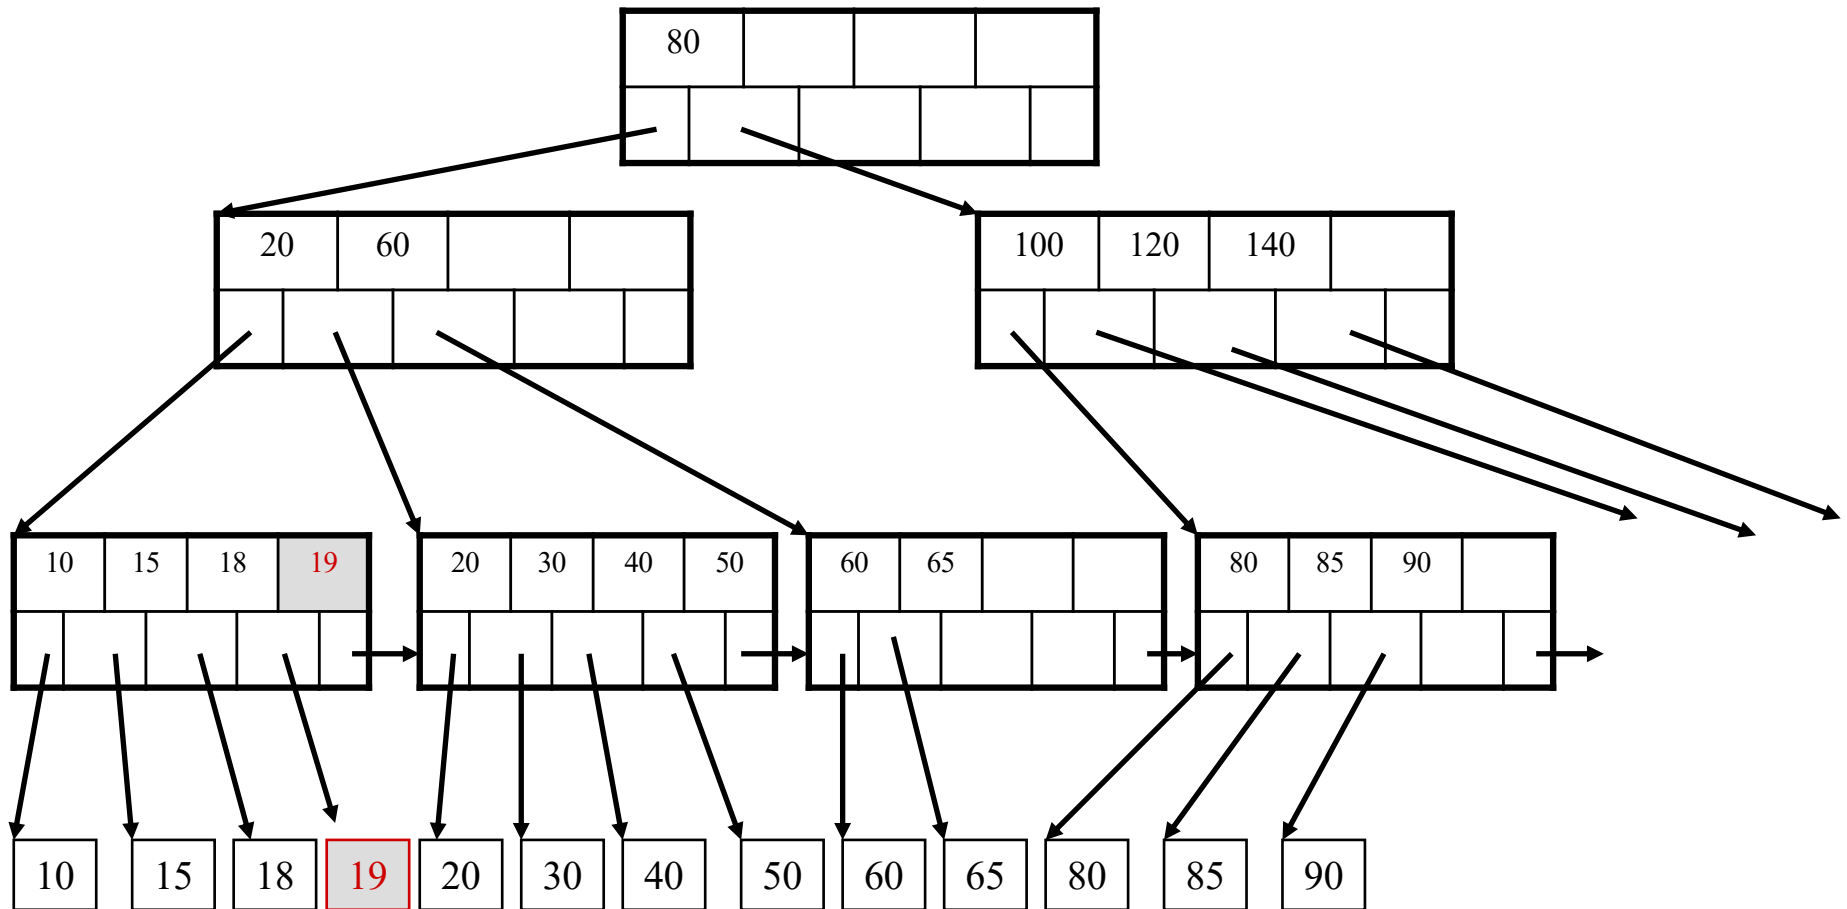
Insertion in a B+ Tree

Insert K=19



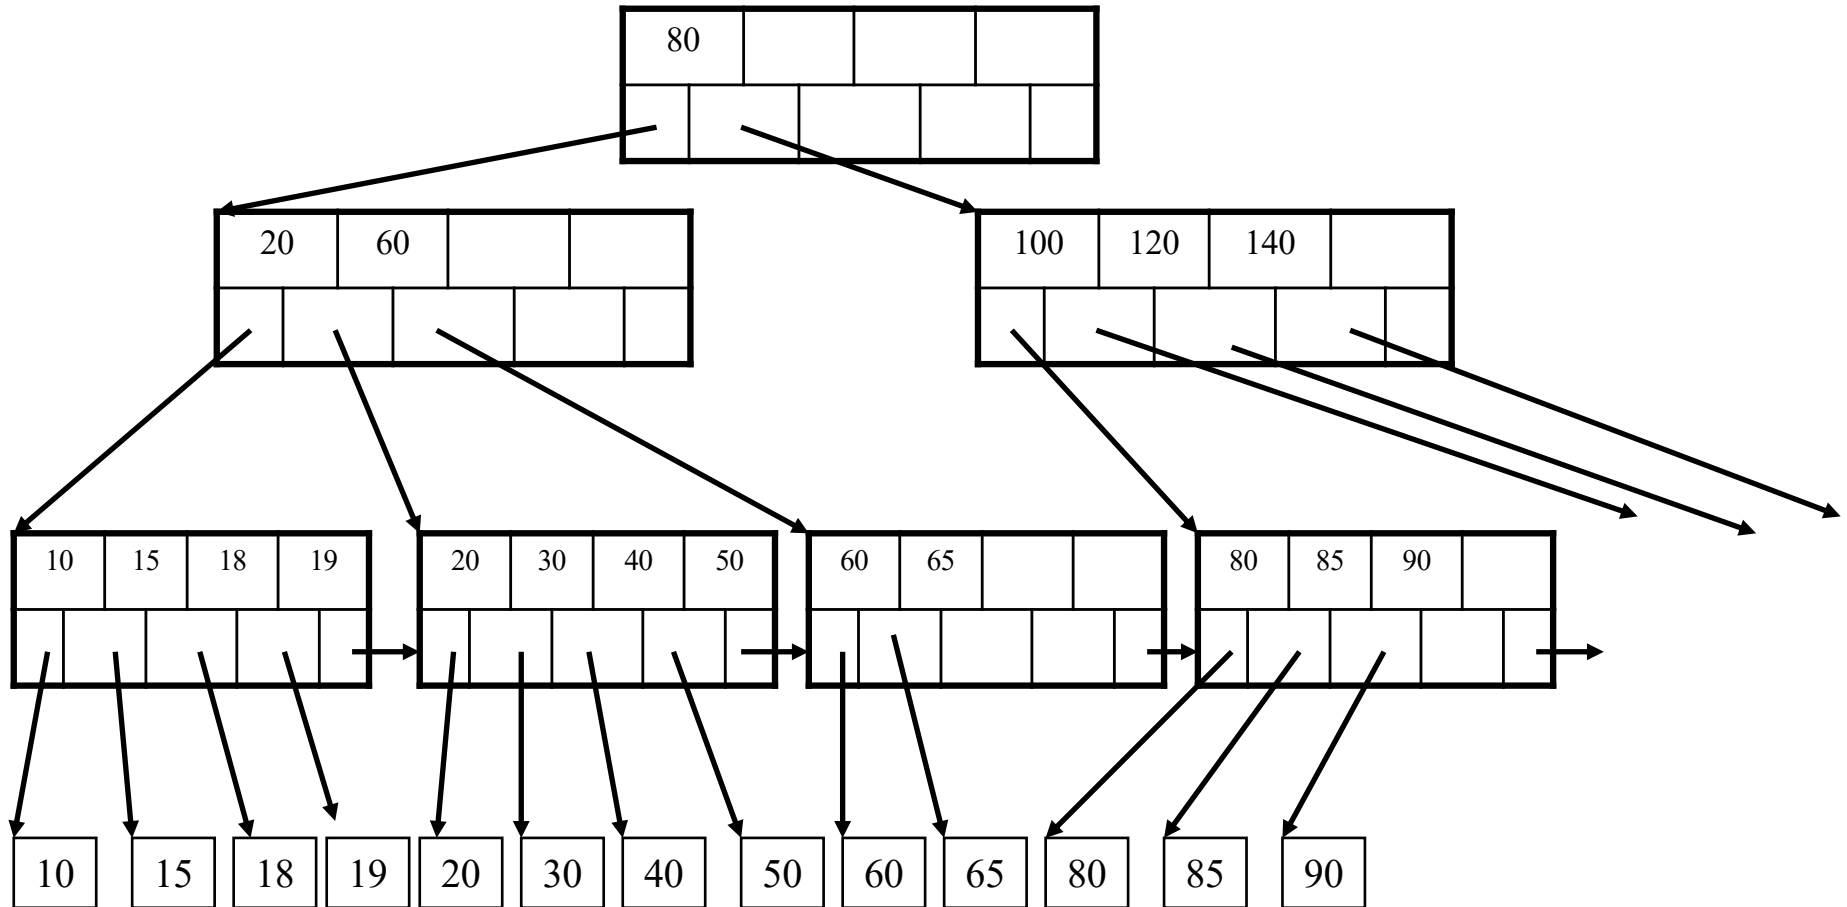
Insertion in a B+ Tree

After insertion



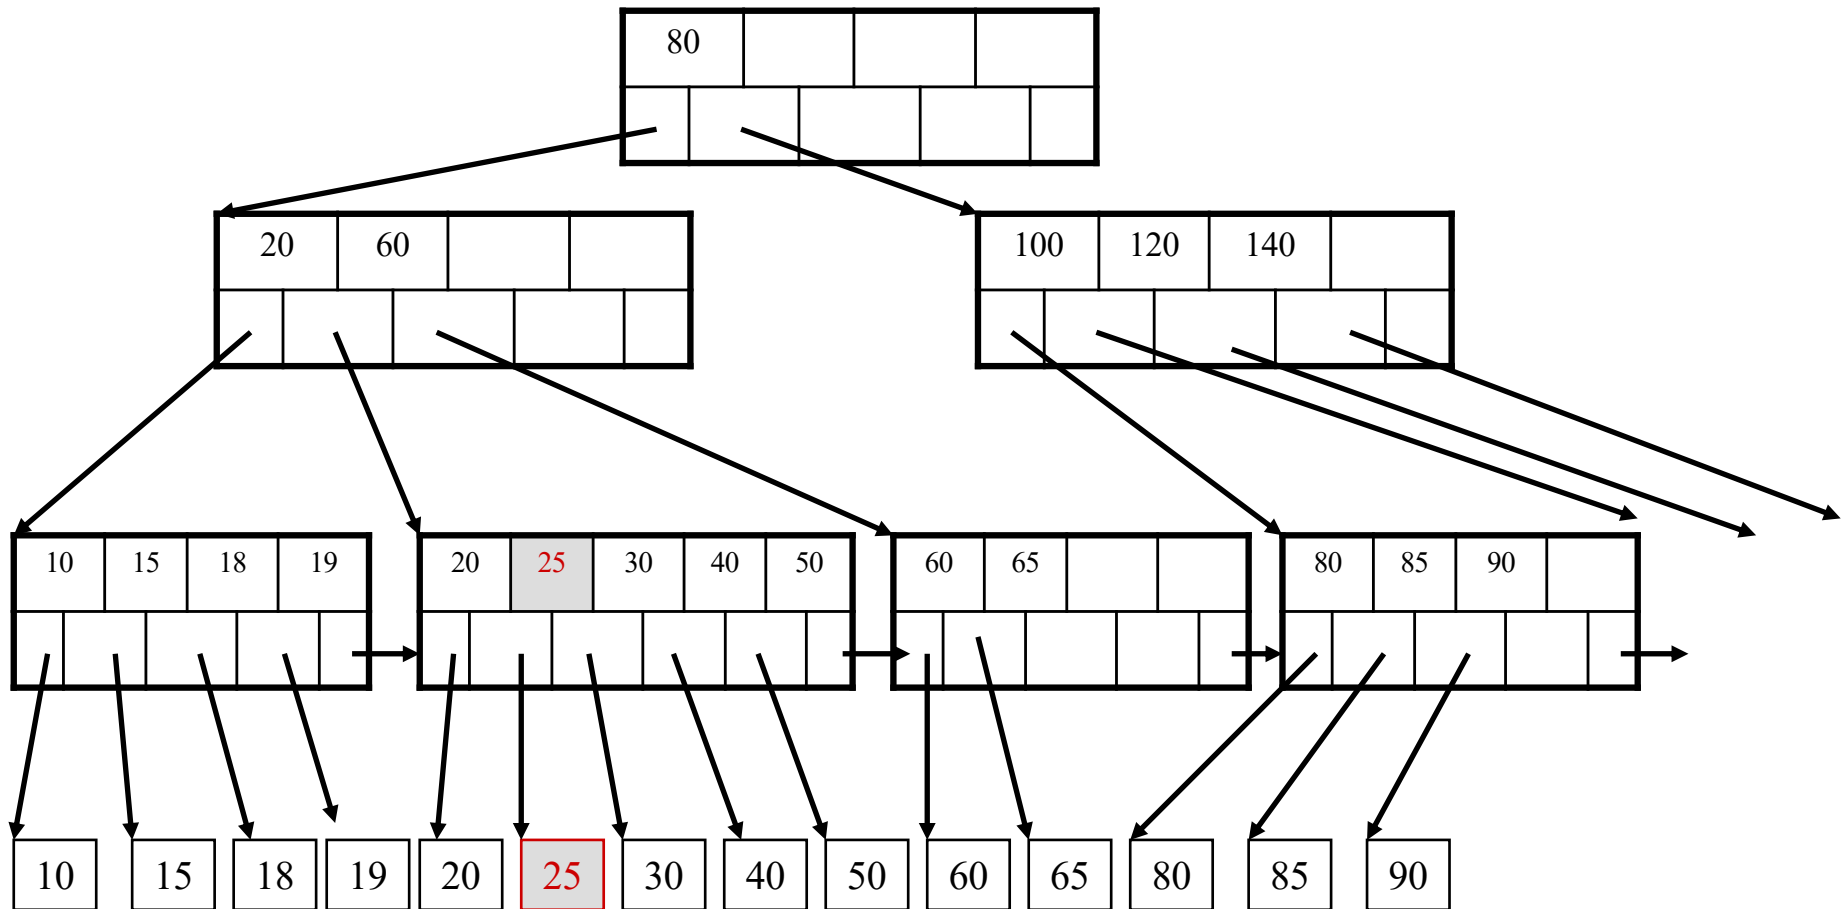
Insertion in a B+ Tree

Now insert 25



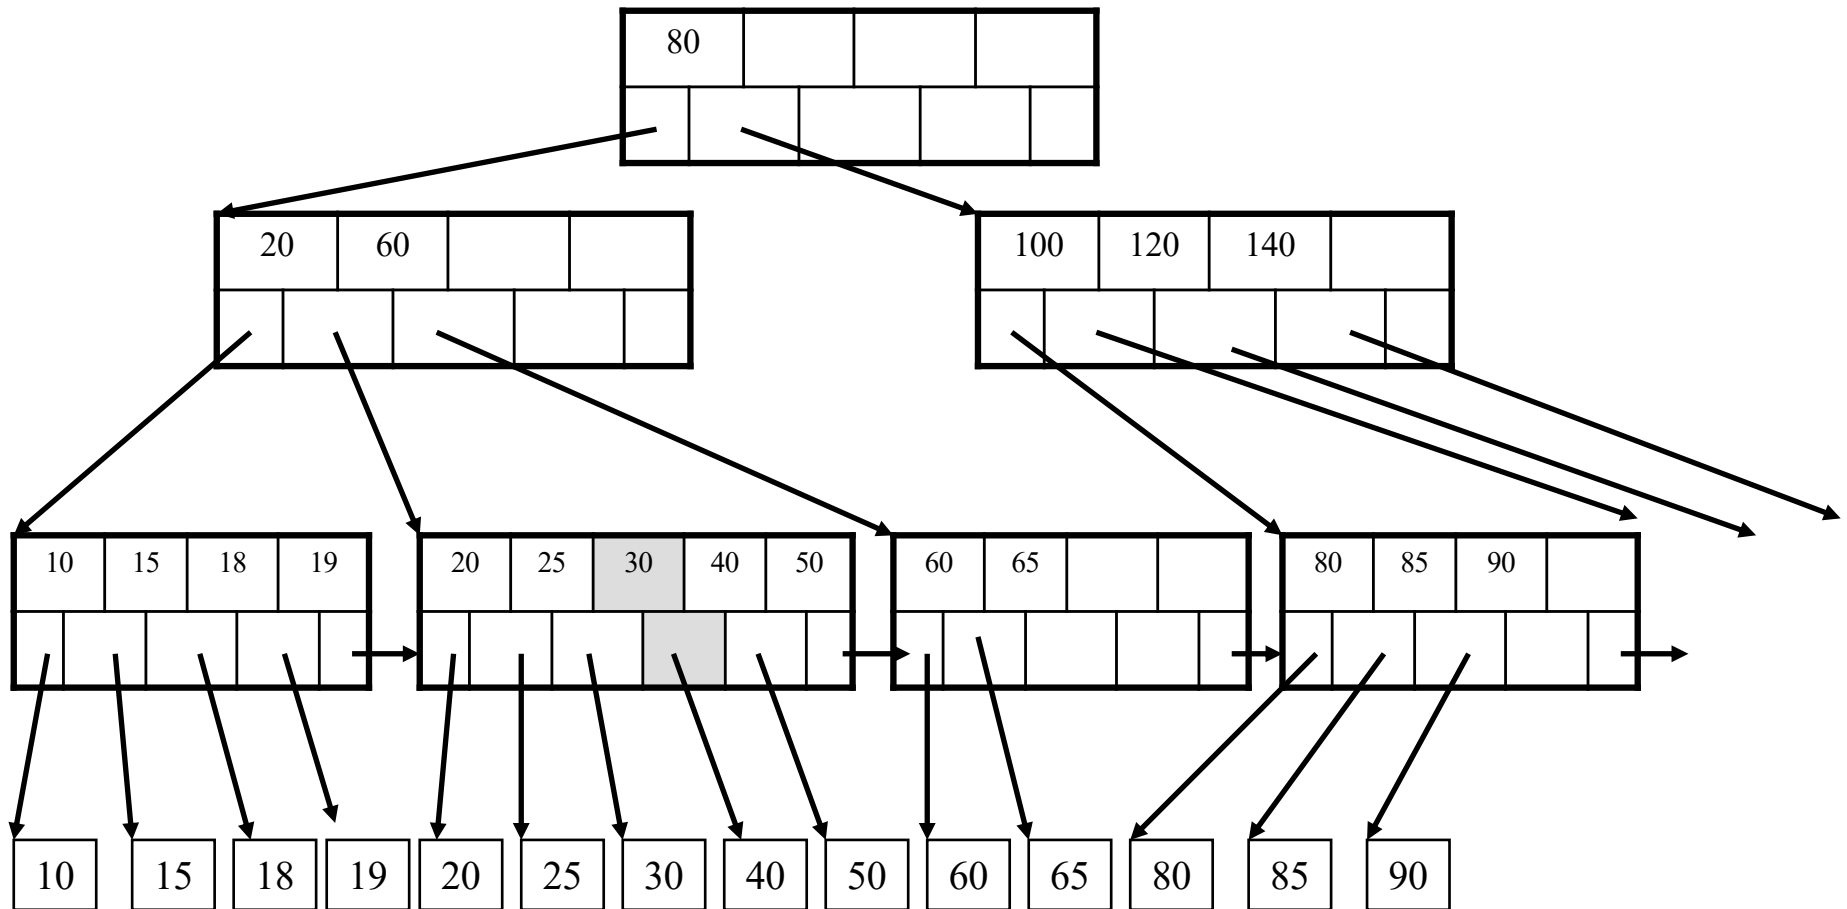
Insertion in a B+ Tree

After insertion




Insertion in a B+ Tree

But now have to split !




Insertion in a B+ Tree

After the split




Deletion from a B+ Tree

Delete 30



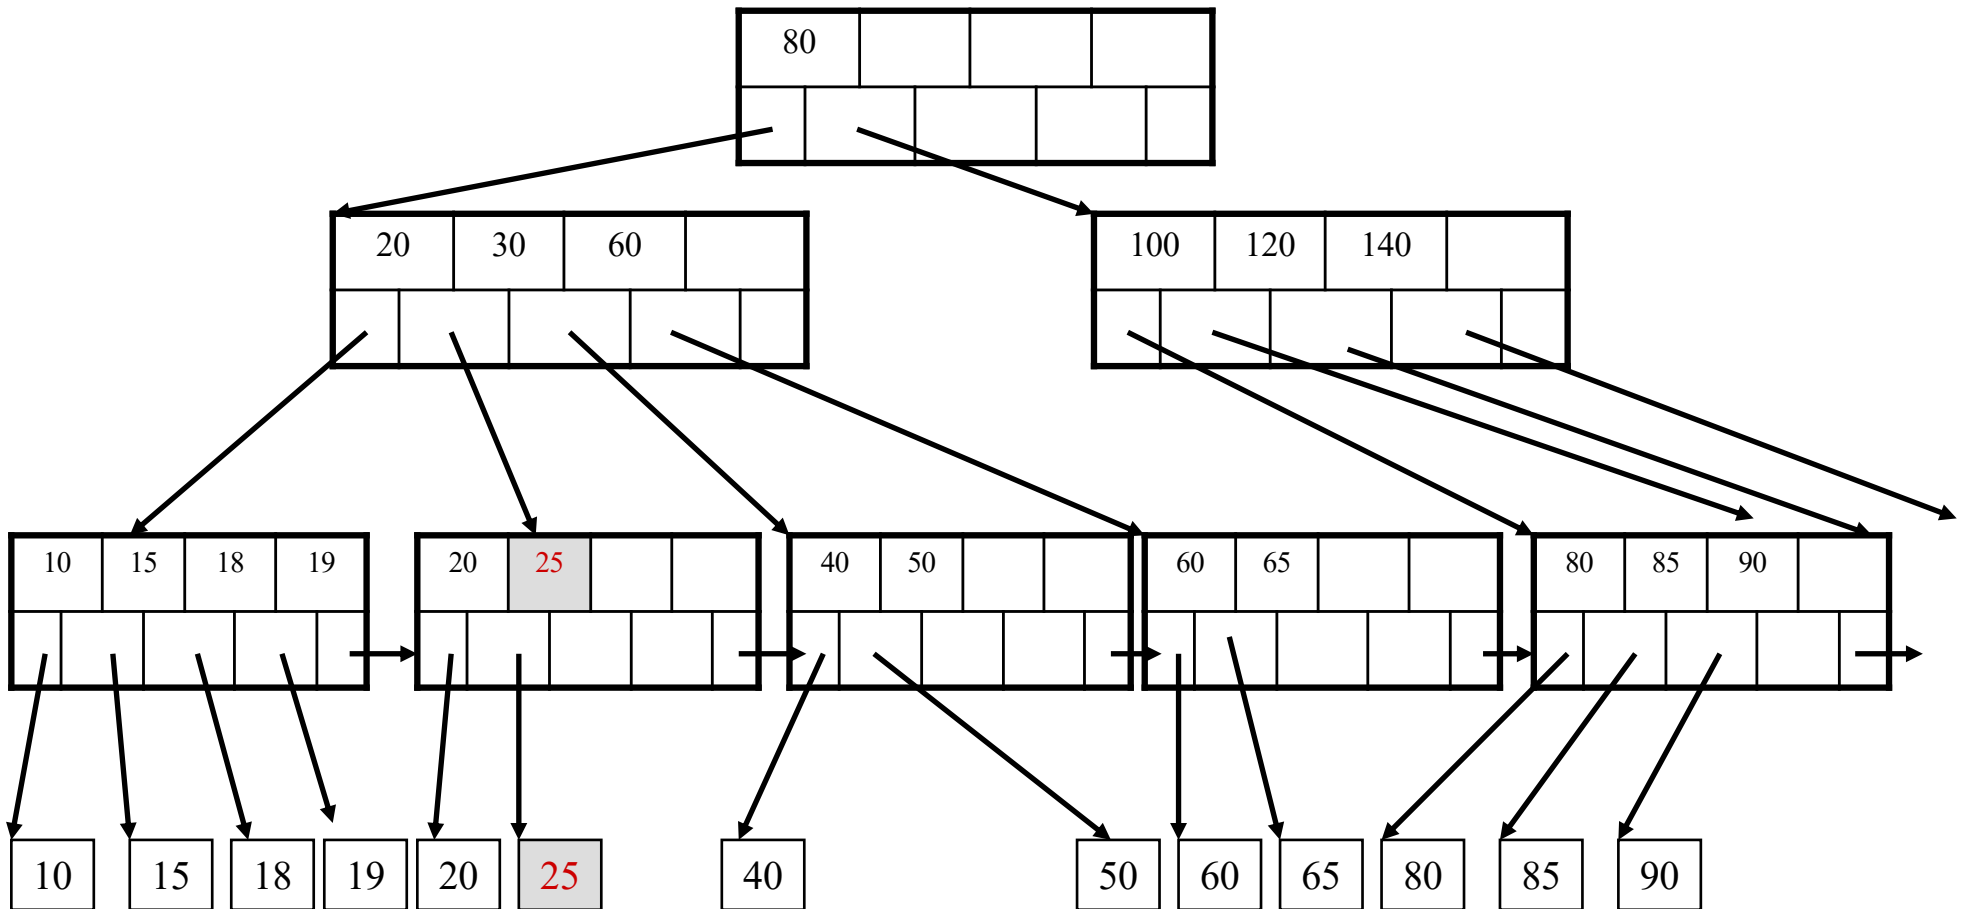
Deletion from a B+ Tree

After deleting 30



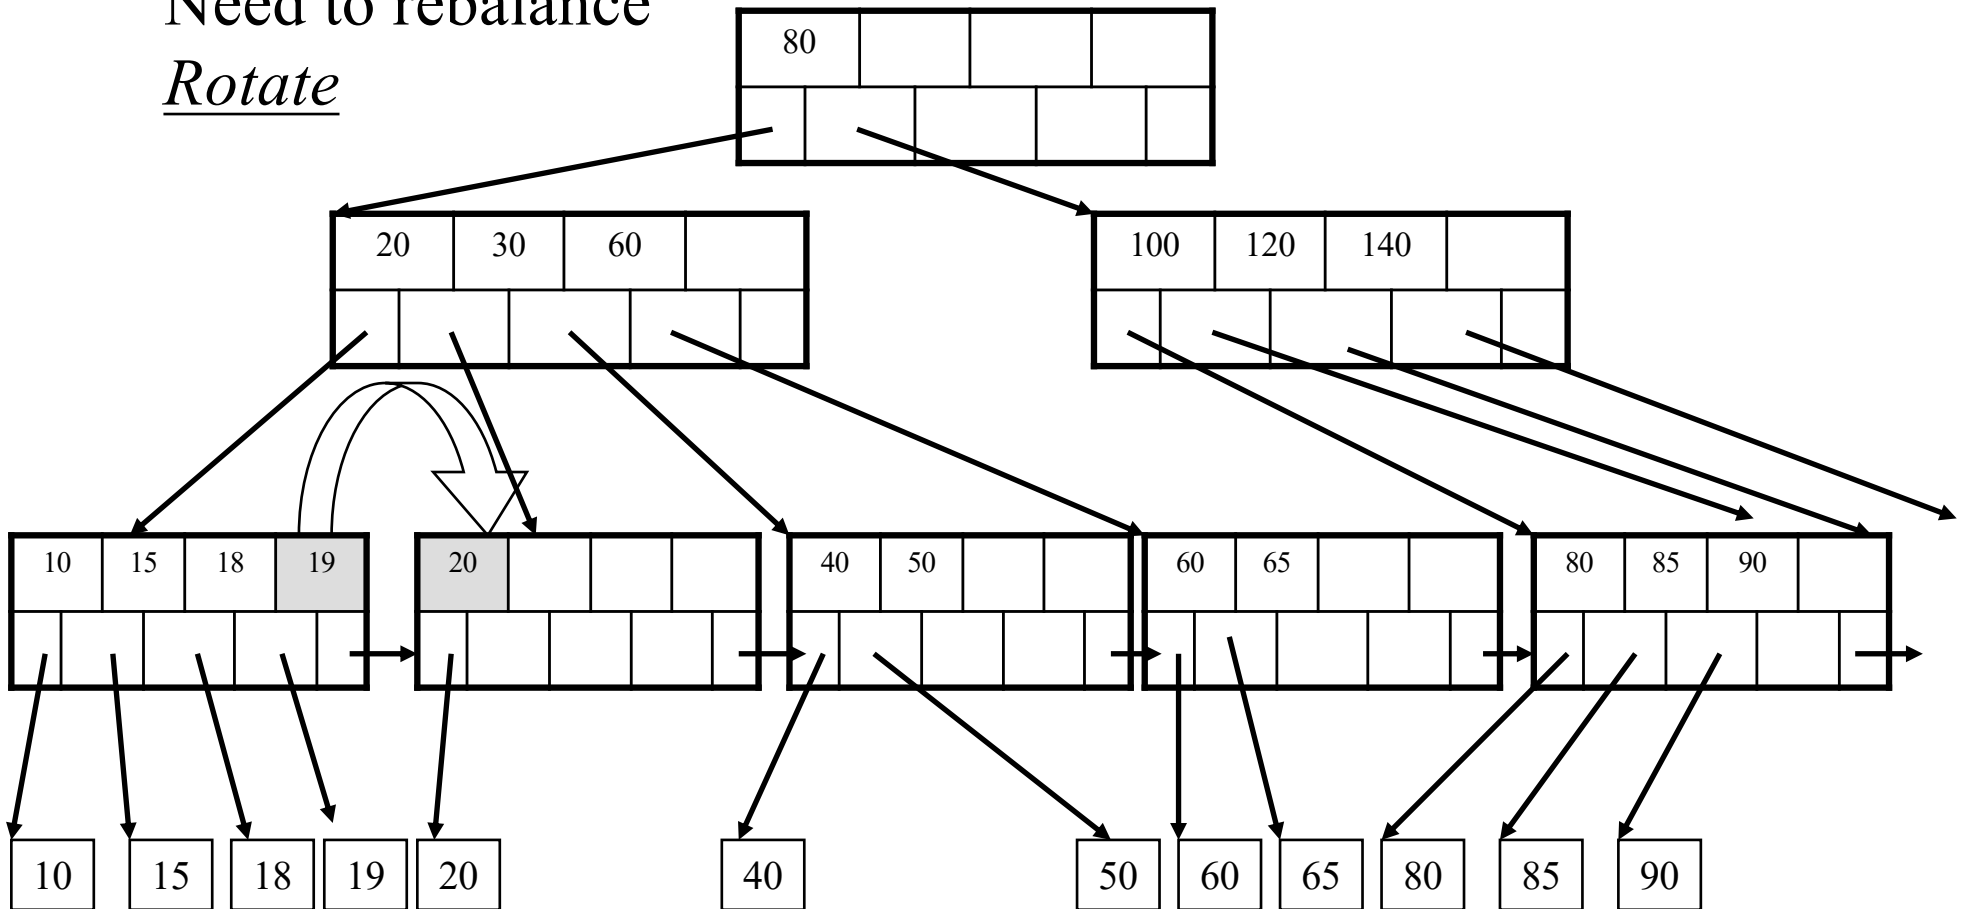
Deletion from a B+ Tree

Now delete 25



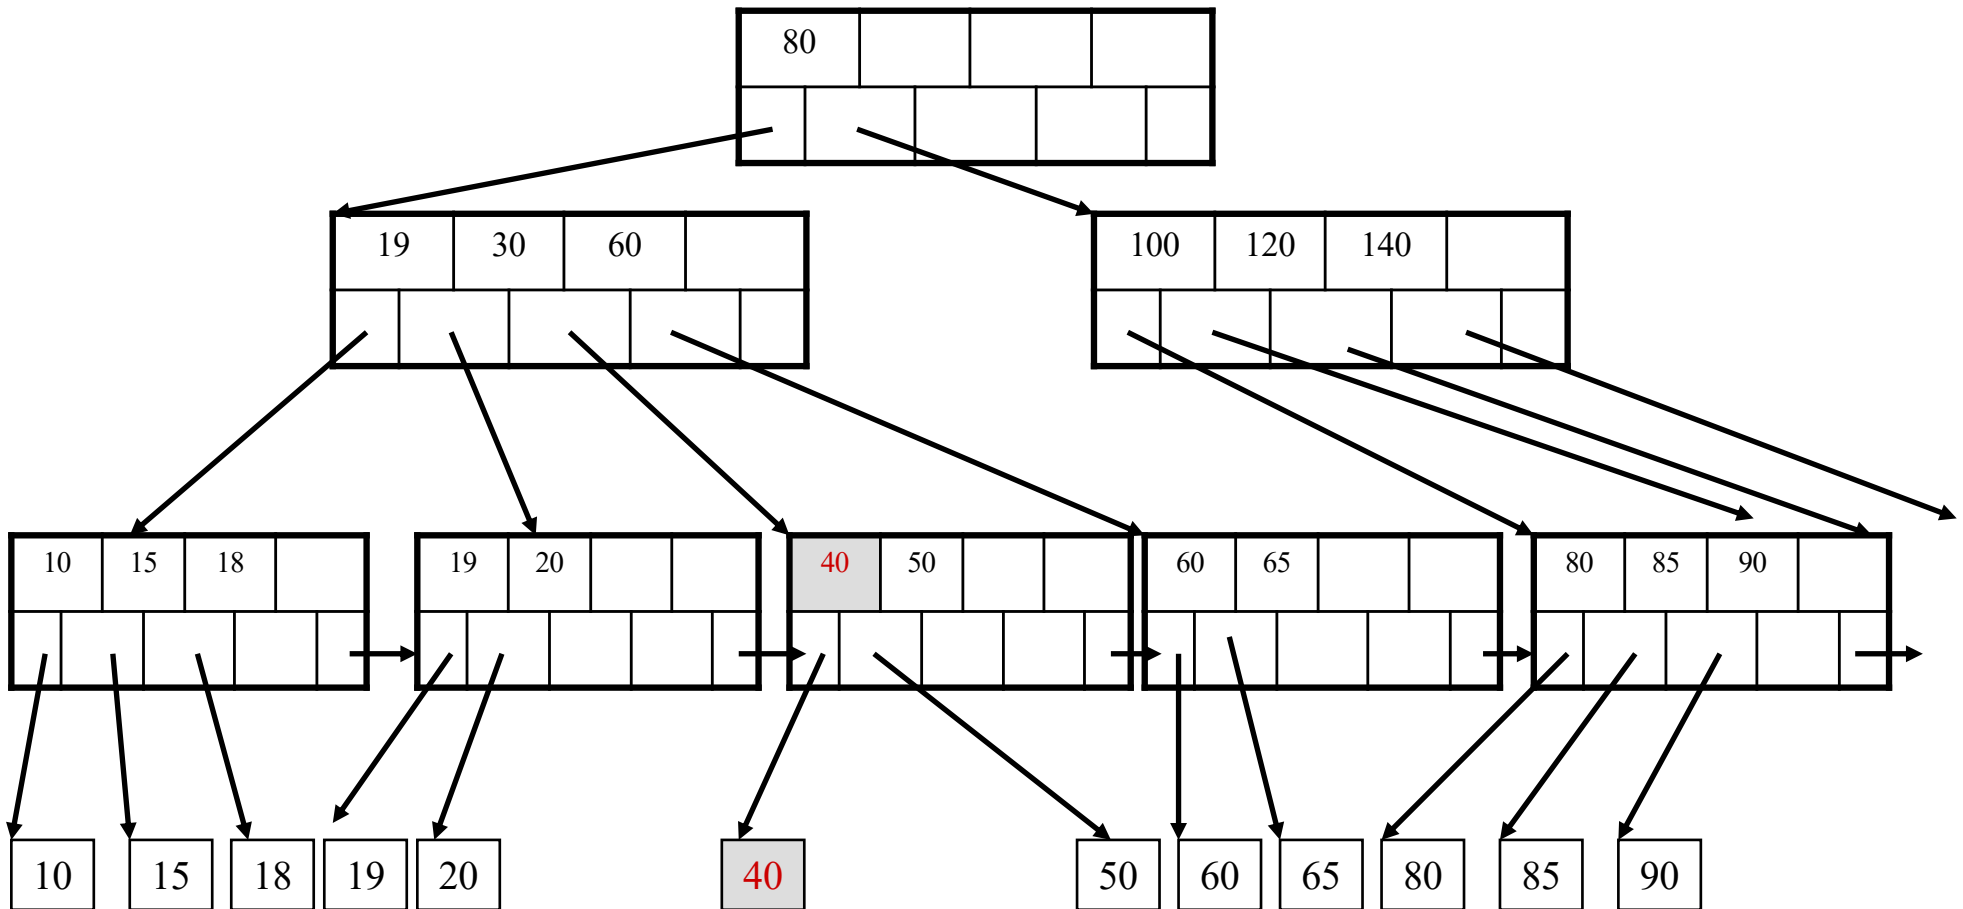
Deletion from a B+ Tree

After deleting 25
Need to rebalance
Rotate



Deletion from a B+ Tree

Now delete 40

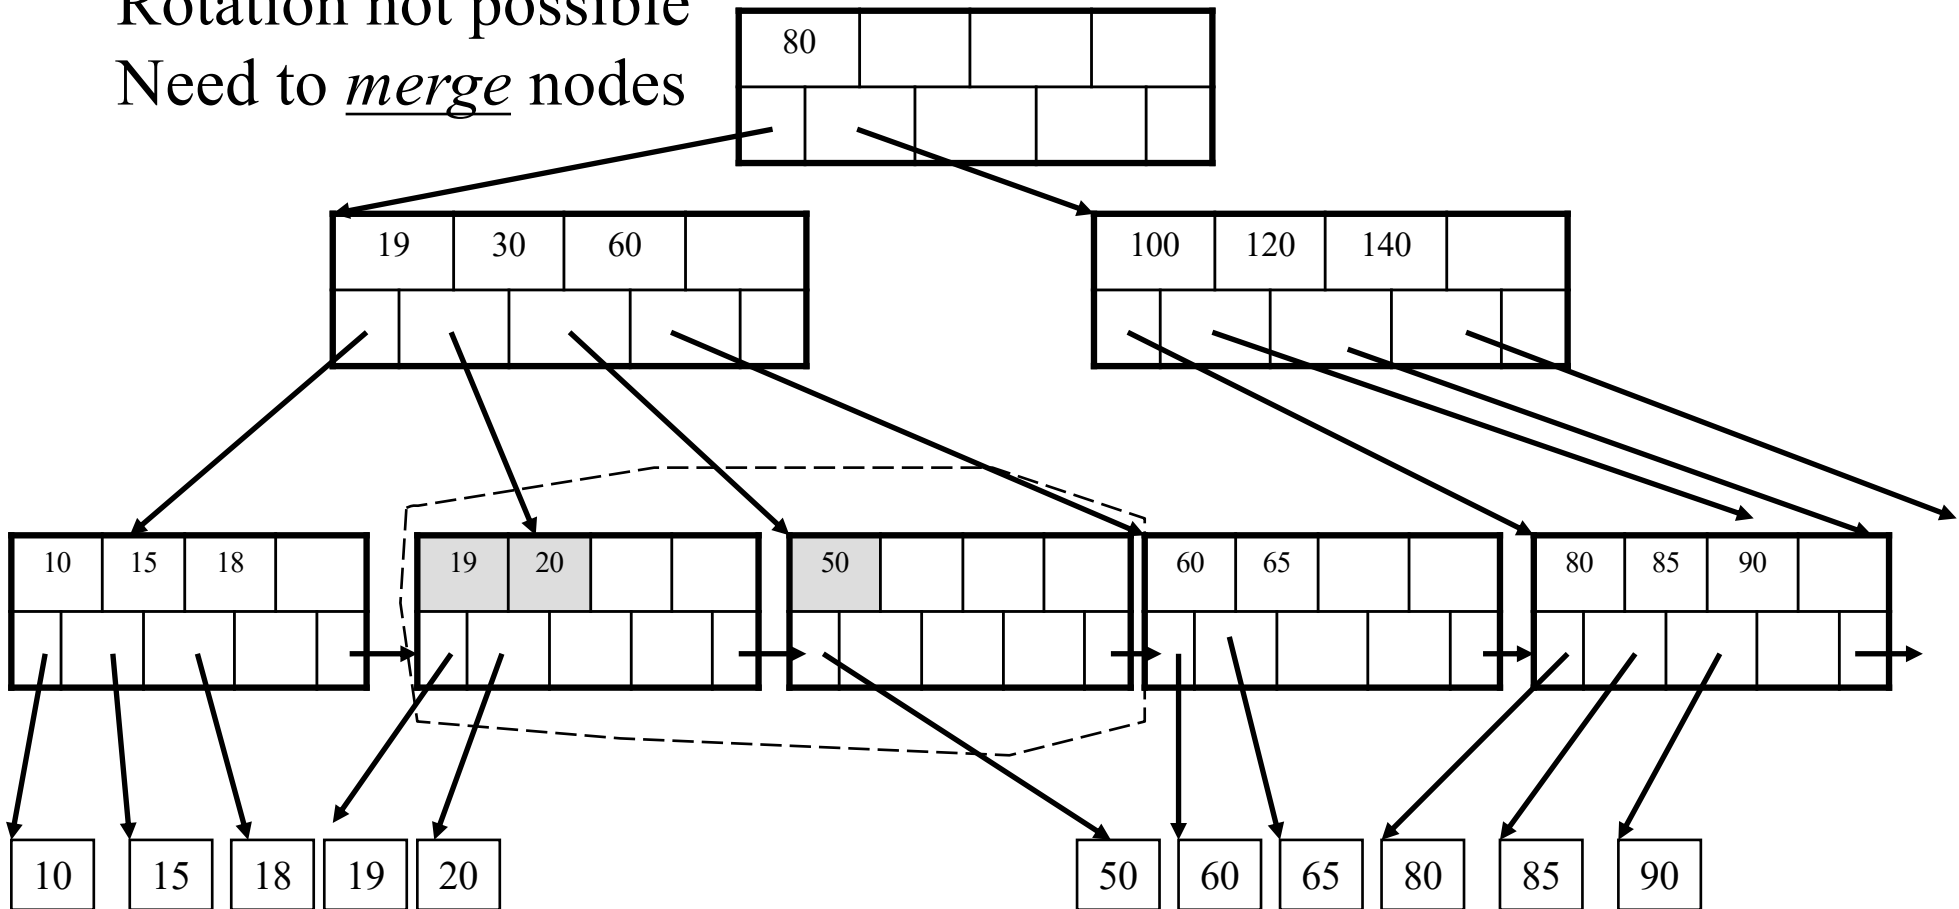


Deletion from a B+ Tree

After deleting 40

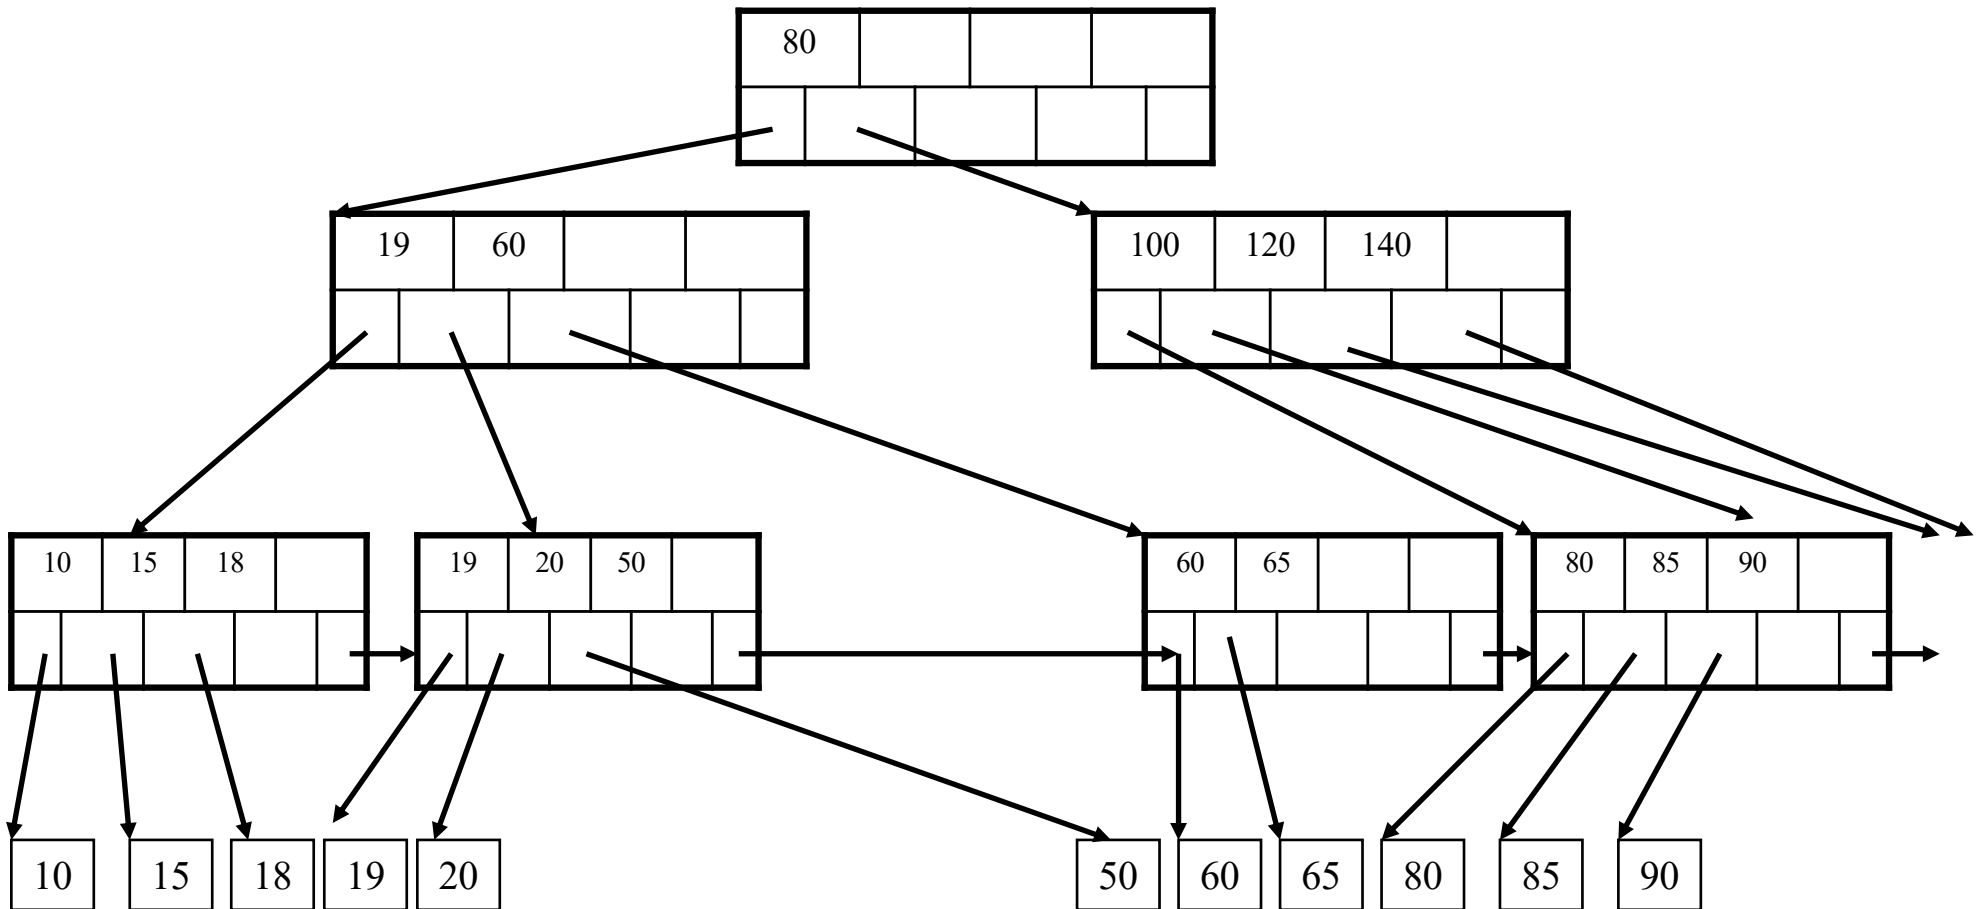
Rotation not possible

Need to merge nodes



Deletion from a B+ Tree

Final tree




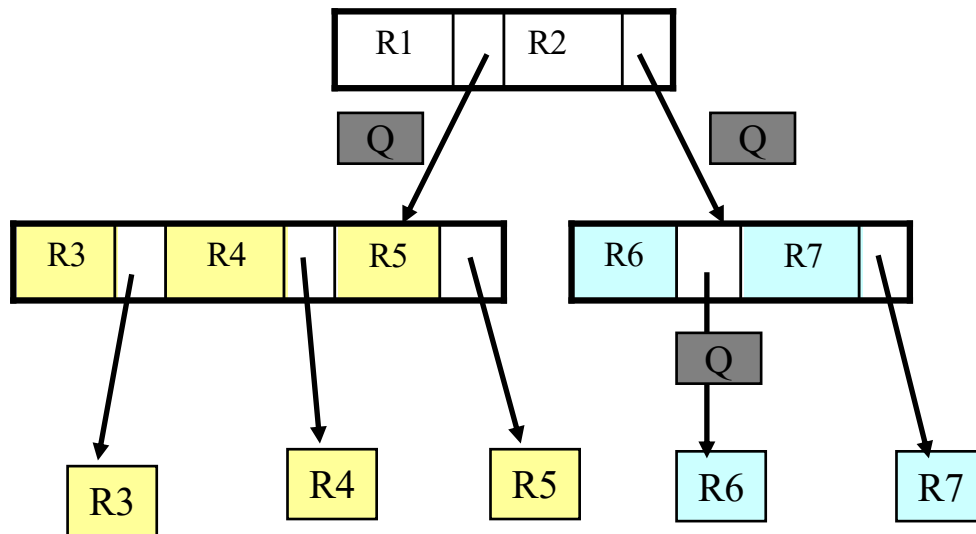
Summary on B+ Trees

- Default index structure on most DBMSs
- Very effective at answering 'point' queries:
 productName = 'gizmo'
- Effective for range queries:
 $50 < \text{price} \text{ AND } \text{price} < 100$
- Less effective for multirange:
 $50 < \text{price} < 100 \text{ AND } 2 < \text{quant} < 20$

R-Tree: a multidimensional B-Tree

Designed for spatial data

Search key values are bounding boxes



For insertion: at each level, choose child whose bounding box needs least enlargement (in terms of area)