

# Big Data Infrastructures


Philippe Cudré-Mauroux

Fall 2018

Lecture 3 – Query Execution & Optimization

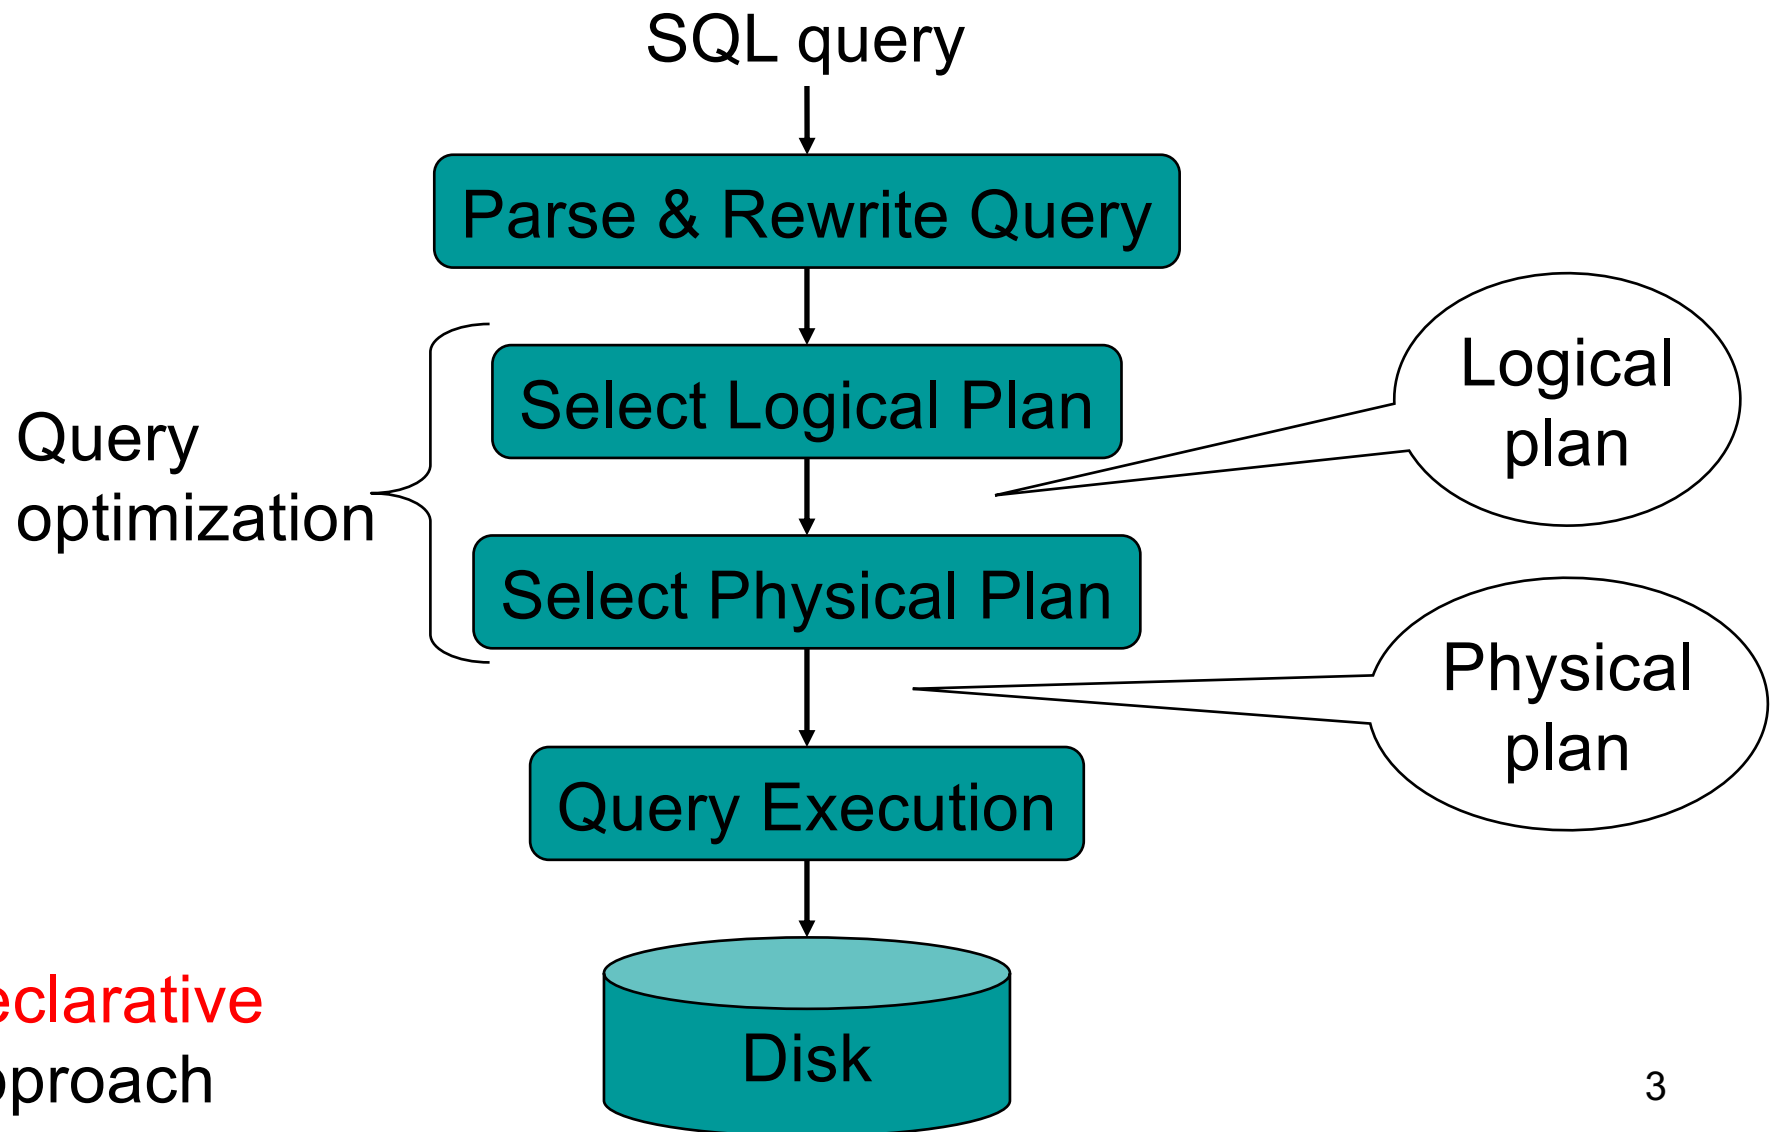
# Outline

---

- **Steps involved in processing a query**
  - Logical query plan
  - Physical query plan
  - Query execution overview
- **Operator implementations**
  - One pass algorithms
  - Two-pass algorithms 
  - Index-based algorithms
- **Query optimization 101**

# Query Evaluation Steps

---



# Example Database Schema

---

Supplier(sno, sname, scity, sstate)

Part(pno, pname, psize, pcolor)

Supply(sno, pno, price)

## View: Suppliers in Seattle

```
CREATE VIEW NearbySupp AS
```

```
SELECT sno, sname
```

```
FROM Supplier
```

```
WHERE scity='Seattle' AND sstate='WA'
```

# Example Query

---

- Find the names of all suppliers in Seattle who supply part number 2

```
SELECT sname FROM NearbySupp
WHERE sno IN ( SELECT sno
                FROM Supplies
                WHERE pno = 2 )
```

# Steps in Query Evaluation

---

- **Step 0: admission control**
  - User connects to the db with username, password
  - User sends query in text format
- **Step 1: Query parsing**
  - Parses query into an internal format
  - Performs various checks using catalog
    - Correctness, authorization, integrity constraints
- **Step 2: Query rewrite**
  - View rewriting, flattening, etc.

# Rewritten Version of Our Query

---

Original query:

```
SELECT sname
FROM NearbySupp
WHERE sno IN ( SELECT sno
                FROM Supplies
                WHERE pno = 2 )
```

Rewritten query:

```
SELECT S.sname
FROM Supplier S, Supplies U
WHERE S.scity='Seattle' AND S.sstate='WA'
AND S.sno = U.sno
AND U.pno = 2;
```

# Continue with Query Evaluation

---

- **Step 3: Query optimization**
  - Find an efficient query plan for executing the query
  - We will spend some time on this topic later
- **A query plan is**
  - **Logical query plan:** an extended relational algebra tree
  - **Physical query plan:** with additional annotations at each node
    - Access method to use for each relation
    - Implementation to use for each relational operator




# Extended Algebra Operators

---

- Union  $\cup$ , intersection  $\cap$ , difference  $-$
- Selection  $\sigma$
- Projection  $\pi$
- Join  $\bowtie$
- Duplicate elimination  $\delta$
- Grouping and aggregation  $\gamma$
- Sorting  $\tau$
- Rename  $\rho$

# Logical Query Plan

---



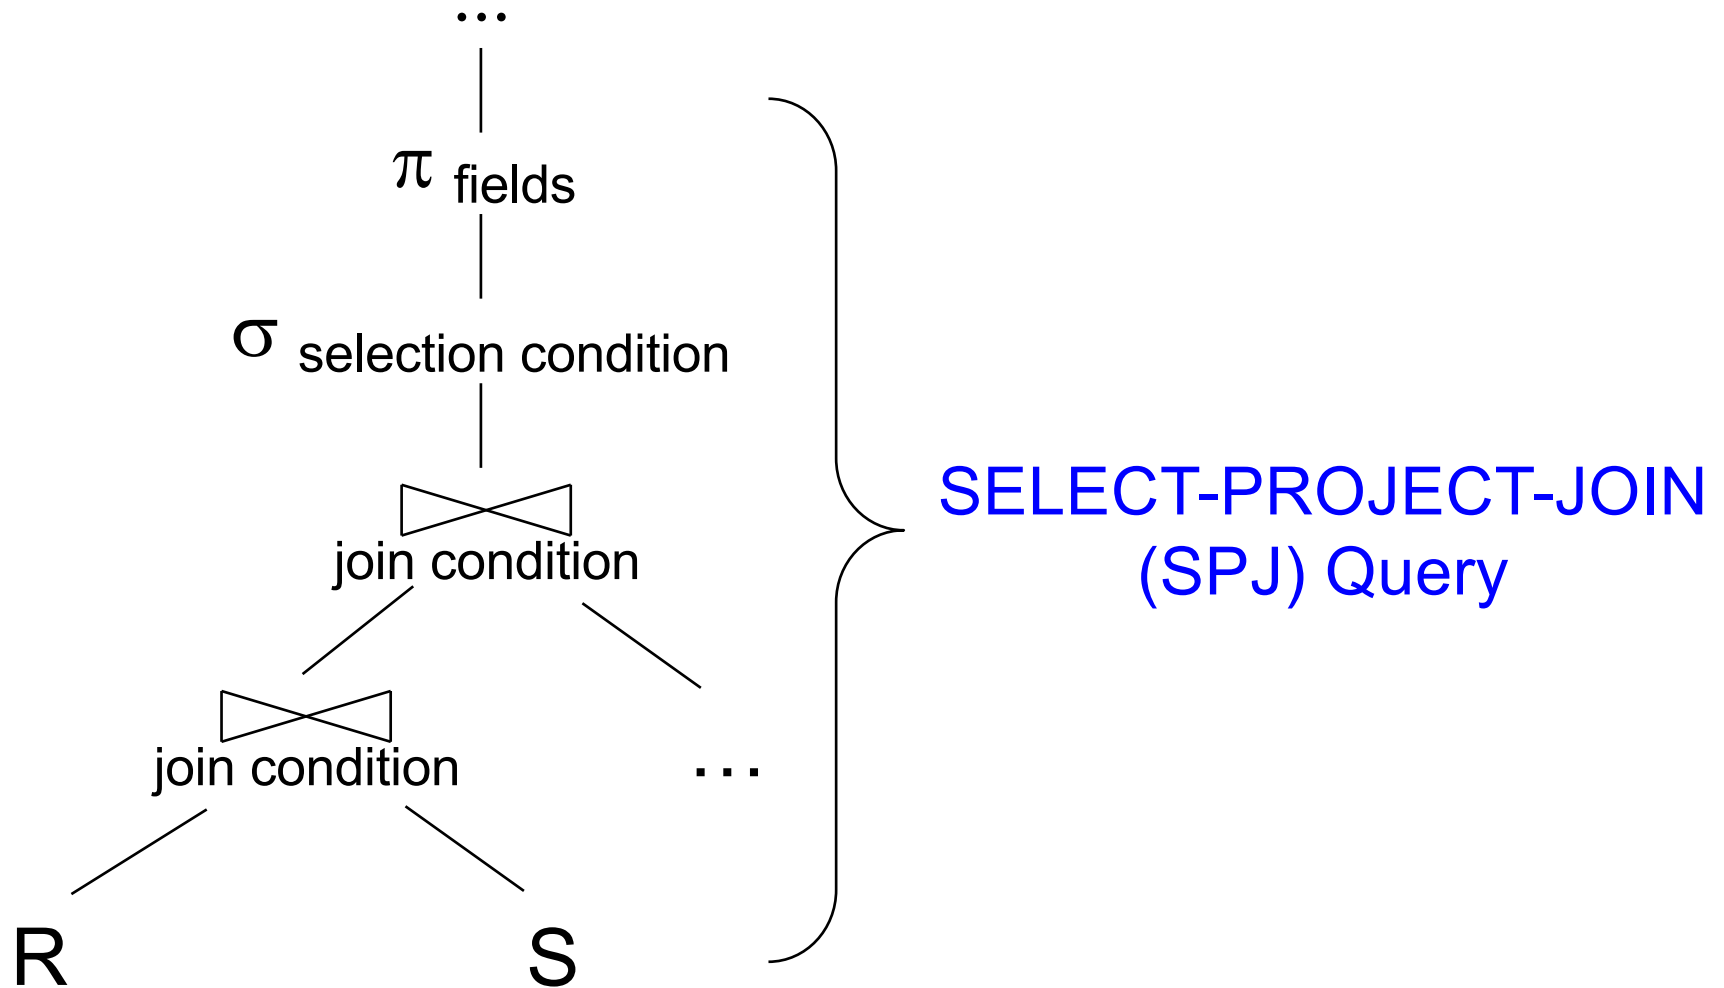
# Query Block

---

- Most optimizers operate on individual query blocks
- A query block is a SQL query with **no nesting**
  - **Exactly one**
    - SELECT clause
    - FROM clause
  - **At most one**
    - WHERE clause
    - GROUP BY clause
    - HAVING clause (behaves like a WHERE for aggregates)

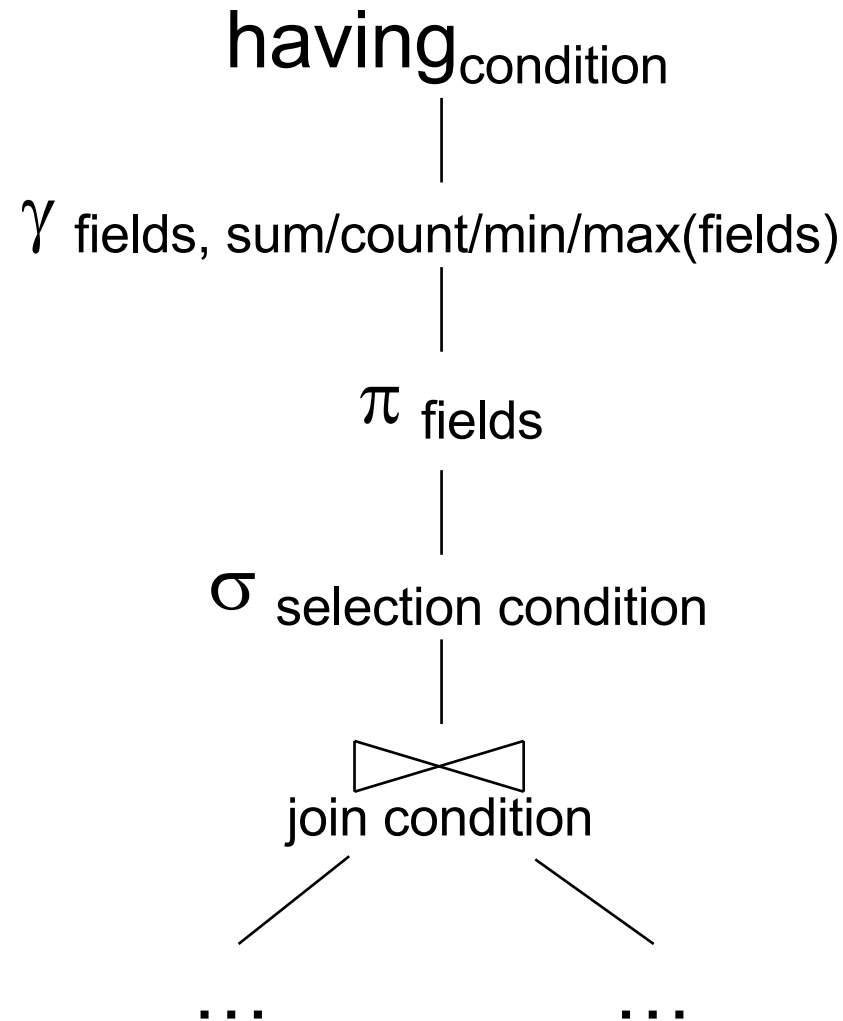
# Typical Plan for Block (1/2)

---



# Typical Plan For Block (2/2)

---



# Physical Query Plan

---

- Logical query plan with extra annotations
- **Access path selection** for each relation
  - Use a file scan or use an index
- **Implementation choice** for each operator
- **Scheduling decisions** for operators

# Physical Query Plan

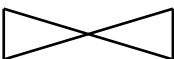
---

(On the fly)

$\pi_{\text{sname}}$

(On the fly)  $\sigma_{\text{sscity}='Seattle' \wedge \text{sstate}='WA' \wedge \text{pno}=2}$

(Nested loop)

  
 $\text{sno} = \text{sno}$

Suppliers  
(File scan)

Supplies  
(File scan)

# Final Step in Query Processing

---

- **Step 4: Query execution**
  - How to **synchronize operators?**
  - How to **pass data between operators?**
- Standard approaches:
  - **Iterator interface**
  - **Pipelined execution**
  - **Intermediate result materialization**



# Iterator Interface

---

- Each **operator implements this interface**
- Interface has only three methods
- **open()**
  - Initializes operator state
  - Sets parameters such as selection condition
- **get\_next()**
  - Operator invokes get\_next() iteratively on its inputs
  - Performs processing and produces an output tuple
- **close()**: clean-up state

# Pipelined Execution

---

- Applies parent operator to tuples directly as they are produced by child operators
- Benefits
  - No operator synchronization issues
  - Saves cost of writing intermediate data to disk
  - Saves cost of reading intermediate data from disk
  - Good resource utilizations on single processor
- This approach is used whenever possible

# Pipelined Execution


---

(On the fly)

$\pi$  sname

(On the fly)  $\sigma$  sscity='Seattle'  $\wedge$  sstate='WA'  $\wedge$  pno=2

(Nested loop)

  
sno = sno

Suppliers  
(File scan)

Supplies  
(File scan)

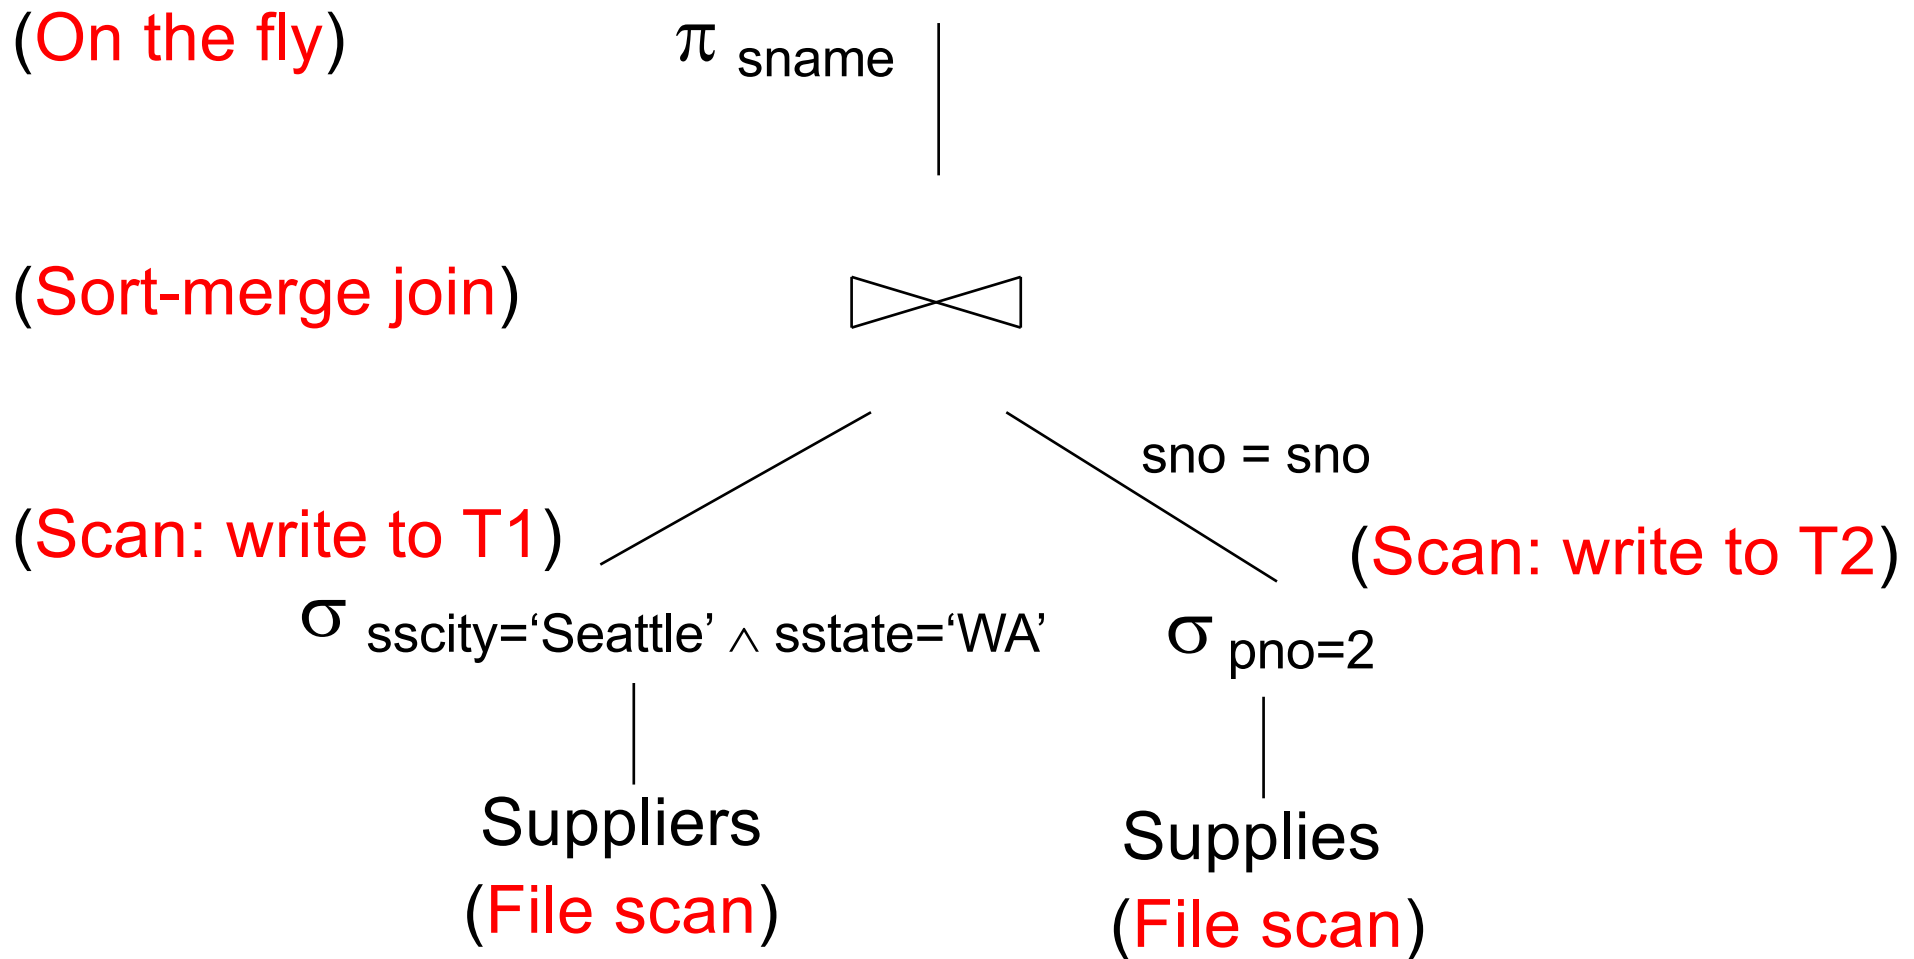
# Intermediate Tuple Materialization

---

- Writes the results of an operator to an intermediate table on disk
- No direct benefit but necessary for some operator implementations
  - E.g., When operator needs to examine the same tuples multiple times

# Intermediate Tuple Materialization

---



# Intermediate Results

---

- The change in available computing resources, e.g. amount of RAM, multiple cores per CPU, may yield to additional changes in result handling
- To achieve best locality of reference, non-pipelined plans can be favored
  - Intermediate results are written only to main memory and not to disk
- Applicability depends on the size of the estimated intermediate result and available buffer size

# Outline

---

- **Steps involved in processing a query**
  - Logical query plan
  - Physical query plan
  - Query execution overview
- **Operator implementations**
  - One pass algorithms
  - Two-pass algorithms
  - Index-based algorithms
- **Query optimization 101**

# Why Learn About Op Algos?

---

- Implemented in commercial DBMSs and Big Data frameworks
- Different DBMSs implement different subsets of these algorithms
- Good algorithms can greatly improve performance
- Need to know about physical operators to understand query optimization



# Cost Parameters

---

- In database systems the data is on disk
- **Cost = total number of I/Os**
  - More complex models possible (which ones?)
  - Different models for main-memory / distributed DBMSs (which ones?)
- Parameters:
  - $B(R)$  = # of blocks (i.e., pages) for relation  $R$
  - $T(R)$  = # of tuples in relation  $R$
  - $V(R, a)$  = # of distinct values of attribute  $a$

# Cost

---

- Cost of an operation = number of disk I/Os to
  - read the operands
  - compute the result
- Cost of writing the result to disk is *not included*
  - Need to count it separately when applicable

# Notions of Clustering

---

- **Co-clustering**

- Tuples of one relation R are placed with a tuple of another relation S with a common value (*uncommon*)

- **Clustered relation**

- Tuples of relation are stored on blocks predominantly devoted to storing *that* relation
- Sometimes also called “clustered file organization”
- We will focus on this in the following

- **Clustered index (aka clustering index)**

- When ordering of data records is close to the ordering of data entries in the index

# Cost Parameters

---

- Clustered relation R:
  - Blocks consists mostly of records from this table
  - $B(R) \approx T(R) / \text{blockSize}$
- Unclustered relation R:
  - Its records are placed on blocks with other tables
  - When R is unclustered:  $B(R) \approx T(R)$
- When a is a key,  $V(R,a) = T(R)$
- When a is not a key,  $V(R,a)$

# Cost of Scanning a Table

---


- Clustered relation:
  - Result may be unsorted:  $B(R)$
  - Result needs to be sorted:  $3B(R)$  (see later)

# One-pass Algorithms

---

Selection  $\sigma(R)$ , projection  $\Pi(R)$

- Both are ***tuple-at-a-time*** algorithms
- Cost:  $B(R)$ , the cost of scanning the relation



# Join Algorithms

---

- Logical operator:
  - $\text{Product}(\text{pname}, \text{cname}) \bowtie \text{Company}(\text{cname}, \text{city})$
- Propose three physical operators for the join, assuming the tables are in main memory:
  - **Hash join**
  - **Nested loop join**
  - **Sort-merge join**

# Hash Join

---

Hash join:  $R \bowtie S$

- Scan R, build buckets in main memory
- Then scan S and join
- Cost:  $B(R) + B(S)$
- One pass algorithm when  $B(R) \leq M$



# Nested Loop Joins

---

- Tuple-based nested loop  $R \bowtie S$
- R is the outer relation, S is the inner relation

```
for each tuple r in R do  
    for each tuple s in S do  
        if r and s join then output (r,s)
```

- Cost:  $B(R) + T(R) B(S)$  when S is clustered

# Page-at-a-time Refinement

---

```
for each page of tuples r in R do  
  for each page of tuples s in S do  
    for all pairs of tuples  
      if r and s join then output (r,s)
```

- Cost:  $B(R) + B(R)B(S)$  if S is clustered

# Nested Loop Joins

---

- We can be cleverer
- How would you compute the join in the following cases ?  
What is the cost ?
  - $B(R) = 1000, B(S) = 2, M = 4$
  - $B(R) = 1000, B(S) = 3, M = 4$
  - $B(R) = 1000, B(S) = 6, M = 4$

# Block Nested Loop Joins

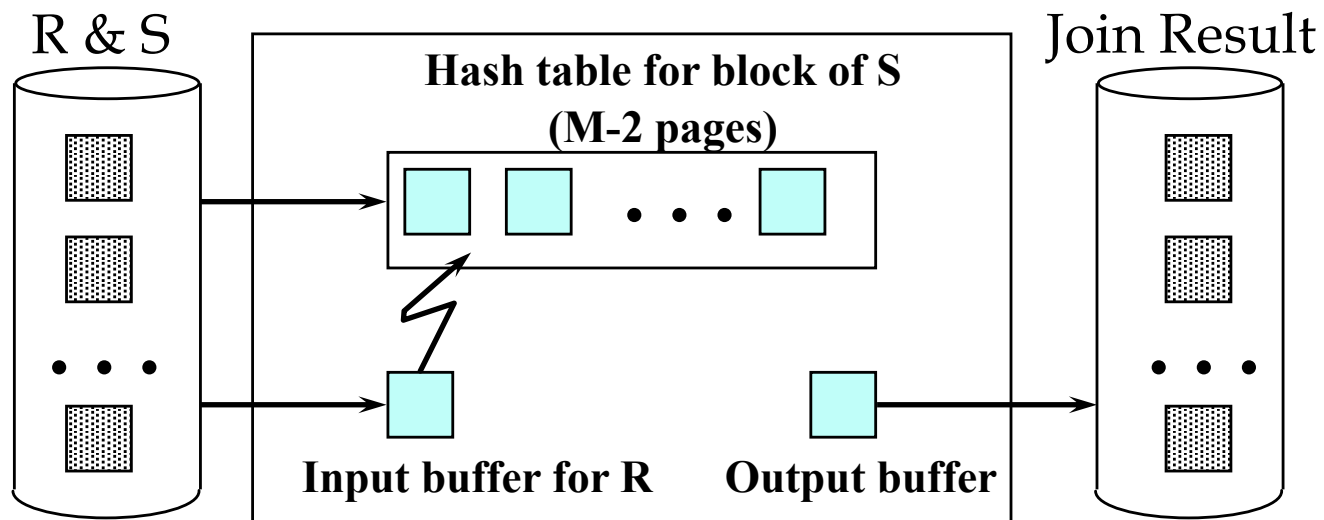
---

- Block Nested Loop Join
- Group of (M-2) pages of S is called a “block”

```
for each (M-2) pages ps of S do  
  for each page pr of R do  
    for each tuple s in ps  
      for each tuple r in pr do  
        if “r and s join” then output(r,s)
```

# Block Nested Loop Joins

---



# Block Nested Loop Joins

---

- Cost of block-based nested loop join
  - Read S once: cost  $B(S)$
  - Outer loop runs  $B(S)/(M-2)$  times, and each time need to read R: costs  $B(S)B(R)/(M-2)$
  - Total cost:  $B(S) + B(S)B(R)/(M-2)$
- Notice: it is better to iterate over the smaller relation first

# Sort-Merge Join

---

Sort-merge join:  $R \bowtie S$

- Scan R and sort in main memory
- Scan S and sort in main memory
- Merge R and S
- Cost:  $B(R) + B(S)$
- One pass algorithm when  $B(S) + B(R) \leq M$
- Typically, this is NOT a one pass algorithm

# More One-pass Algorithms

---

Duplicate elimination  $\delta(R)$

- Need to keep tuples in memory
- When new tuple arrives, need to compare it with previously seen tuples
- Balanced search tree or hash table
- Cost:  $B(R)$
- Assumption:  $B(\delta(R)) \leq M$



# Even More One-pass Algorithms

---

Grouping:

$\gamma_{\text{department, sum(quantity)}}(\text{Product})$

How can we compute this in main memory ?

# Even More One-pass Algorithms

---

- Grouping:  $\gamma_{\text{department, sum(quantity)}}(R)$
- Need to store all departments in memory
- Also store the  $\text{sum(quantity)}$  for each department
- Balanced search tree or hash table
- Cost:  $B(R)$
- Assumption: number of depts fits in memory

# Outline

---

- **Steps involved in processing a query**
  - Logical query plan
  - Physical query plan
  - Query execution overview
- **Operator implementations**
  - One pass algorithms
  - Two-pass algorithms
  - Index-based algorithms
- **Query optimization 101**


# Two-Pass Algorithms

---

- What if data does not fit in memory?
- Need to process it in multiple passes
- Two key techniques
  - Hashing
  - Sorting

# Two Pass Algorithms Based on Hashing

- Idea: partition a relation R into partitions (buckets), on disk
- Each bucket has size approx.  $B(R)/M$



- Does each partition fit in main memory ?
  - Yes if  $B(R)/M \leq M$ , i.e.  $B(R) \leq M^2$

# Hash Based Algorithms for $\delta$

---

- Recall:  $\delta(R)$  = duplicate elimination
- Step 1. Partition  $R$  into buckets
- Step 2. Apply  $\delta$  to each bucket
- Cost:  $3B(R)$
- Assumption:  $B(R) \leq M^2$

# Partitioned (Grace) Hash Join


---

$R \bowtie S$

- Step 1:
  - Hash S into M-1 buckets
  - Send all buckets to disk
- Step 2
  - Hash R into M-1 buckets
  - Send all buckets to disk
- Step 3
  - Join every pair of buckets

# Partitioned Hash Join

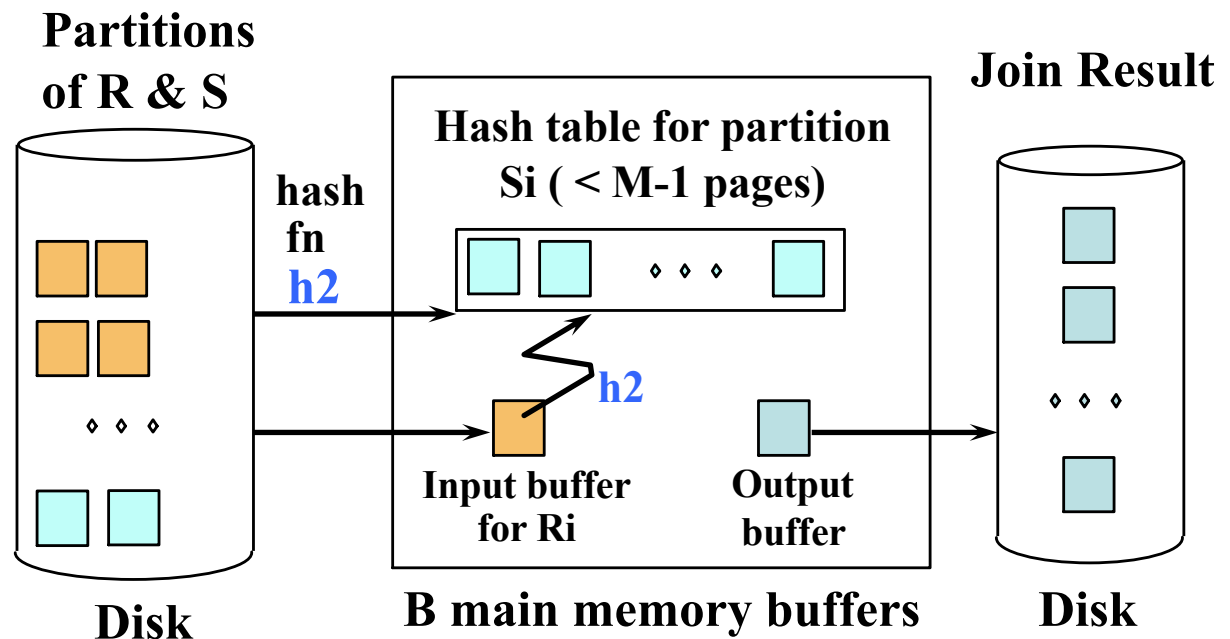
- Partition both relations using hash fn **h**
- R tuples in partition i will only match S tuples in partition i.





# Partitioned Hash Join

- Read in partition of R, hash it using  $h_2$  ( $\neq h$ )
  - Build phase
- Scan matching partition of S, search for matches
  - Probe phase



# Partitioned Hash Join

---

- Cost:  $3B(R) + 3B(S)$
- Assumption:  $\min(B(R), B(S)) \leq M^2$

# External Sorting


---

- Problem: Sort a file of size  $B$  with memory  $M$
- Where we need this:
  - ORDER BY in SQL queries
  - Several physical operators
  - Bulk loading of B+-tree indexes.
- Will discuss only 2-pass sorting, for when  $B < M^2$

# External Merge-Sort: Step 1

---


- Phase one: load  $M$  pages in memory, sort



# External Merge-Sort: Step 2

---

- Merge  $M - 1$  (max, due to memory) runs into a new run



If  $B(R) < M^2$  then we have  $M-1$  runs max and we are done (slightly more complex otherwise)<sup>53</sup>

# External Merge-Sort

---

- Cost:
  - Read+write+read =  $3B(R)$
  - Assumption:  $B(R) \leq M^2$
- Other considerations
  - In general, a lot of optimizations are possible

# Two-Pass Algorithms Based on Sorting

---

Duplicate elimination  $\delta(R)$

- Trivial idea: sort first, then eliminate duplicates
- Step 1: sort chunks of size  $M$ , write
  - cost  $2B(R)$
- Step 2: merge  $M-1$  runs, but include each tuple only once
  - cost  $B(R)$
- Total cost:  $3B(R)$ , Assumption:  $B(R) \leq M^2$

# Two-Pass Algorithms Based on Sorting

---

Grouping:  $\gamma_{a, \text{sum}(b)}(R)$

- Same as before: sort, then compute the  $\text{sum}(b)$  for each group of  $a$ 's
- Total cost:  $3B(R)$
- Assumption:  $B(R) \leq M^2$



# Two-Pass Algorithms Based on Sorting

---

Join  $R \bowtie S$

- Start by sorting both  $R$  and  $S$  on the join attribute:
  - Cost:  $4B(R)+4B(S)$  (because need to write to disk)
- Read both relations in sorted order, match tuples
  - Cost:  $B(R)+B(S)$
- Total cost:  $5B(R)+5B(S)$
- Assumption:  $B(R) \leq M^2$ ,  $B(S) \leq M^2$

# Two-Pass Algorithms Based on Sorting

---

Join  $R \bowtie S$

- If the number of tuples in  $R$  matching those in  $S$  is small (or vice versa) (e.g., If  $B(R) + B(S) \leq M^2$ ) we can compute the join during the merge phase
- Total cost:  $3B(R) + 3B(S)$

# Outline

---

- **Steps involved in processing a query**
  - Logical query plan
  - Physical query plan
  - Query execution overview
- **Operator implementations**
  - One pass algorithms
  - Two-pass algorithms
  - Index-based algorithms
- **Query optimization 101**

# Review: Access Methods

---

- **Heap file**
  - Scan tuples one at the time
- **Hash-based index**
  - Efficient selection on equality predicates
- **Tree-based index**
  - Efficient selection on equality or range predicates

# Index Based Selection

---

- Selection on equality:  $\sigma_{a=v}(R)$
- $V(R, a) = \#$  of distinct values of attribute  $a$
- Clustered index on  $a$ : cost  $B(R)/V(R,a)$
- Unclustered index on  $a$ : cost  $T(R)/V(R,a)$
- Note: we ignored the I/O cost for the index pages

# Index Based Selection

---

- Example:

$$\begin{aligned}B(R) &= 2000 \\T(R) &= 100,000 \\V(R, a) &= 20\end{aligned}$$

$$\text{cost of } \sigma_{a=v}(R) = ?$$

- Table scan (assuming R is clustered)
  - $B(R) = 2,000$  I/Os
- Index based selection
  - If index is clustered:  $B(R)/V(R,a) = 100$  I/Os
  - If index is unclustered:  $T(R)/V(R,a) = 5,000$  I/Os
- Lesson
  - Don't build unclustered indexes when  $V(R,a)$  is small !

# Index Nested Loop Join

---

$R \bowtie S$

- Assume  $S$  has an index on the join attribute
- Iterate over  $R$ , for each tuple fetch corresponding tuple(s) from  $S$
- Cost:
  - Assuming  $R$  is clustered
  - If index on  $S$  is clustered:  $B(R) + T(R)B(S)/V(S,a)$
  - If index on  $S$  is unclustered:  $B(R) + T(R)T(S)/V(S,a)$

# Summary of External Join Algorithms

---

- Block Nested Loop Join:  $B(R) + B(R) \cdot B(S) / M$
- Sort-Merge Join:  $3B(R) + 3B(S)$   
Assuming  $B(R) + B(S) \leq M^2$
- Index Nested Loop Join:  $B(R) + T(R)B(S)/V(S,a)$   
Assuming R is clustered and S has clustered index on a



# Summary of Query Execution

---

- For each logical query plan
  - There exist many physical query plans
  - Each plan has a different cost
  - Cost depends on the data
- Additionally, for each query
  - There exist several logical plans

# Outline

---

- **Steps involved in processing a query**
  - Logical query plan
  - Physical query plan
  - Query execution overview
- **Operator implementations**
  - One pass algorithms
  - Two-pass algorithms
  - Index-based algorithms
- **Query optimization 101**

# Query Optimization Algorithm

---

- For a query
  - There exists many physical query plans
  - Query optimizer needs to pick a good one
- Basic query optimization algorithm
  - Enumerate alternative plans
  - Compute estimated cost of each plan
    - Compute number of I/Os
    - Optionally take into account other resources
  - Choose plan with lowest cost
  - This is called cost-based optimization

# Estimating Cost of a Query Plan

---

- We already know how to
  - Compute the cost of different operations in terms of number IOs
- We still need to
  - Compute cost of retrieving tuples from disk with different access paths (for more sophisticated predicates than equality)
  - Compute cost of a complete plan

# Access Path Selection

---

- `Supplier(sid,sname,scity,sstate)`
- Selection condition: `sid > 300 ∧ scity='Seattle'`
- Indexes: B+-tree on `sid` and B+-tree on `scity`
- Which access path should we use?
- We should pick the **most selective** access path

# Access Path Selectivity

---

- **Access path selectivity is the number of pages retrieved if we use this access path**
  - Most selective retrieves fewest pages
- As we saw earlier, **for equality predicates**
  - Selection on equality:  $\sigma_{a=v}(R)$
  - $V(R, a) = \#$  of distinct values of attribute  $a$
  - $1/V(R,a)$  is thus the reduction factor
  - Clustered index on  $a$ : cost  $B(R)/V(R,a)$
  - Unclustered index on  $a$ : cost  $T(R)/V(R,a)$
  - (we **are ignoring I/O cost of index pages** for simplicity)

# Selectivity for Range Predicates

---

Selection on range:  $\sigma_{a>v}(R)$

- How to compute the selectivity?
- Assume values are uniformly distributed
- Reduction factor  $X$
- $X = (\text{Max}(R,a) - v) / (\text{Max}(R,a) - \text{Min}(R,a))$
- Clustered index on  $a$ : cost  $B(R)*X$
- Unclustered index on  $a$ : cost  $T(R)*X$

# Example

---

- Selection condition: **sid > 300**
  - Index I1: B+-tree on sid clustered
- Selection condition: **scity = 'Fribourg'**
  - Index I2: B+-tree on scity unclustered
- Let's assume
  - $V(\text{Supplier}, \text{scity}) = 20$
  - $\text{Max}(\text{Supplier}, \text{sid}) = 1000, \text{Min}(\text{Supplier}, \text{sid}) = 1$
  - $B(\text{Supplier}) = 100, T(\text{Supplier}) = 1000$
- **Cost I1:  $B(R) * (\text{Max}-v)/(\text{Max}-\text{Min}) = 100 * 700 / 999 \approx 70$**
- **Cost I2:  $T(R) * 1/V(\text{Supplier}, \text{scity}) = 1000 / 20 = 50$**



# Selectivity with Multiple Conditions

---

What if we have an index on multiple attributes?

- Example selection  $\sigma_{a=v1 \wedge b=v2}(R)$  and index on  $\langle a, b \rangle$

How to compute the selectivity?

- Assume attributes are independent
- $X = 1 / (V(R, a) * V(R, b))$
- Clustered index on  $\langle a, b \rangle$ : cost  $B(R) * X$
- Unclustered index on  $\langle a, b \rangle$ : cost  $T(R) * X$

# Back to Estimating Cost of a Query Plan

---

- We already know how to
  - Compute the cost of different operations
  - Compute cost of retrieving tuples from disk with different access paths (for more sophisticated predicates than equality)
- We still need to
  - Compute cost of a complete plan

# Computing the Cost of a Plan

---

- Collect statistical summaries of stored data
- Compute cost in a bottom-up fashion
- For each operator compute
  - Estimate cost of executing the operation
  - Estimate statistical summary of the output data

# Statistics on the Output Data

---

- Most important piece of information
  - **Size of operator result**
  - I.e., the number of output tuples
- **Projection**: output size same as input size
- **Selection**: multiply input size by reduction factor
  - Similar to what we did for estimating access path selectivity
  - Assume independence between conditions in the predicate
  - (use product of the reduction factors for the terms)

# Estimating Result Sizes

---

- For **joins**  $R \bowtie S$ 
  - Take product of cardinalities of relations R and S
  - Apply reduction factors for each term in join condition
  - Terms are of the form: column1 = column2
  - Reduction:  $1 / (\text{MAX}(V(R, \text{column1}), V(S, \text{column2})))$
  - **Assumes each value in smaller set has a matching value in the larger set**

# Our Example

---

- Suppliers(sno,sname,scity,sstate)
- Supplies(pno,sno,quantity)
- Some statistics
  - T(Supplier) = 1000 records
  - B(Supplier) = 100 pages
  - T(Supplies) = 10,000 records
  - B(Supplies) = 100 pages
  - V(Supplier,scity) = 20, V(Supplier,state) = 10
  - V(Supplies,pno) = 3,000
  - Both relations are clustered

# Physical Query Plan 1

---

(On the fly)


$\pi_{\text{sname}}$

Selection and project on-the-fly  
-> No additional cost.

(On the fly)

$\sigma_{\text{scity}='Seattle' \wedge \text{sstate}='WA' \wedge \text{pno}=2}$

(Nested loop)

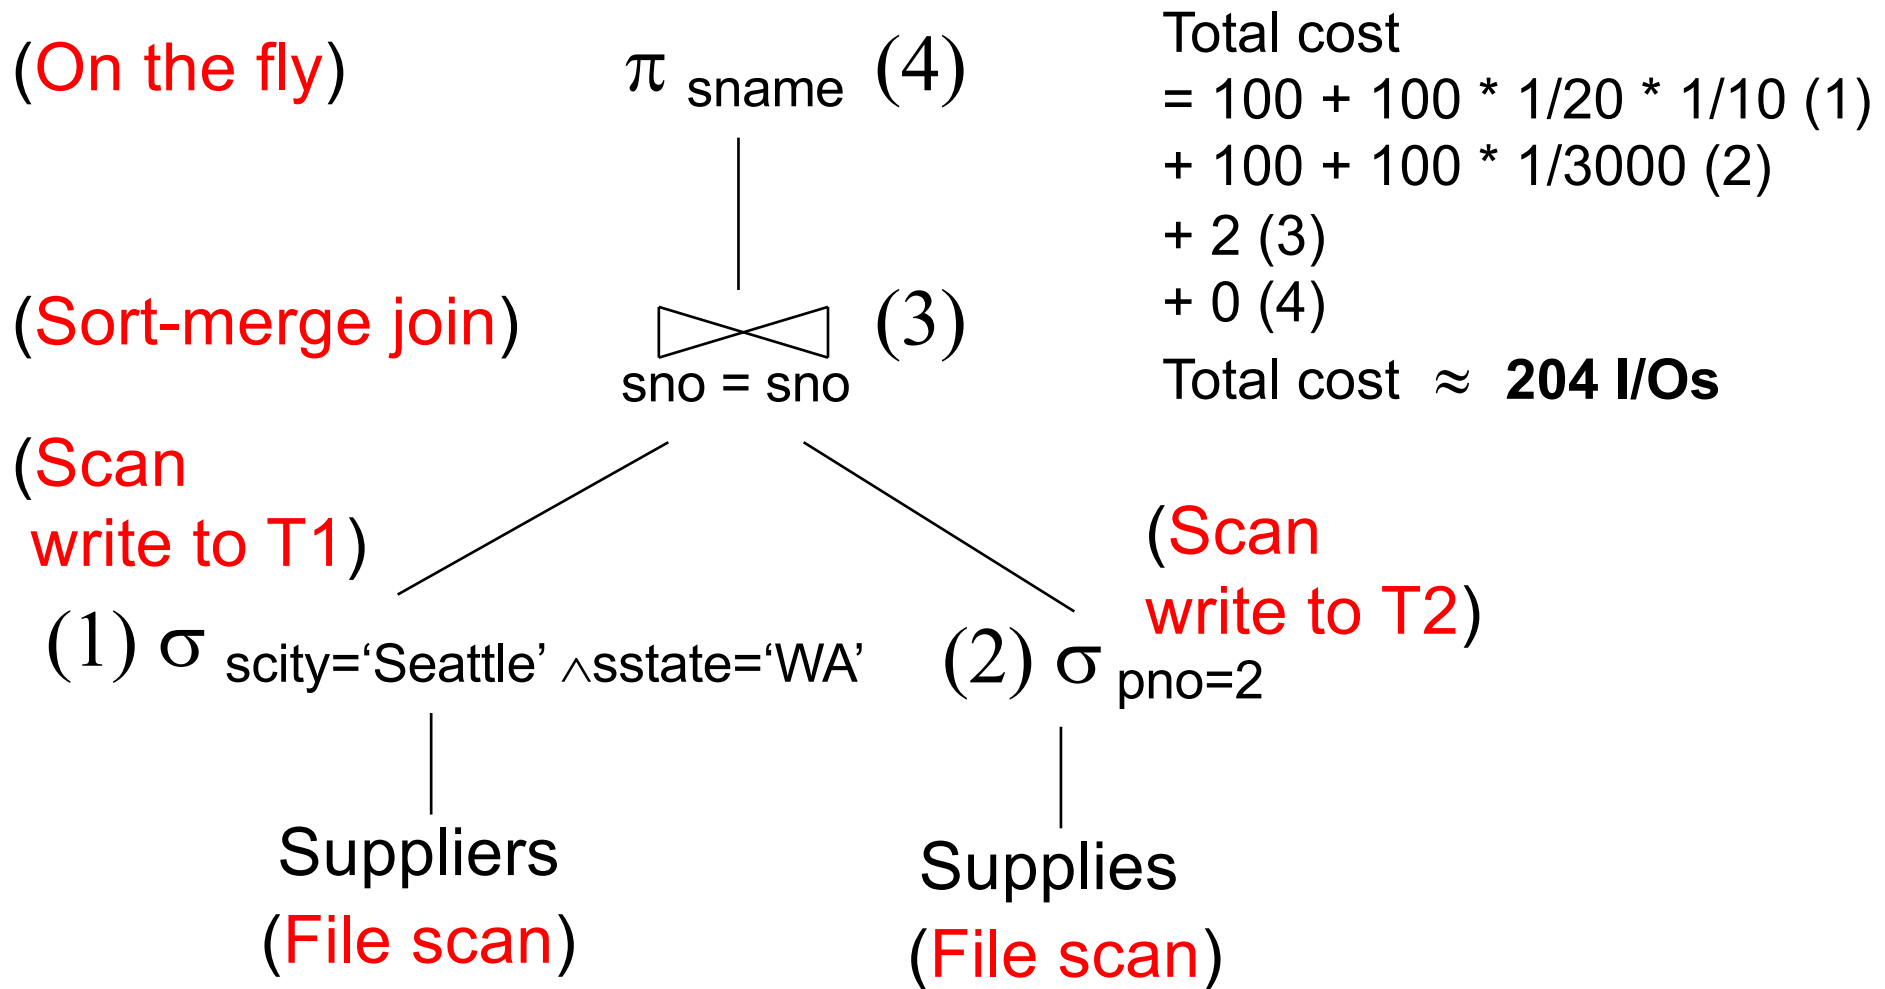
  
 $\text{sno} = \text{sno}$

Total cost of plan is thus cost of join:  
=  
 $B(\text{Supplier}) + B(\text{Supplier}) * B(\text{Supplies})$   
 $= 100 + 100 * 100$   
 **$= 10,100 \text{ I/Os}$**

Suppliers  
(File scan)

Supplies  
(File scan)

# Physical Query Plan 2





# Plan 2 with Different Numbers

What if we had:

10K pages of Suppliers

10K pages of Supplies


(Sort-merge join)

(Scan  
write to T1)

(1)  $\sigma_{\text{scity}='Seattle' \wedge \text{sstate}='WA'}$

Suppliers  
(File scan)

$\pi_{\text{sname}}$  (4)



(3)

(Scan  
write to T2)

(2)  $\sigma_{\text{pno}=2}$

Supplies  
(File scan)

Total cost

= 10000 + 50 (1)

+ 10000 + 4 (2)

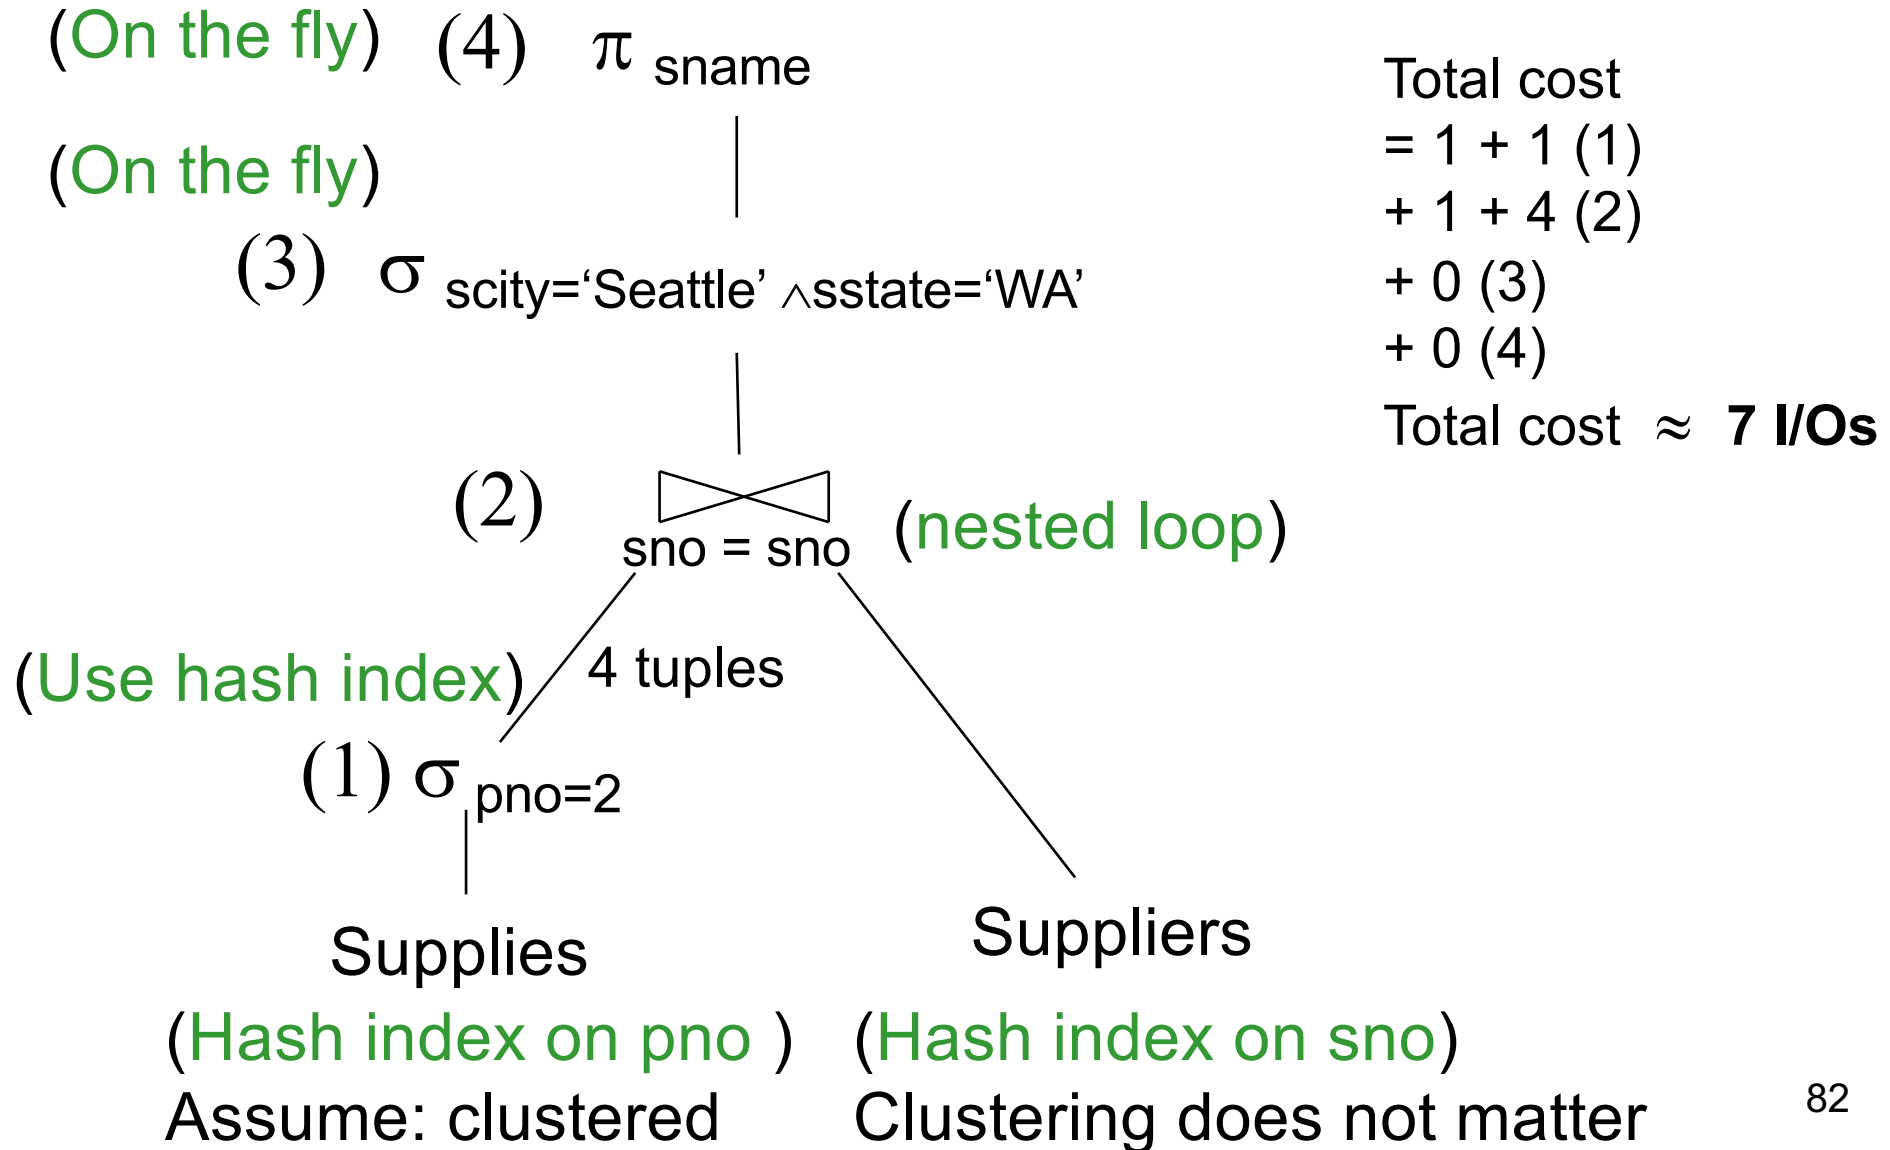
+ 5\*50 + 5\*4 (3)

+ 0 (4)

Total cost  $\approx$  20,324 I/Os

Assuming naive  
two-pass sort  
algorithm

# Physical Query Plan 3



# Simplifications

---

- In the previous examples, we assumed that all index pages were in memory
- When this is not the case, we need to add the cost of fetching index pages from disk

# Relational Algebra Equivalences

---

- Selections

- Commutative:  $\sigma_{c_1}(\sigma_{c_2}(R))$  same as  $\sigma_{c_2}(\sigma_{c_1}(R))$
- Cascading:  $\sigma_{c_1 \wedge c_2}(R)$  same as  $\sigma_{c_2}(\sigma_{c_1}(R))$

- Projections


- Cascading

- Joins


- Commutative :  $R \bowtie S$  same as  $S \bowtie R$
- Associative:  $R \bowtie (S \bowtie T)$  same as  $(R \bowtie S) \bowtie T$

# Left-Deep Plans and Bushy Plans

---



Left-deep plan



Bushy plan

# Search Space Challenges

---

- Search space is huge!
  - Many possible equivalent trees
  - Many implementations for each operator
  - Many access paths for each relation
- Cannot consider ALL plans
- Want a search space that includes low-cost plans

# System R Search Space

---

- Only left-deep plans
  - Enable dynamic programming for enumeration
  - Facilitate tuple pipelining from outer relation
- Consider plans with all “interesting orders”
- Perform cross-products after all other joins (heuristic)
- Only consider nested loop & sort-merge joins
- Consider both file scan and indexes
- Try to evaluate predicates early
- Many other search strategies possible