



Technical Deep-Dive in a Column-Oriented In-Memory Database

Slides by Martin Grund
Presented today by Alberto Lerner

eXascale Infolab
Department of Informatics
University of Fribourg Switzerland



Goals

Deep technical understanding of a column-oriented, dictionary-encoded in-memory database and its application in enterprise computing

Chapters

- The future of enterprise computing
- Foundations of database storage techniques
- In-memory database operators
- Advanced database storage techniques
- Implications on Application Development



Enterprise Application Characteristics



OLTP vs. OLAP

Online Transaction Processing

Online Analytical Processing

- Modern enterprise resource planning (ERP) systems are challenged by **mixed workloads**, including OLAP-style queries. For example:
 - OLTP-style: create sales order, invoice, accounting documents, display customer master data or sales order
 - OLAP-style: dunning, available-to-promise, cross selling, operational reporting (list open sales orders)
- But: Today's data management systems are optimized **either** for daily **transactional** or **analytical** workloads storing their data along rows or columns



Drawbacks of the Separation

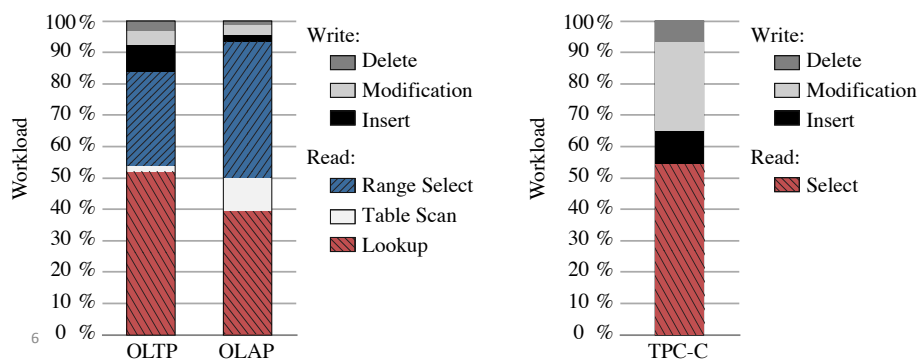
- ❑ OLAP systems do not have the **latest** data
- ❑ OLAP systems only have **predefined subset** of the data
- ❑ **Cost-intensive ETL** processes have to sync both systems
- ❑ Separation introduces data **redundancy**
- ❑ **Different data schemas** introduce complexity for applications combining sources

5





Enterprise Workloads are Read-Mostly

- Workload in enterprise applications constitutes of
 - Mainly read queries (OLTP 83%, OLAP 94%)
 - Many queries access large sets of data



6



 UNIVERSITAS FRIBURGENSIS

Vision

Combine OLTP and OLAP data
 using **modern** hardware and database systems
 to create a **single source of truth**,
 enable **real-time analytics** and
simplify applications and database structures.

Additionally,

- ☐ Extraction, transformation, and loading (ETL) processes
- ☐ Pre-computed aggregates and materialized views

become (almost) obsolete.

8



Enterprise Data Characteristics

- Many columns are not used even once
- Many columns have a low cardinality of values
- NULL values/default values are dominant
- Sparse distribution facilitates high compression

Standard enterprise software data is **sparse and wide**.


9



Many Columns are not Used Even Once

55% unused columns per company in average

40% unused columns across all analyzed companies



10




Changes in Hardware



Changes in Hardware...

... give an opportunity to re-think the assumptions of yesterday because of what is possible today.

- Multi-Core Architecture (96 cores per server)
- Large main memories: 2TB/blade
- One blade ~\$50.000 = 1 enterprise class server
- Parallel scaling across blades



- 64 bit address space
- 2TB in current servers
- Cost-performance ratio rapidly declining
- Memory hierarchies


- Main Memory becomes **cheaper and larger**



In the Meantime Research has come up with...

... several advances in software for processing data

- Column-oriented data organization (the column store)
 - **Sequential** scans allow best bandwidth utilization between CPU cores and memory
 - **Independence** of tuples allows easy partitioning and therefore parallel processing
- Lightweight Compression
 - Reducing data amount, while..
 - Increasing processing speed through late materialization
- And more, e.g., parallel scan/join/aggregation



13



Foundations of Database Storage Techniques



Data Layout in Main Memory




Memory Basics (I)

- ❑ Memory in todays computers has a linear address layout: addresses start at 0x0 and go to 0xFFFFFFFFFFFFFFFF for 64bit
- ❑ Each process has its own virtual address space
- ❑ Virtual memory allocated by the program can distribute over multiple physical memory locations
- ❑ Address translation is done in hardware by the CPU



Memory Basics (2)

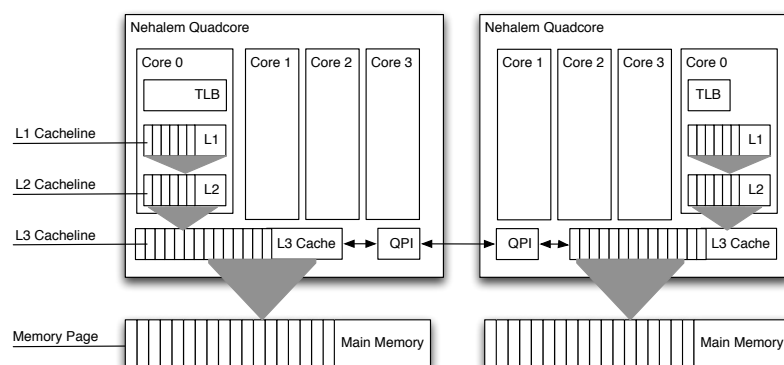


- ❑ Memory layout is **only linear**
- ❑ Every higher-dimensional access (like two-dimensional database tables) is mapped to this linear band

17



Memory Hierarchy





18



Physical Data Representation

- **Row store:**
 - Rows are stored consecutively
 - Optimal for row-wise access (e.g. SELECT *)
- **Column store:**
 - Columns are stored consecutively
 - Optimal for attribute focused access (e.g. SUM, GROUP BY)
- Note: concept is **independent** from storage type




19



Row Data Layout

- Data is stored tuple-wise
- Leverage co-location of attributes for a single tuple
- Low cost for reconstruction, but higher cost for sequential scan of a single attribute




20




Columnar Data Layout

- Data is stored attribute-wise
- Leverage sequential scan-speed in main memory for predicate evaluation
- Tuple reconstruction is more expensive

Column Operation



Row Operation



21



Dictionary Encoding



Motivation

- Main memory access is the new bottleneck
- Idea: **Trade** CPU time to compress and decompress data
- Compression reduces number of memory accesses
- Leads to **less** cache misses due to more information on a cache line
- Operation directly on compressed data
- Offsetting with bit-encoded fixed-length data types
- Based on limited value domain

23



Dictionary Encoding Example

8 billion humans

- Attributes
 - first name
 - last name
 - gender
 - country
 - city
 - birthday
 - → 200 byte
- Each attribute is stored dictionary encoded



24



Sample Data

rec ID	fname	lname	gender	city	country	birthday
...
39	John	Smith	m	Chicago	USA	12.03.1964
40	Mary	Brown	f	London	UK	12.05.1964
41	Jane	Doe	f	Palo Alto	USA	23.04.1976
42	John	Doe	m	Palo Alto	USA	17.06.1952
43	Peter	Schmidt	m	Potsdam	GER	11.11.1975
...

25



Dictionary Encoding a Column

- A column is split into a dictionary and an attribute vector
- Dictionary stores all distinct values with implicit value ID
- Attribute vector stores value IDs for all entries in the column
- Position is implicit, not stored explicitly
- Enables offsetting with fixed-length data types

Rec ID	fname	Dictionary for "fname"		Attribute Vector for "fname"	
...	...	Value ID	Value	position	Value ID
39	John
40	Mary	23	John	39	23
41	Jane	24	Mary	40	24
42	John	25	Jane	41	25
43	Peter	26	Peter	42	23
...	43	26

26



Sorted Dictionary

- Dictionary entries are sorted either by their numeric value or lexicographically
 - Dictionary lookup complexity: $O(\log(n))$ instead of $O(n)$
- Dictionary entries can be compressed to reduce the amount of required storage
- Selection criteria with ranges are less expensive (order-preserving dictionary)

27




Data Size Examples

Column	Cardinality	Bits Needed	Item Size	Plain Size	Size with Dictionary (Dictionary + Column)	Compression Factor
First names	5 million	23 bit	50 Byte	373GB	238.4MB + 21.4GB	≈17
Last names	8million	23 bit	50 Byte	373GB	381.5MB + 21.4GB	≈17
Gender	2	1 bit	1 Byte	7GB	2.0b + 953.7MB	≈8
City	1million	20 bit	50 Byte	373GB	47.7MB + 18.6GB	≈20
Country	200	8 bit	47 Byte	350GB	9.2KB + 7.5GB	≈47
Birthday	40000	16 bit	2 Byte	15GB	78.1KB + 14.9GB	≈1
Totals			200 Byte	≈ 1.6TB	≈ 92GB	≈17

28

Compression



Compression Techniques

- ❑ Heavy weight vs. light weight techniques
- ❑ Focus on light weight techniques for databases

- ❑ For attribute vector
 - Prefix encoding
 - Run length encoding
 - Cluster encoding
 - Sparse encoding
 - Indirect encoding

- ❑ For dictionary
 - Delta compression for strings
 - Other data types are stored as sorted arrays



Example Table

recID	fname	lname	gender	country	city	birthday	2nd_nationality
0	Martin	Albrecht	m	GER	Berlin	08-05-1955	n/a
1	Michael	Berg	m	GER	Berlin	03-05-1970	n/a
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968	n/a
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992	US
4	Ulrike	Schulze	f	GER	Potsdam	09-03-1977	n/a
5	Martin	Schulz	m	GER	Mainz	06-04-1980	GER
6	Sushi	Pao	f	CN	Peking	09-12-1954	n/a
7	Chen	Su Wong	m	CN	Shanghai	27-06-1999	n/a
...

- 200 countries = 8 bit
- 1 million cities = 20 bit
- 100 different 2nd nationalities = 7 bit
- 5 million first names = 23 bit

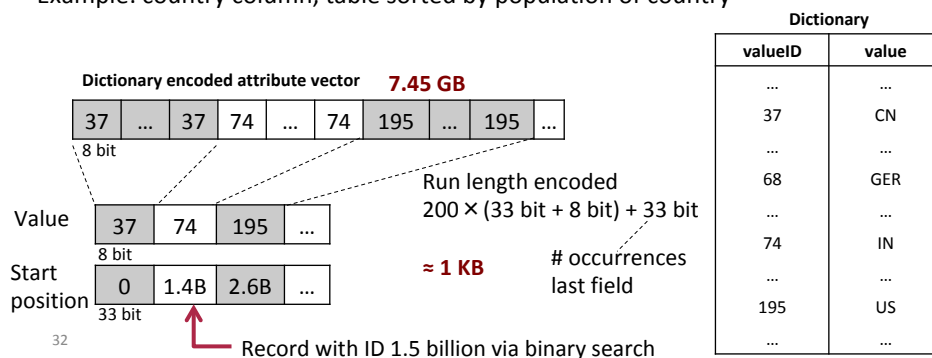
31



Run Length Encoding

- Replace sequence of the same value with a single instance of the value and
 - a) Its number of occurrences
 - b) Its start position (shown below) **Direct access!**
- Variant b) speeds up access compared to a)

Example: country column, table sorted by population of country






Sparse Encoding

- Remove the value v that appears most often
- A bit vector indicates at which positions v was removed from the original sequence

Example: 2nd nationality column, regardless of sorting order of table



33



Keep in Mind

- Most compression techniques require sorted sets, but a table can only be sorted by one column or cascading
- No direct access to rows in some cases, but offset has to be computed

34



In-Memory Database Operators



Scan Performance (I)

8 billion humans

☐ Attributes

- First Name
- Last Name
- Gender
- Country
- City
- Birthday
- 200 byte

☐ Question: How many men/women?

☐ Assumed scan speed: 2MB/ms/core





Scan Performance (2)

Row Store – Layout

Table: humans

	First Name	Last Name	Gender	Country	City	Birthday
Row 1						
Row 2						
Row 3						
...						
Row 8×10^9						

- Table size =
8 billion tuples x
200 bytes per tuple
→ **≈ 1.6 TB**
- Scan through all rows
with 2MB/ms/core
→ **≈ 800 seconds**
with 1 core

37



Scan Performance (3)

Row Store – Full Table Scan

Table: humans

	First Name	Last Name	Gender	Country	City	Birthday
Row 1						
Row 2						
Row 3						
...						
Row 8×10^9						

- Table size =
8 billion tuples x
200 bytes per tuple
→ **≈ 1.6 TB**
- Scan through all rows
with 2MB/ms/core
→ **≈ 800 seconds**
with 1 core

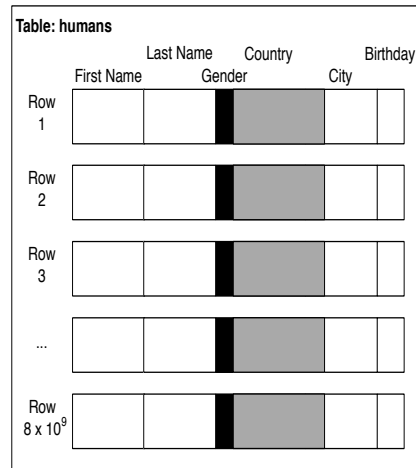
38

Data loaded and used
 Data loaded but not used



Scan Performance (4)

Row Store – Stride Access “Gender”



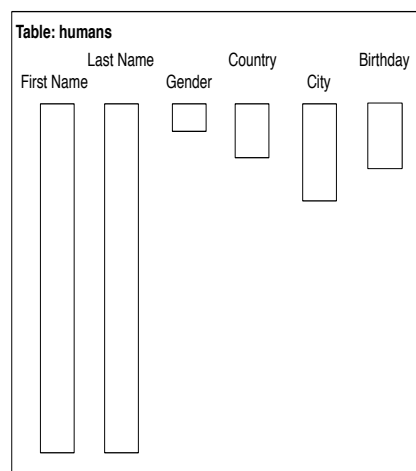
39

- 8 billion cache accesses à 64 byte
→ **≈512 GB**
- Read with 2MB/ms/core
→ **≈256 seconds**
with 1 core



Scan Performance (5)

Column Store – Layout



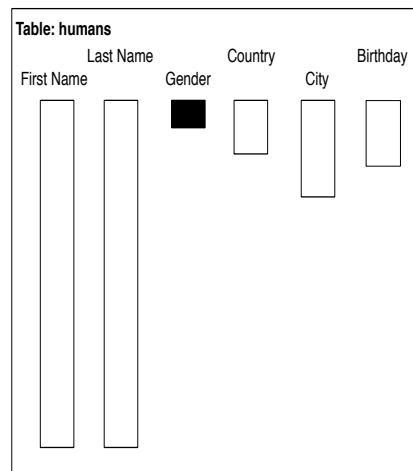
40

- Table size
 - Attribute vectors: **≈91 GB**
 - Dictionaries: **≈700 MB**
 - Total: **≈92 GB**
- Compression factor: **≈17**



Scan Performance (6)

Column Store – Full Column Scan on “Gender”

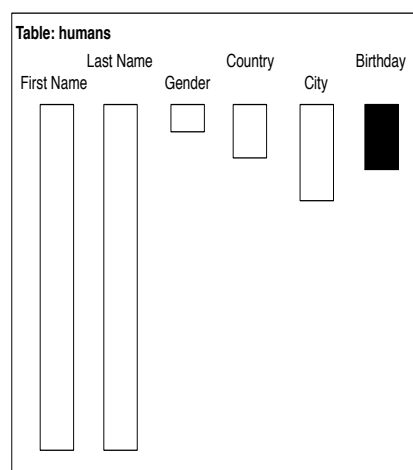


- Size of attribute vector “gender” =
8 billion tuples x
1 bit per tuple
→ **≈ 1 GB**
- Scan through attribute vector with
2MB/ms/core →
≈ 0.5 seconds with 1 core



Scan Performance (7)

Column Store – Full Column Scan on “Birthday”



- Size of attribute vector “birthday” =
8 billion tuples x
2 byte per tuple =
≈ 16 GB
- Scan through column with
2MB/ms/core →
≈ 8 seconds with 1 core



Scan Performance – Summary

□ How many women, how many men?

	Column Store	Row Store	
		Full table scan	Stride access
Time in seconds	0.5	800	256
		1,600x slower	512x slower

43



Tuple Reconstruction



Tuple Reconstruction (1)

Accessing a record in a row store

Table: world_population

	First Name	Last Name	Gender	Country	City	Birthday
Row 1						
Row 2						
Row 3						
Row 4						
...						
Row 8×10^9						

45

Data not loaded
 Data loaded and used
 Data loaded but not used

- All attributes are stored consecutively
- 200 byte \rightarrow 4 cache accesses à 64 byte \rightarrow 256 byte
- Read with 2MB/ms/core $\rightarrow \approx 0.128 \mu\text{s}$ with 1 core



Tuple Reconstruction (2)

Virtual record IDs

Table: world_population

First Name	Last Name	Gender	Country	City	Birthday
1	1	2	2	2	1
2	2	3	3	3	2
3	3	4	4	4	3
4	4	5	5	5	4
5	5	6	6	6	5
6	6	7	7	7	6
7	7	8	8	8	7
8	8	9	9	9	8
9	9	10	10	10	9
10	10	11	11	11	10
11	11	12	12	12	11
12	12	13	13	13	12
13	13	14	14	14	13
14	14	15	15	15	14
15	15	16	16	16	15
16	16	17	17	17	16
17	17	18	18	18	17
18	18	19	19	19	18
19	19	20	20	20	19
20	20	21	21	21	20
21	21	22	22	22	21
22	22	23	23	23	22
23	23	24	24	24	23
24	24	25	25	25	24
25	25	26	26	26	25
26	26	27	27	27	26
27	27	28	28	28	27
28	28	29	29	29	28
29	29	30	30	30	29
30	30	31	31	31	30
31	31	32	32	32	31
32	32	33	33	33	32
33	33	34	34	34	33
34	34	35	35	35	34
35	35	36	36	36	35
36	36	37	37	37	36
37	37	38	38	38	37
38	38	39	39	39	38
39	39	40	40	40	39
40	40	41	41	41	40
41	41	42	42	42	41
42	42	43	43	43	42
43	43	44	44	44	43
44	44	45	45	45	44
45	45	46	46	46	45
46	46	47	47	47	46
47	47	48	48	48	47
48	48	49	49	49	48
49	49	50	50	50	49
50	50	51	51	51	50
51	51	52	52	52	51
52	52	53	53	53	52
53	53	54	54	54	53
54	54	55	55	55	54
55	55	56	56	56	55
56	56	57	57	57	56
57	57	58	58	58	57
58	58	59	59	59	58
59	59	60	60	60	59
60	60	61	61	61	60
61	61	62	62	62	61
62	62	63	63	63	62
63	63	64	64	64	63
64	64	65	65	65	64
65	65	66	66	66	65
66	66	67	67	67	66
67	67	68	68	68	67
68	68	69	69	69	68
69	69	70	70	70	69
70	70	71	71	71	70
71	71	72	72	72	71
72	72	73	73	73	72
73	73	74	74	74	73
74	74	75	75	75	74
75	75	76	76	76	75
76	76	77	77	77	76
77	77	78	78	78	77
78	78	79	79	79	78
79	79	80	80	80	79
80	80	81	81	81	80
81	81	82	82	82	81
82	82	83	83	83	82
83	83	84	84	84	83
84	84	85	85	85	84
85	85	86	86	86	85
86	86	87	87	87	86
87	87	88	88	88	87
88	88	89	89	89	88
89	89	90	90	90	89
90	90	91	91	91	90
91	91	92	92	92	91
92	92	93	93	93	92
93	93	94	94	94	93
94	94	95	95	95	94
95	95	96	96	96	95
96	96	97	97	97	96
97	97	98	98	98	97
98	98	99	99	99	98
99	99	100	100	100	99

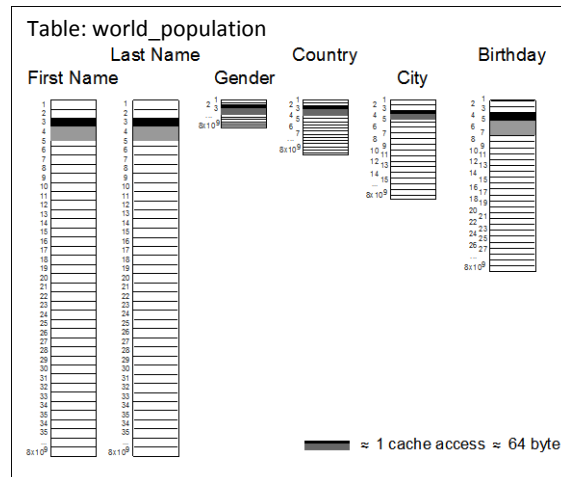
46

- All attributes are stored in separate columns
- Implicit record IDs are used to reconstruct rows



Tuple Reconstruction (3)

Virtual record IDs



- 1 cache access for each attribute
- 6 cache accesses à 64 byte
→ **384 byte**
- Read with 2MB/ms/core
→ **$\approx 0.192 \mu s$** with 1 core

47



Materialization Strategies



Tuple Reconstruction

Strategies:

- ❑ **Early** materialization
 - Decompress and decode data early and operate on strings
 - Create a row-wise data representation early on
- ❑ **Late** materialization
 - Operate on integers as long as possible
 - Operate on columns as long as possible

49



Example

Query:

List the number of Male inhabitants per city in Germany

```
SELECT city, COUNT(*)
FROM world_population
WHERE gender = "m"
AND country= "GER"
GROUP BY city
```

Table "world_population"

fname	lname	gender	country	city	birthday
Martin	Albrecht	m	GER	Berlin	08-05-1955
Michael	Berg	m	GER	Berlin	03-05-1970
Hanna	Schulze	f	GER	Bonn	04-04-1968
Ulrich	Schulze	m	GER	Bonn	10-20-1992
...

Dictionary encoded attribute vectors

53946	10435	0	68	357	15556
54368	25063	0	68	357	20882
30145	99645	1	68	443	20182
99312	99645	0	68	443	29147
fname	lname	gender	country	city	birthday

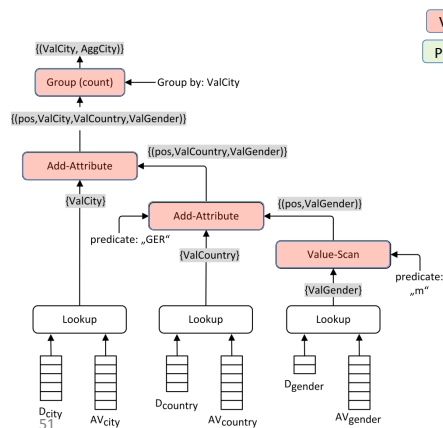
50



Comparison

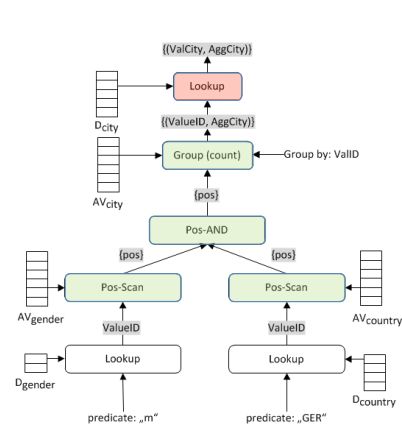
Early Materialization

Early dictionary lookup → working on materialized values



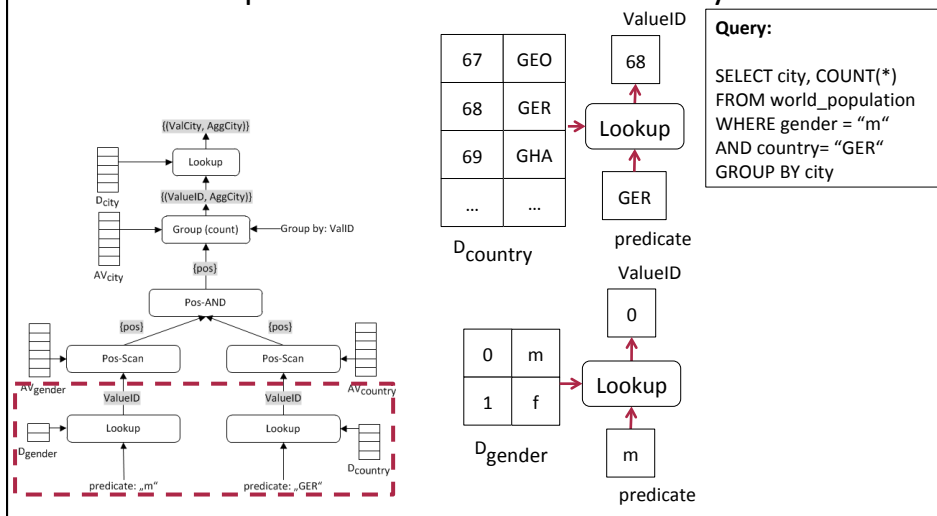
Late Materialization

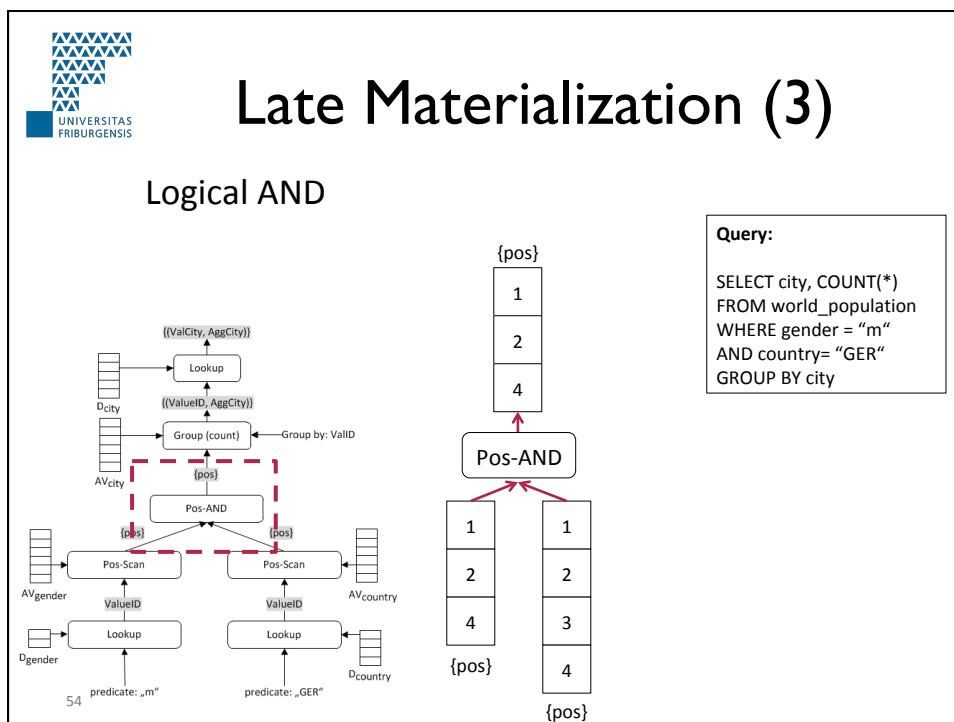
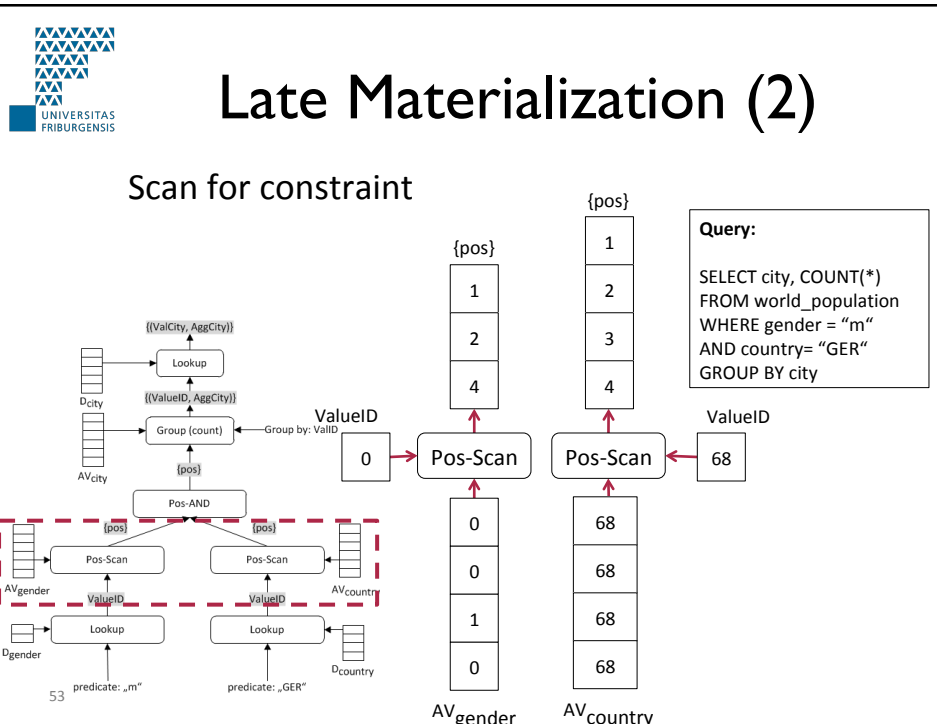
Late dictionary lookup → working with positional information



Late Materialization (I)

Look up constraint values in dictionary

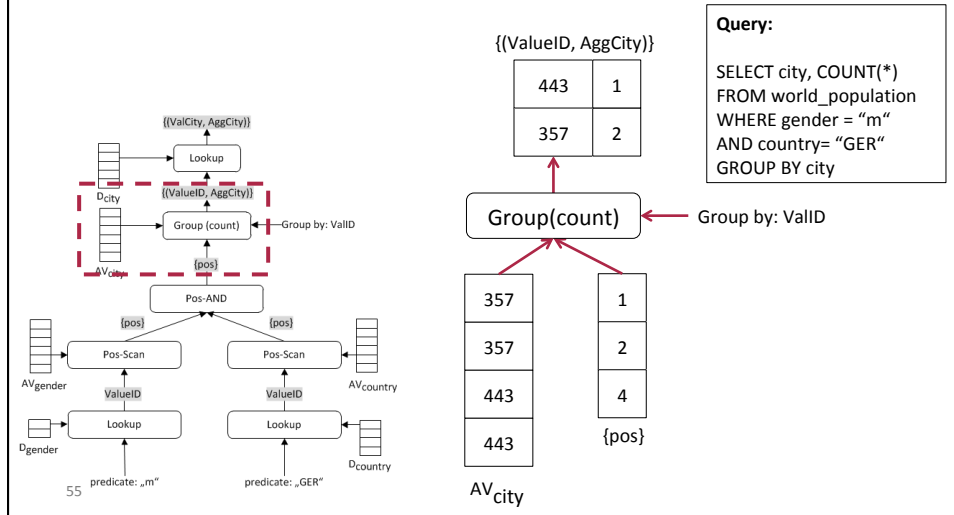






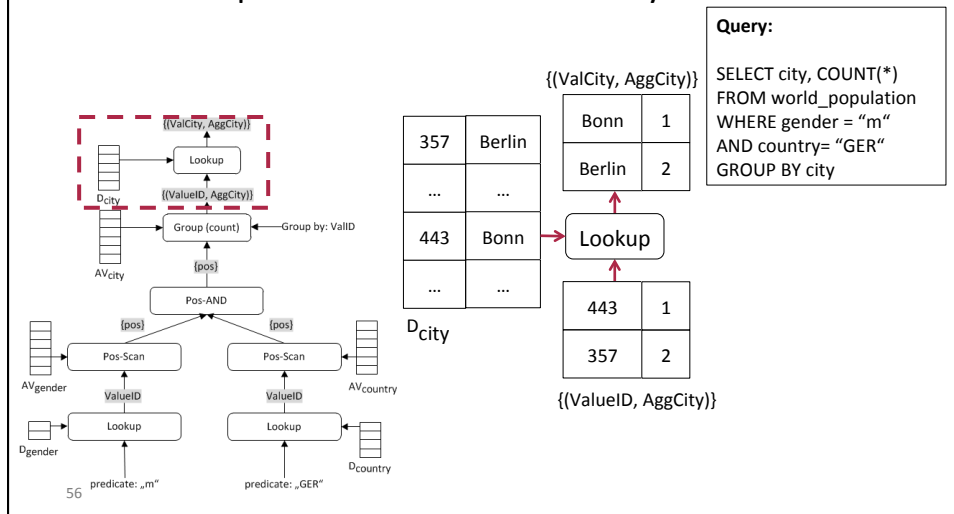
Late Materialization (4)

Filter attribute vector



Late Materialization (5)

Lookup values from the dictionary





Select



SELECT Example

```
SELECT fname, lname FROM world_population
WHERE country="Italy" and gender="m"
```

fname	lname	country	gender
Gianluigi	Buffon	Italy	m
Lena	Gercke	Germany	f
Mario	Balotelli	Italy	m
Manuel	Neuer	Germany	m
Lukas	Podolski	Germany	m
Klaas-Jan	Huntelaar	Netherlands	m



Query Plan

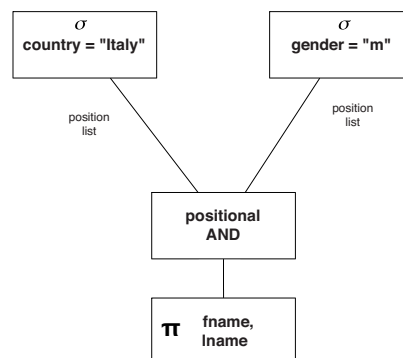
- ❑ Multiple plans exist to execute query
 - ❑ Query Optimizer decides which is executed
 - ❑ Based on cost model, statistics and other parameters
- ❑ Alternatives
 - ❑ Scan "country" and "gender", positional AND
 - ❑ Scan over "country" and probe into "gender"
 - ❑ Indices might be used
 - ❑ Decision depends on data and query parameters like e.g. selectivity

```
SELECT fname, lname FROM world_population
WHERE country="Italy" and gender="m"
```

59



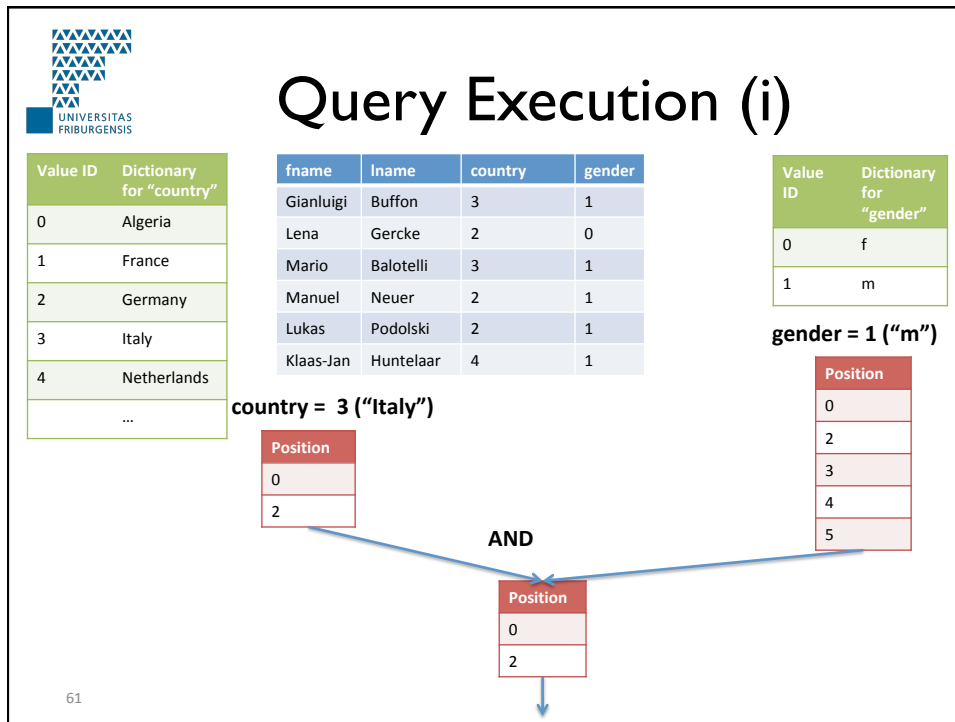
Query Plan (i)



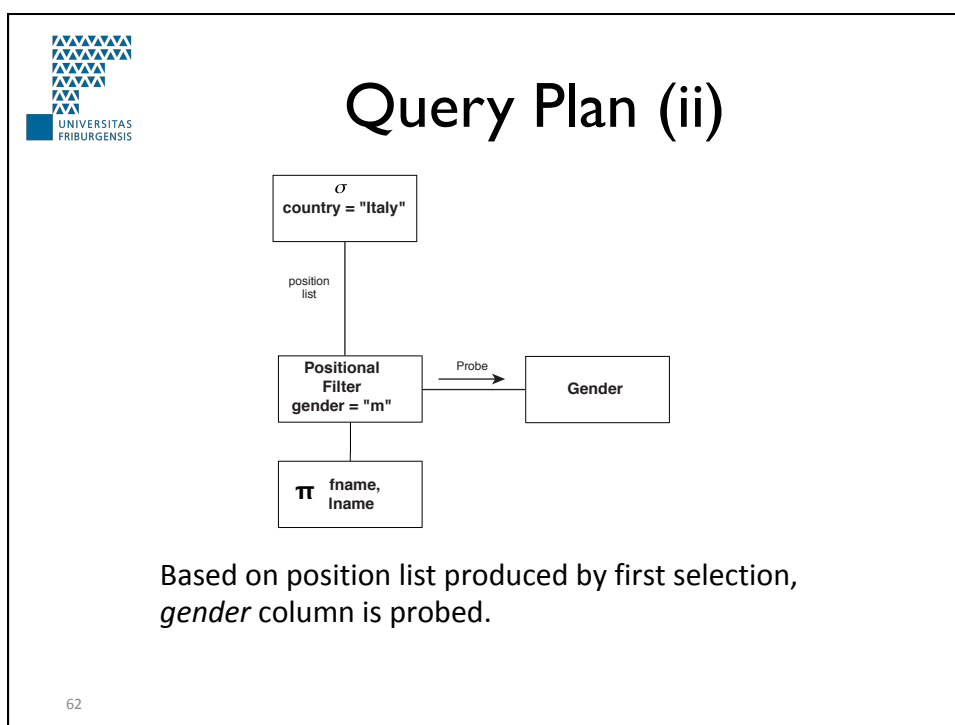
Positional AND:

- ❑ Predicates are evaluated and generate position lists
- ❑ Intermediate position lists are logically combined
- ❑ Final position list is used for materialization

60



61



62



Insert



Insert

- Insert is the dominant modification operation (Delete/Update can be modeled as Inserts as well)
- Inserting into a compressed in-memory persistence can be expensive
 - Updating sorted sequences (e.g. dictionaries) is a challenge
 - Inserting into columnar storages is generally more expensive than inserting into row storages



Insert Example

world_population

rowID	fname	lname	gender	country	city	birthday
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Sophie	Schulze	f	GER	Potsdam	09-03-1977
...

INSERT INTO world_population
VALUES (Karen, Schulze, f, GER, Rostock, 11-15-2012)

65



INSERT (I) w/o new Dictionary entry

INSERT INTO world_population VALUES (Karen, **Schulze**, f, GER, Rostock, 11-15-2012)

AV	D
0	0
1	1
2	3
3	2
4	3

	fname	lname	gender	country	city	birthday
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Sophie	Schulze	f	GER	Potsdam	09-03-1977
...

Attribute Vector (AV)
Dictionary (D)

66



INSERT (I) w/o new Dictionary entry

INSERT INTO world_population VALUES (Karen, **Schulze**, f, GER, Rostock, , 11-15-2012)

AV	D
0 0	0 Albrecht
1 1	1 Berg
2 3	2 Meyer
3 2	3 Schulze
4 3	

	fname	Iname	gender	country	city	birthday
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Sophie	Schulze	f	GER	Potsdam	09-03-1977
...

1. Look-up on D → entry found

Attribute Vector (AV)
Dictionary (D)

67



INSERT (I) w/o new Dictionary entry

INSERT INTO world_population VALUES (Karen, **Schulze**, f, GER, Rostock, , 11-15-2012)

AV	D
0 0	0 Albrecht
1 1	1 Berg
2 3	2 Meyer
3 2	3 Schulze
4 3	
5 3	

	fname	Iname	gender	country	city	birthday
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Sophie	Schulze	f	GER	Potsdam	09-03-1977
5		Schulze				
...

1. Look-up on D → entry found
2. Append ValueID to AV

Attribute Vector (AV)
Dictionary (D)

68



INSERT (2) with new Dictionary Entry

INSERT INTO world_population VALUES (Karen, Schulze, f, GER, Rostock, 11-15-2012)

AV		D
0	2	0 Anton
1	3	1 Hanna
2	1	2 Martin
3	0	3 Michael
4	4	4 Sophie

	fname	lname	gender	country	city	birthday
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Sophie	Schulze	f	GER	Potsdam	09-03-1977
5		Schulze			Rostock	
...

Attribute Vector (AV)
Dictionary (D)

1. Look-up on D → **no** entry found

69



INSERT (2) with new Dictionary Entry

INSERT INTO world_population VALUES (Karen, Schulze, f, GER, Rostock, 11-15-2012)

AV		D
0	2	0 Anton
1	3	1 Hanna
2	1	2 Karen
3	0	3 Martin
4	4	4 Michael
		5 Sophie

	fname	lname	gender	country	city	birthday
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Sophie	Schulze	f	GER	Potsdam	09-03-1977
5		Schulze			Rostock	
...

Attribute Vector (AV)
Dictionary (D)

1. Look-up on D → **no** entry found
2. Insert new value to D

70



INSERT (2) with new Dictionary Entry

INSERT INTO world_population VALUES (Karen, Schulze, f, GER, Rostock, 11-15-2012)

AV	D (old)	D (new)
0 2	0 Anton	0 Anton
1 3	1 Hanna	1 Hanna
2 1	2 Martin	2 Karen
3 0	3 Michael	3 Martin
4 4	4 Sophie	4 Michael
		5 Sophie

Look-up on D → **no entry found**

Insert new value to D

fname	lname	gender	country	city	birthday	
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Sophie	Schulze	f	GER	Potsdam	09-03-1977
5		Schulze			Rostock	
...

Attribute Vector (AV)
Dictionary (D)

1. Look-up on D → **no entry found**
2. Insert new value to D

71



INSERT (2) with new Dictionary Entry

INSERT INTO world_population VALUES (Karen, Schulze, f, GER, Rostock, 11-15-2012)

AV (old)		AV (new)		D (new)								
0	2	0	3	0	Anton		fname	lname	gender	country	city	birthday
1	3	1	4	1	Hanna	0	Martin	Albrecht	m	GER	Berlin	08-05-1955
2	1	2	1	2	Karen	1	Michael	Berg	m	GER	Berlin	03-05-1970
3	0	3	0	3	Martin	2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
4	4	4	5	4	Michael	3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
				5	Sophie	4	Sophie	Schulze	f	GER	Potsdam	09-03-1977
						5		Schulze			Rostock	
					

Look-up on D → **no** entry found

Insert new value to D

Attribute Vector (AV)
Dictionary (D)

1. Look-up on D → **no entry found**
2. Insert new value to D
3. Change ValueIDs in AV

72



INSERT (2) with new Dictionary Entry

INSERT INTO world_population VALUES (Karen, Schulze, f, GER, Rostock, 11-15-2012)

Changed Value IDs	AV	D
→ 0	3	0 Anton
→ 1	4	1 Hanna
2	1	2 Karen
3	0	3 Martin
→ 4	5	4 Michael
5	2	5 Sophie

	fname	Iname	gender	country	city	birthday
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Sophie	Schulze	f	GER	Potsdam	09-03-1977
5	Karen	Schulze			Rostock	
...

Attribute Vector (AV)
Dictionary (D)

1. Look-up on D → **no** entry found
2. Insert new value to D
3. Change ValueIDs in AV
4. Append new ValueID to AV

73



RESULT

world_population

rowID	fname	Iname	gender	country	city	birthday
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Ulrike	Schulze	f	GER	Potsdam	09-03-1977
5	Karen	Schulze	f	GER	Rostock	11-15-2012

INSERT INTO world_population
VALUES (Karen, Schulze, f, GER, Rostock, 11-15-2012)

74



Advanced Database Storage Techniques



Differential Buffer



Motivation

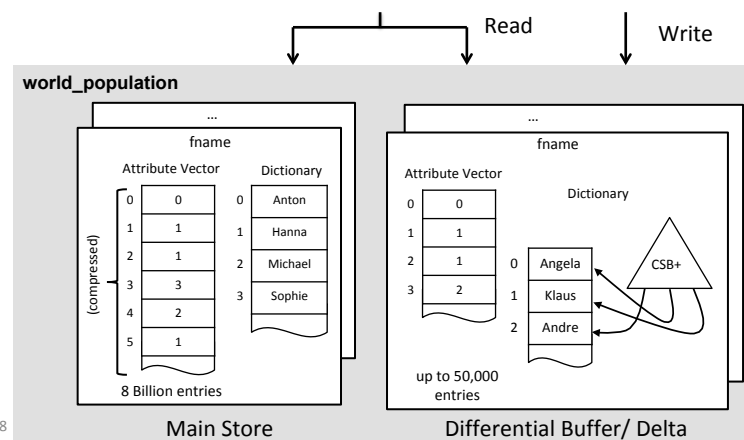
- Inserting new tuples directly into a compressed structure can be expensive
 - Especially when using sorted structures
 - New values can require reorganizing the dictionary
 - Number of bits required to encode all dictionary values can change, attribute vector has to be reorganized

77



Differential Buffer

- New values are written to a dedicated differential buffer (Delta)
- Cache Sensitive B+ Tree (CSB+) used for fast search on Delta



78



Differential Buffer

- Inserts of new values are fast, because dictionary and attribute vector do not need to be resorted
- Range selects on differential buffer are expensive
 - Unsorted dictionary allows no direct comparison of values
 - Scans with range selection need to lookup values in dictionary for comparisons
- Differential Buffer requires more memory:
 - No attribute vector compression
 - Additional CSB+ Tree for dictionary

79



Tuple Lifetime

Michael moves from Berlin to Potsdam

Main Table: world_population

recid	fname	lname	gender	country	city	birthday
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Ulrike	Schulze	f	GER	Potsdam	09-03-1977
5	Sophie	Schulze	f	GER	Rostock	06-20-2012
...
8 * 10 ⁹	Zacharias	Perdopolus	m	GRE	Athen	03-12-1979

Main Store

Differential Buffer

```
UPDATE 'world_population'
SET city='Potsdam'
WHERE fname="Michael" AND lname="Berg"
```

80



Tuple Lifetime

Michael moves from Berlin to Potsdam

Main Table: world_population

recid	fname	lname	gender	country	city	birthday
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Ulrike	Schulze	f	GER	Potsdam	09-03-1977
5	Sophie	Schulze	f	GER	Rostock	06-20-2012
...
8 * 10 ⁹	Zacharias	Perdopolus	m	GRE	Athen	03-12-1979

Main Store

Differential Buffer

```
UPDATE 'world_population'
SET city='Potsdam'
WHERE fname="Michael" AND lname="Berg"
```

81



Tuple Lifetime

Michael moves from Berlin to Potsdam

Main Table: world_population

recid	fname	lname	gender	country	city	birthday
0	Martin	Albrecht	m	GER	Berlin	08-05-1955
1	Michael	Berg	m	GER	Berlin	03-05-1970
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992
4	Ulrike	Schulze	f	GER	Potsdam	09-03-1977
5	Sophie	Schulze	f	GER	Rostock	06-20-2012
...
8 * 10 ⁹	Zacharias	Perdopolus	m	GRE	Athen	03-12-1979
0	Michael	Berg	m	GER	Potsdam	03-05-1970

Main Store

Differential Buffer

```
UPDATE 'world_population'
SET city='Potsdam'
WHERE fname="Michael" AND lname="Berg"
```

82



Tuple Lifetime

- Tuples are now available in Main Store and Differential Buffer
- Tuples of a table are marked by a validity vector to reduce the required amount of reorganization steps
 - Like an attribute vector for validity
 - 1 bit required per database tuple
- Invalidated tuples stay in the database table, until the next reorganization takes place
- Query results
 - Main and delta have to be queried
 - Results are filtered using the validity vector

83



Tuple Lifetime

Michael moves from Berlin to Potsdam

Main Table: world_population

recid	fname	lname	gender	country	city	birthday	valid
0	Martin	Albrecht	m	GER	Berlin	08-05-1955	1
1	Michael	Berg	m	GER	Berlin	03-05-1970	0
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968	1
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992	1
4	Ulrike	Schulze	f	GER	Potsdam	09-03-1977	1
5	Sophie	Schulze	f	GER	Rostock	06-20-2012	1
...
8 * 10 ⁹	Zacharias	Perdopolus	m	GRE	Athen	03-12-1979	1
0	Michael	Berg	m	GER	Potsdam	03-05-1970	1

Main Store

Differential Buffer

```
UPDATE 'world_population'
SET city='Potsdam'
WHERE fname="Michael" AND lname="Berg"
```

84



Tuple Lifetime

Michael moves from Berlin to Potsdam

Main Table: world_population

recId	fname	lname	gender	country	city	birthday	valid
0	Martin	Albrecht	m	GER	Berlin	08-05-1955	1
1	Michael	Berg	m	GER	Berlin	03-05-1970	0
2	Hanna	Schulze	f	GER	Hamburg	04-04-1968	1
3	Anton	Meyer	m	AUT	Innsbruck	10-20-1992	1
4	Ulrike	Schulze	f	GER	Potsdam	09-03-1977	1
5	Sophie	Schulze	f	GER	Rostock	06-20-2012	1
...
8 * 10 ⁹	Zacharias	Perdopolus	m	GRE	Athen	03-12-1979	1
0	Michael	Berg	m	GER	Potsdam	03-05-1970	1

Main Store

Differential Buffer

```
UPDATE 'world_population'
SET city='Potsdam'
WHERE fname="Michael" AND lname="Berg"
```

85



Merge



Handling Write Operations

- All Write operations (INSERT, UPDATE) are stored within a differential buffer (delta) first
- Read-operations on differential buffer are more expensive than on main store
- Differential buffer is merged periodically with the main store
 - To avoid performance degradation based on large delta
 - Merge is performed asynchronously

87



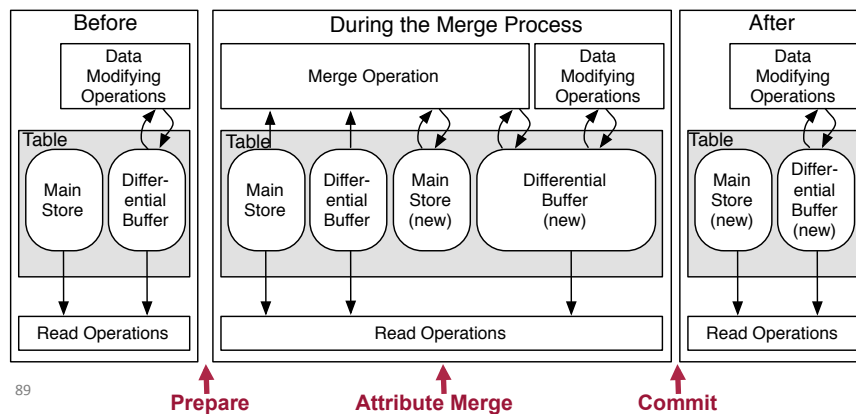
Merge Overview I/II

- The merge process is triggered for single tables
- Is triggered by:
 - Amount of tuples in buffer
 - Cost model to
 - Schedule
 - Take query cost into account
 - Manually

88

Merge Overview II/II

- Working on data copies allows asynchronous merge
- Very limited interruption due to short lock
- At least twice the memory of the table needed!



89

Attribute Merge

Step 1: Dictionary Merge

- Merge main and delta dictionary
 - Optionally remove unused values
 - Merge of two sorted arrays
- Create mapping if valueIDs changed

Step 2: Update Attribute Vector

- Create new merged main partition
- Update valueIDs reflecting changed dictionary

90



Example

Main Store

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Albert	Berlin	0	2	0	0	1
1	Michael	London	1	1	1	1	1
2	Nadja		2	0	0	2	1

Delta

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid

91



Example

Michael moves from London to Berlin

Main Store

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Albert	Berlin	0	2	0	0	1
1	Michael	London	1	1	1	1	0
2	Nadja		2	0	0	2	1

Delta

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Michael	Berlin	0	0	0	0	1

92



Example

Nadja moves from Berlin to Potsdam

Main Store

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Albert	Berlin	0	2	0	0	0
1	Michael	London	1	1	1	1	0
2	Nadja		2	0	0	2	1

Delta

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Michael	Berlin	0	0	0	0	1
1	Nadja	Potsdam	1	1	1	1	1

93



Example

Michael moves from Berlin to Potsdam

Main Store

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Albert	Berlin	0	2	0	0	0
1	Michael	London	1	1	1	1	0
2	Nadja		2	0	0	2	1

Delta

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Michael	Berlin	0	0	0	0	0
1	Nadja	Potsdam	1	1	1	1	1
			2	0	1	2	1

94



First Step: Validity Detection

Main Store

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Albert	Berlin	0	2	0	0	0
1	Michael	London	1	1	1	1	0
2	Nadja		2	0	0	2	1

Delta

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Michael	Berlin	0	0	0	0	0
1	Nadja	Potsdam	1	1	1	1	1
			2	0	1	2	1

Will be deleted / moved into history partition

95



First Step: Dictionary Merge

Main Store

Dictionaries		
valueID	fname	city
0	Albert	Berlin
1	Michael	London
2	Nadja	

New Combined Main Store

Dictionaries		
valueID	fname	city
0	Albert	Berlin
1	Michael	Potsdam
2	Nadja	

Delta

Dictionaries		
valueID	fname	city
0	Michael	Berlin
1	Nadja	Potsdam

Mappings

fname: old to new Main

valueID (old)	0	1	2
valueID (new)	0	1	2

fname: Delta to Main

valueID (Delta)	0	1
valueID (Main)	1	2

city: old to new Main

valueID (old)	0	1
valueID (new)	0	NA

city: Delta to Main

valueID (Delta)	0	1
valueID (Main)	0	1



Second Step: New Main Column

Main Store

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Albert	Berlin	0	0	0	0	1
1	Michael	London					
2	Nadja						

Delta

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Michael	Berlin	0	1	1	0	1
1	Nadja	Potsdam	1	0	1	1	1

Mappings

fname: old to new Main				fname: Delta to Main		city: old to new Main			city: Delta to Main			
valueID (old)	0	1	2	valueID (Delta)	0	1	valueID (old)	0	1	valueID (Delta)	0	1
valueID (new)	0	1	2	valueID (Main)	1	2	valueID (new)	0	NA	valueID (Main)	0	1



Second Step: New Main Column

Main Store

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid
0	Albert	Berlin	0	0	0	0	1
1	Michael	Potsdam	1	2	1	1	1
2	Nadja		2	1	1	2	1


Delta

Dictionaries			Attribute Vectors			Validity Vector	
valueID	fname	city	recID	fname	city	recID	valid



How does it all come together?

1. Mixed Workload combining OLTP and analytic-style queries
 - **Column-Stores** are best suited for analytic-style queries
 - **In-memory** databases enable fast tuple re-construction
 - In-memory column store allows aggregation **on-the-fly**
2. Sparse enterprise data
 - Lightweight **compression** schemes are optimal
 - Increases query execution
 - Improves feasibility of in-memory databases
3. Mostly read workload
 - Read-optimized stores provide best throughput
 - i.e. compressed in-memory column-store
 - Write-optimized store as delta partition to handle data changes is sufficient



99



Want More?

□ Master Thesis Topics

- **Distributed In-Memory Key-Value Store on Ultra-Low-Power Hardware**



Raspberry PI Cluster

- **Graph Analytics and User Tracking in In-Memory Databases**
 - Data analysis + query execution + compression und huge servers
 - E.g. 32-cores, 412 GB RAM
 - Based on HYRISE: <http://github.com/hyrise/hyrise>

100



Technical Deep-Dive in a Column-Oriented In-Memory Database

Slides by Martin Grund

Presented today by Alberto Lerner

alberto.lerner@unifr.ch

eXascale Infolab

Department of Informatics
University of Fribourg Switzerland